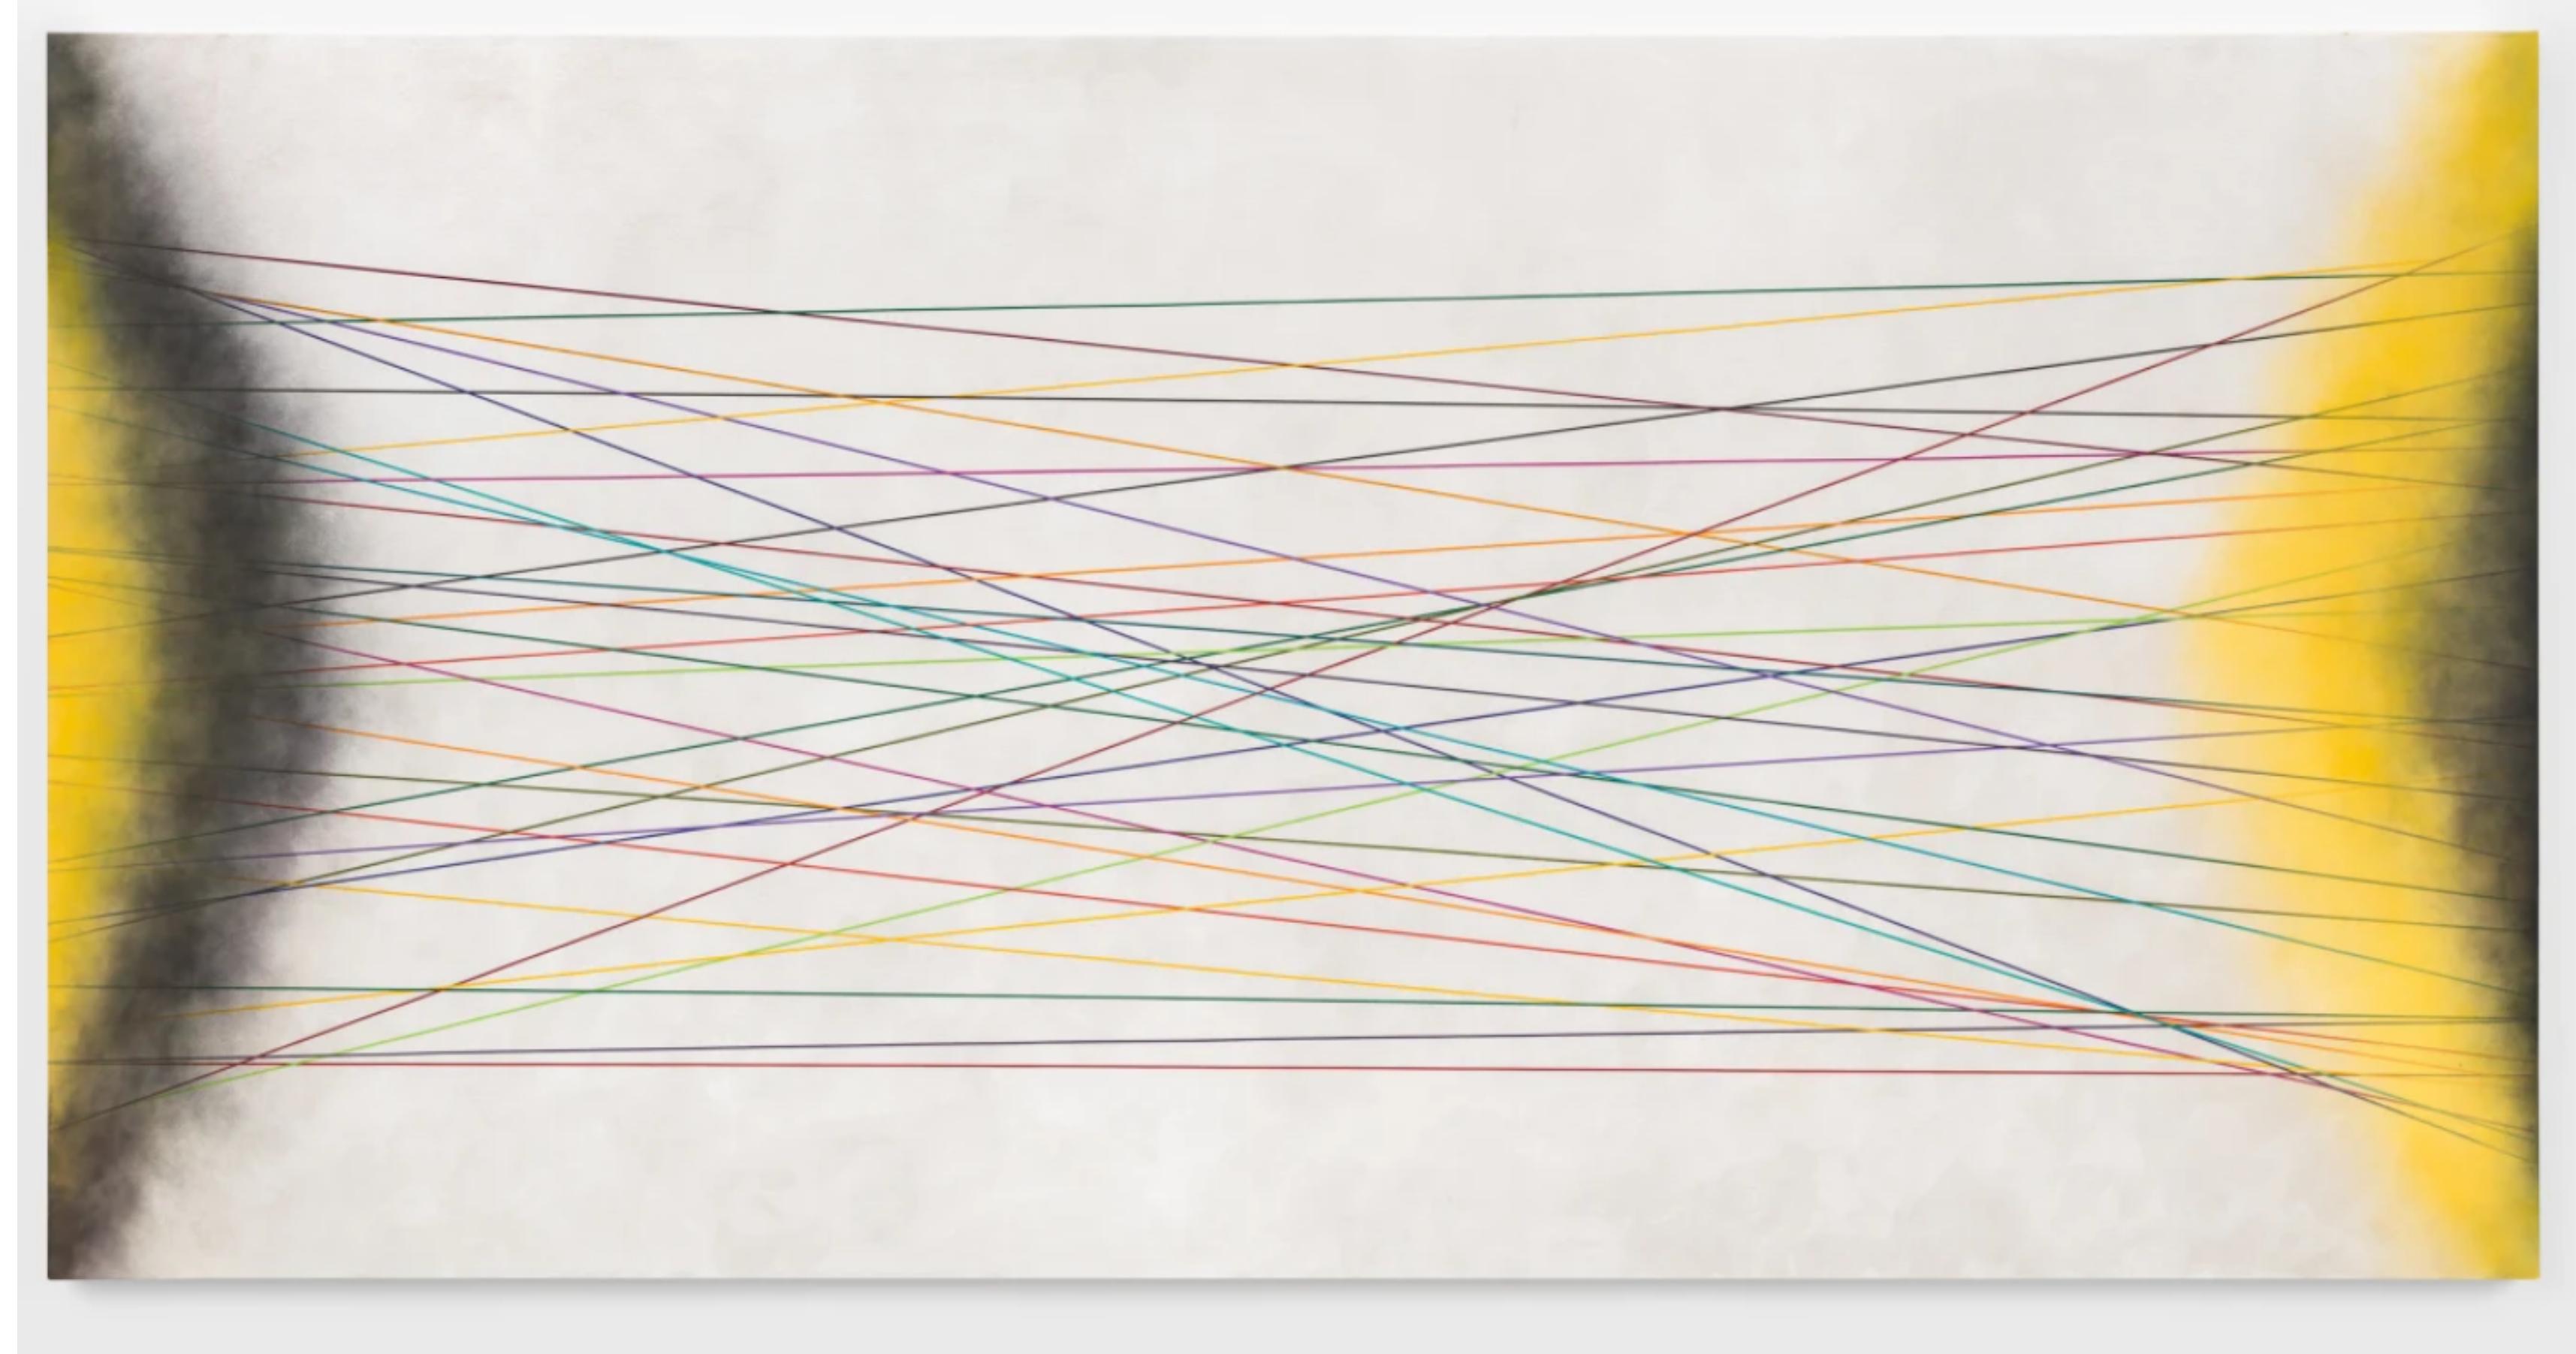


Optimization of derivation jobs and modernization of VO integration tests for the ATLAS Experiment

March 19th,
2025



Virginia Jamarillo, "Quanta", 2021

Jack Kraus

Advised by: Jahred Adelman (Northern Illinois University)
Serhan Mete, Peter van Gemmeren (Argonne National Labs)



An outline...

ATLAS Experiment Introduction

**Toy Model Investigation into TBranch
basket buffer sizes**

**Applying changes to derivation of real
AODs**

**Creating new unit tests covering new I/O
core functionality**

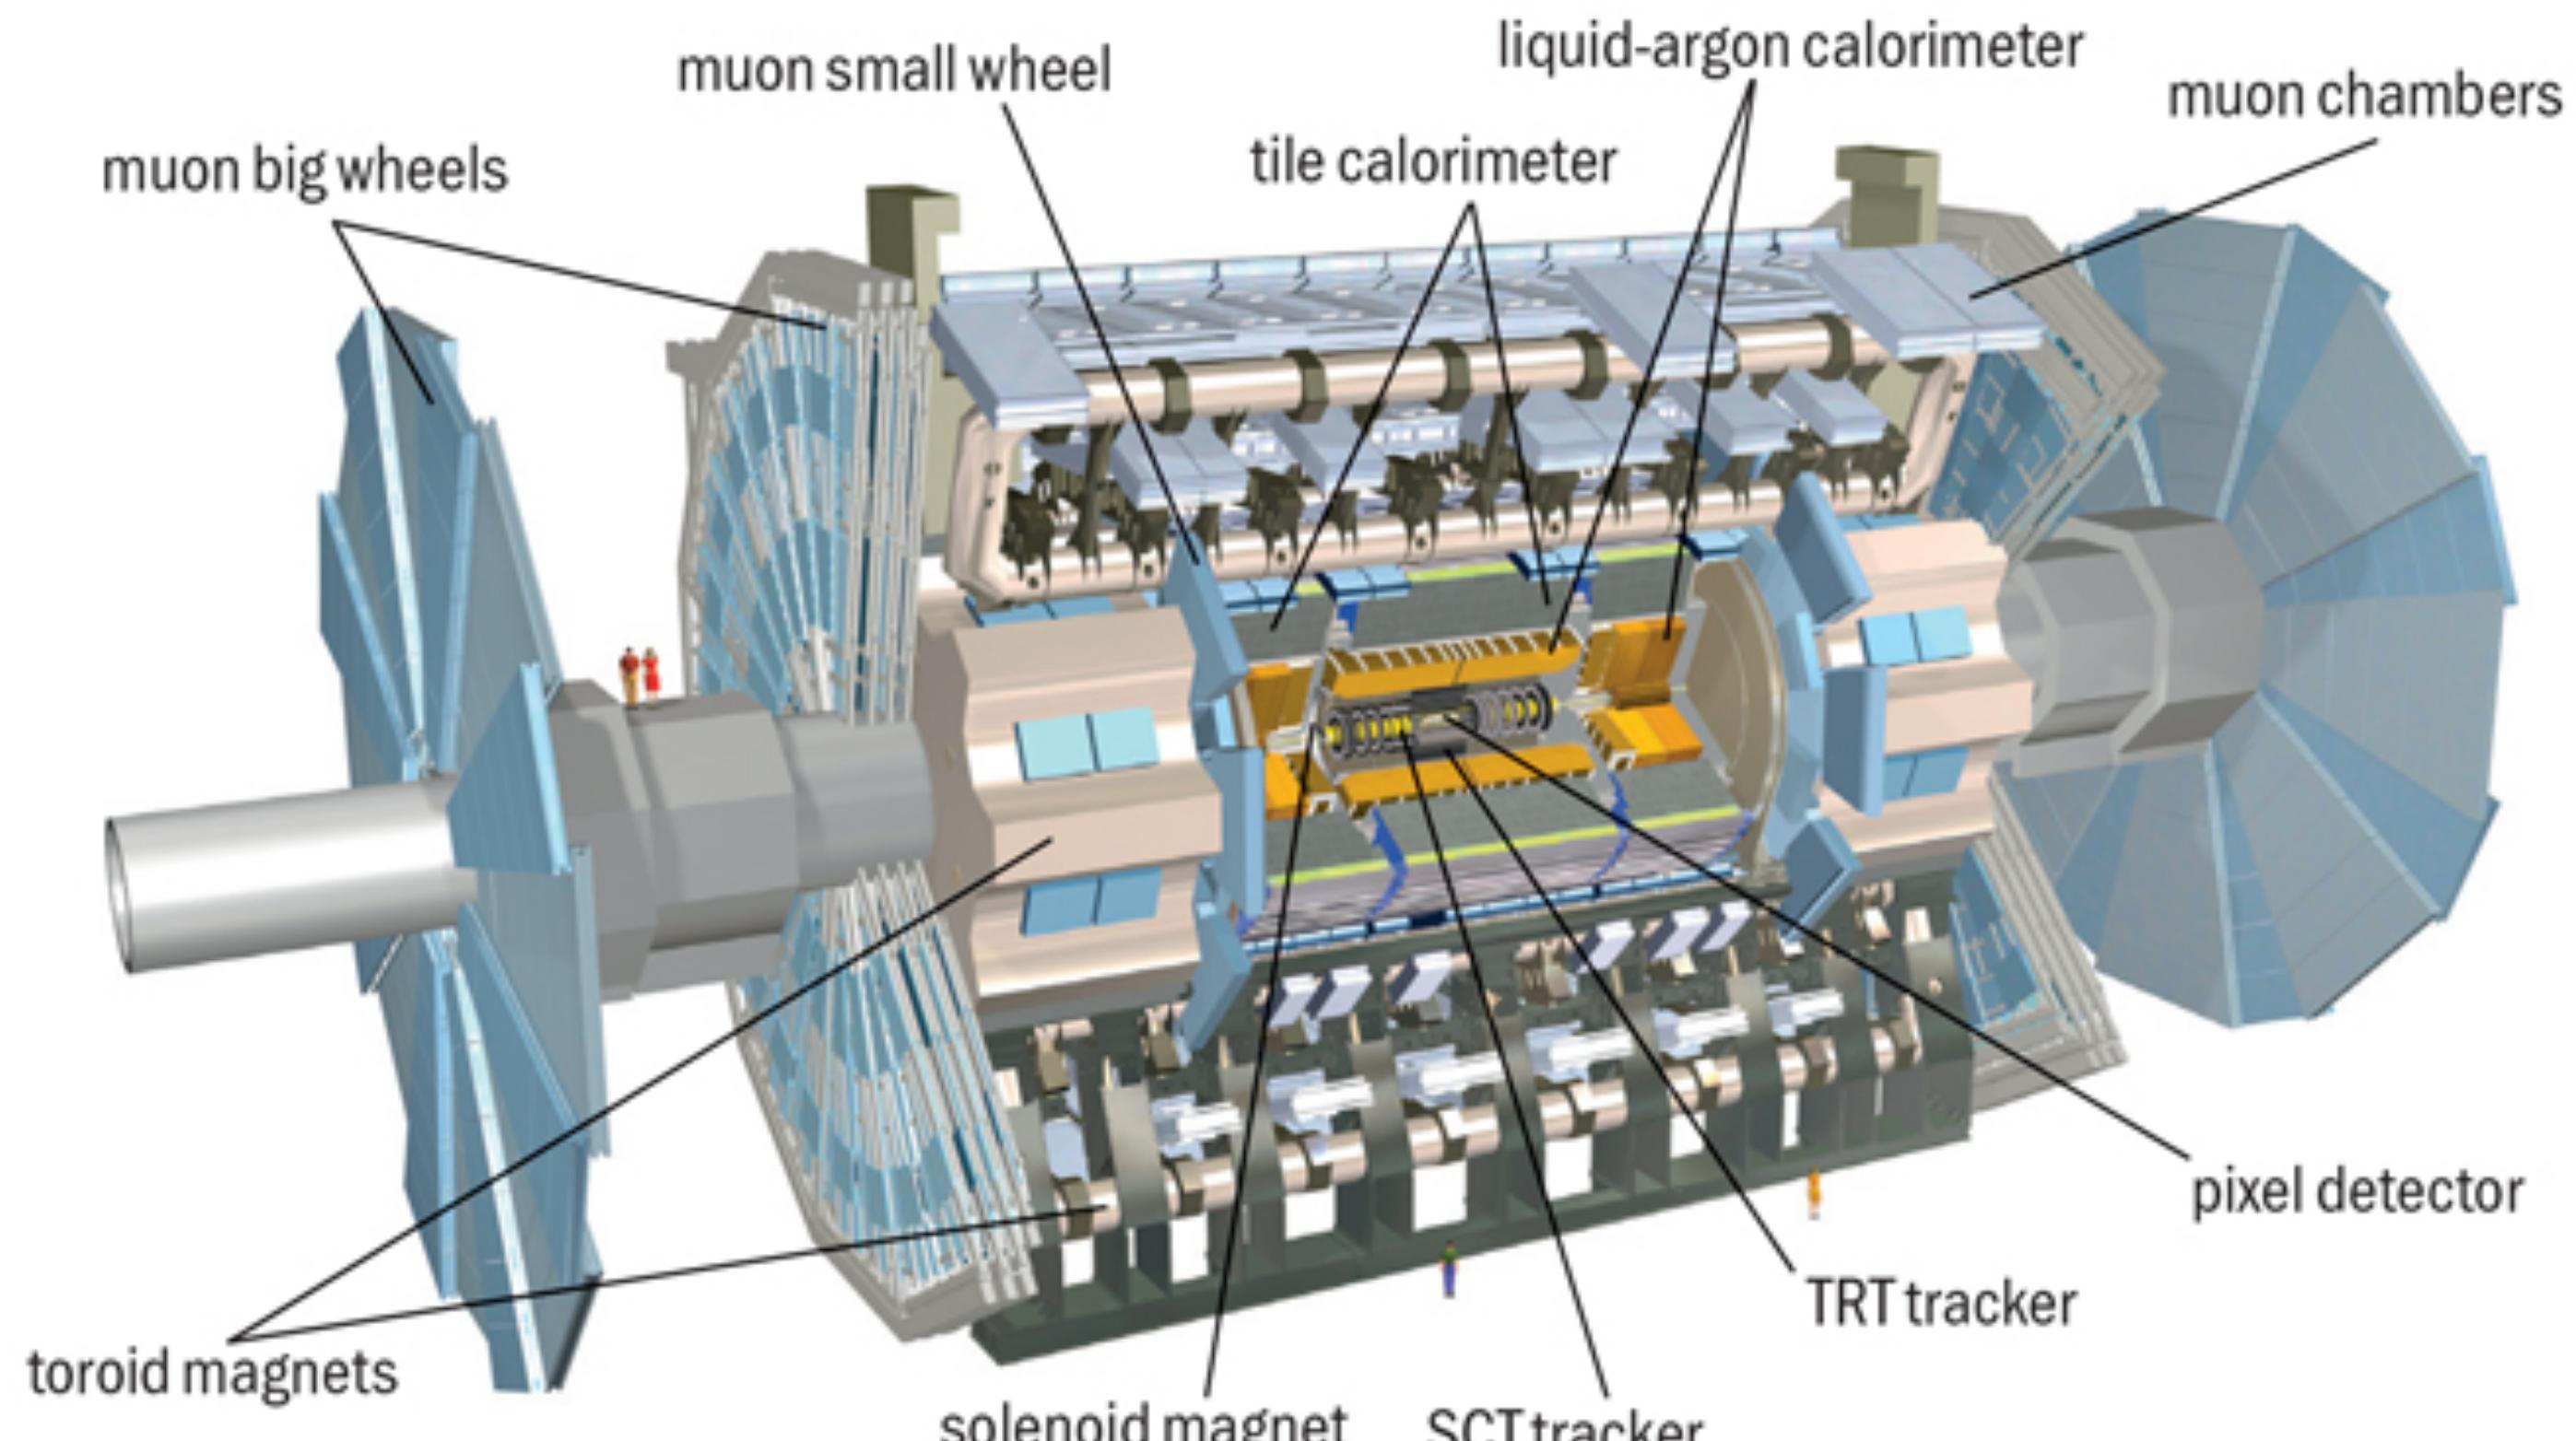
Starting with:

ATLAS Experiment Introduction

The ATLAS Detector

Inner Detector (ID)

- **Pixel detector** is the first point of contact, ionizes a one-sided silicon wafer to create an electron-hole pair drift, generating an electric current.
- **SemiConductor Tracker (SCT)** uses silicon readout strips, works with the pixel detector in measuring hits to build up tracks.
- **Transition radiation tracker (TRT)** is a collection of tubes filled with a combination of gasses to discriminate between heavier and lighter particles
- Measures direction, momenta, and charge of electrically charged particles.
- Reconstruction of **tracks** can be done from the various **hits** to each part of the ID



The ATLAS Detector

Calorimeters

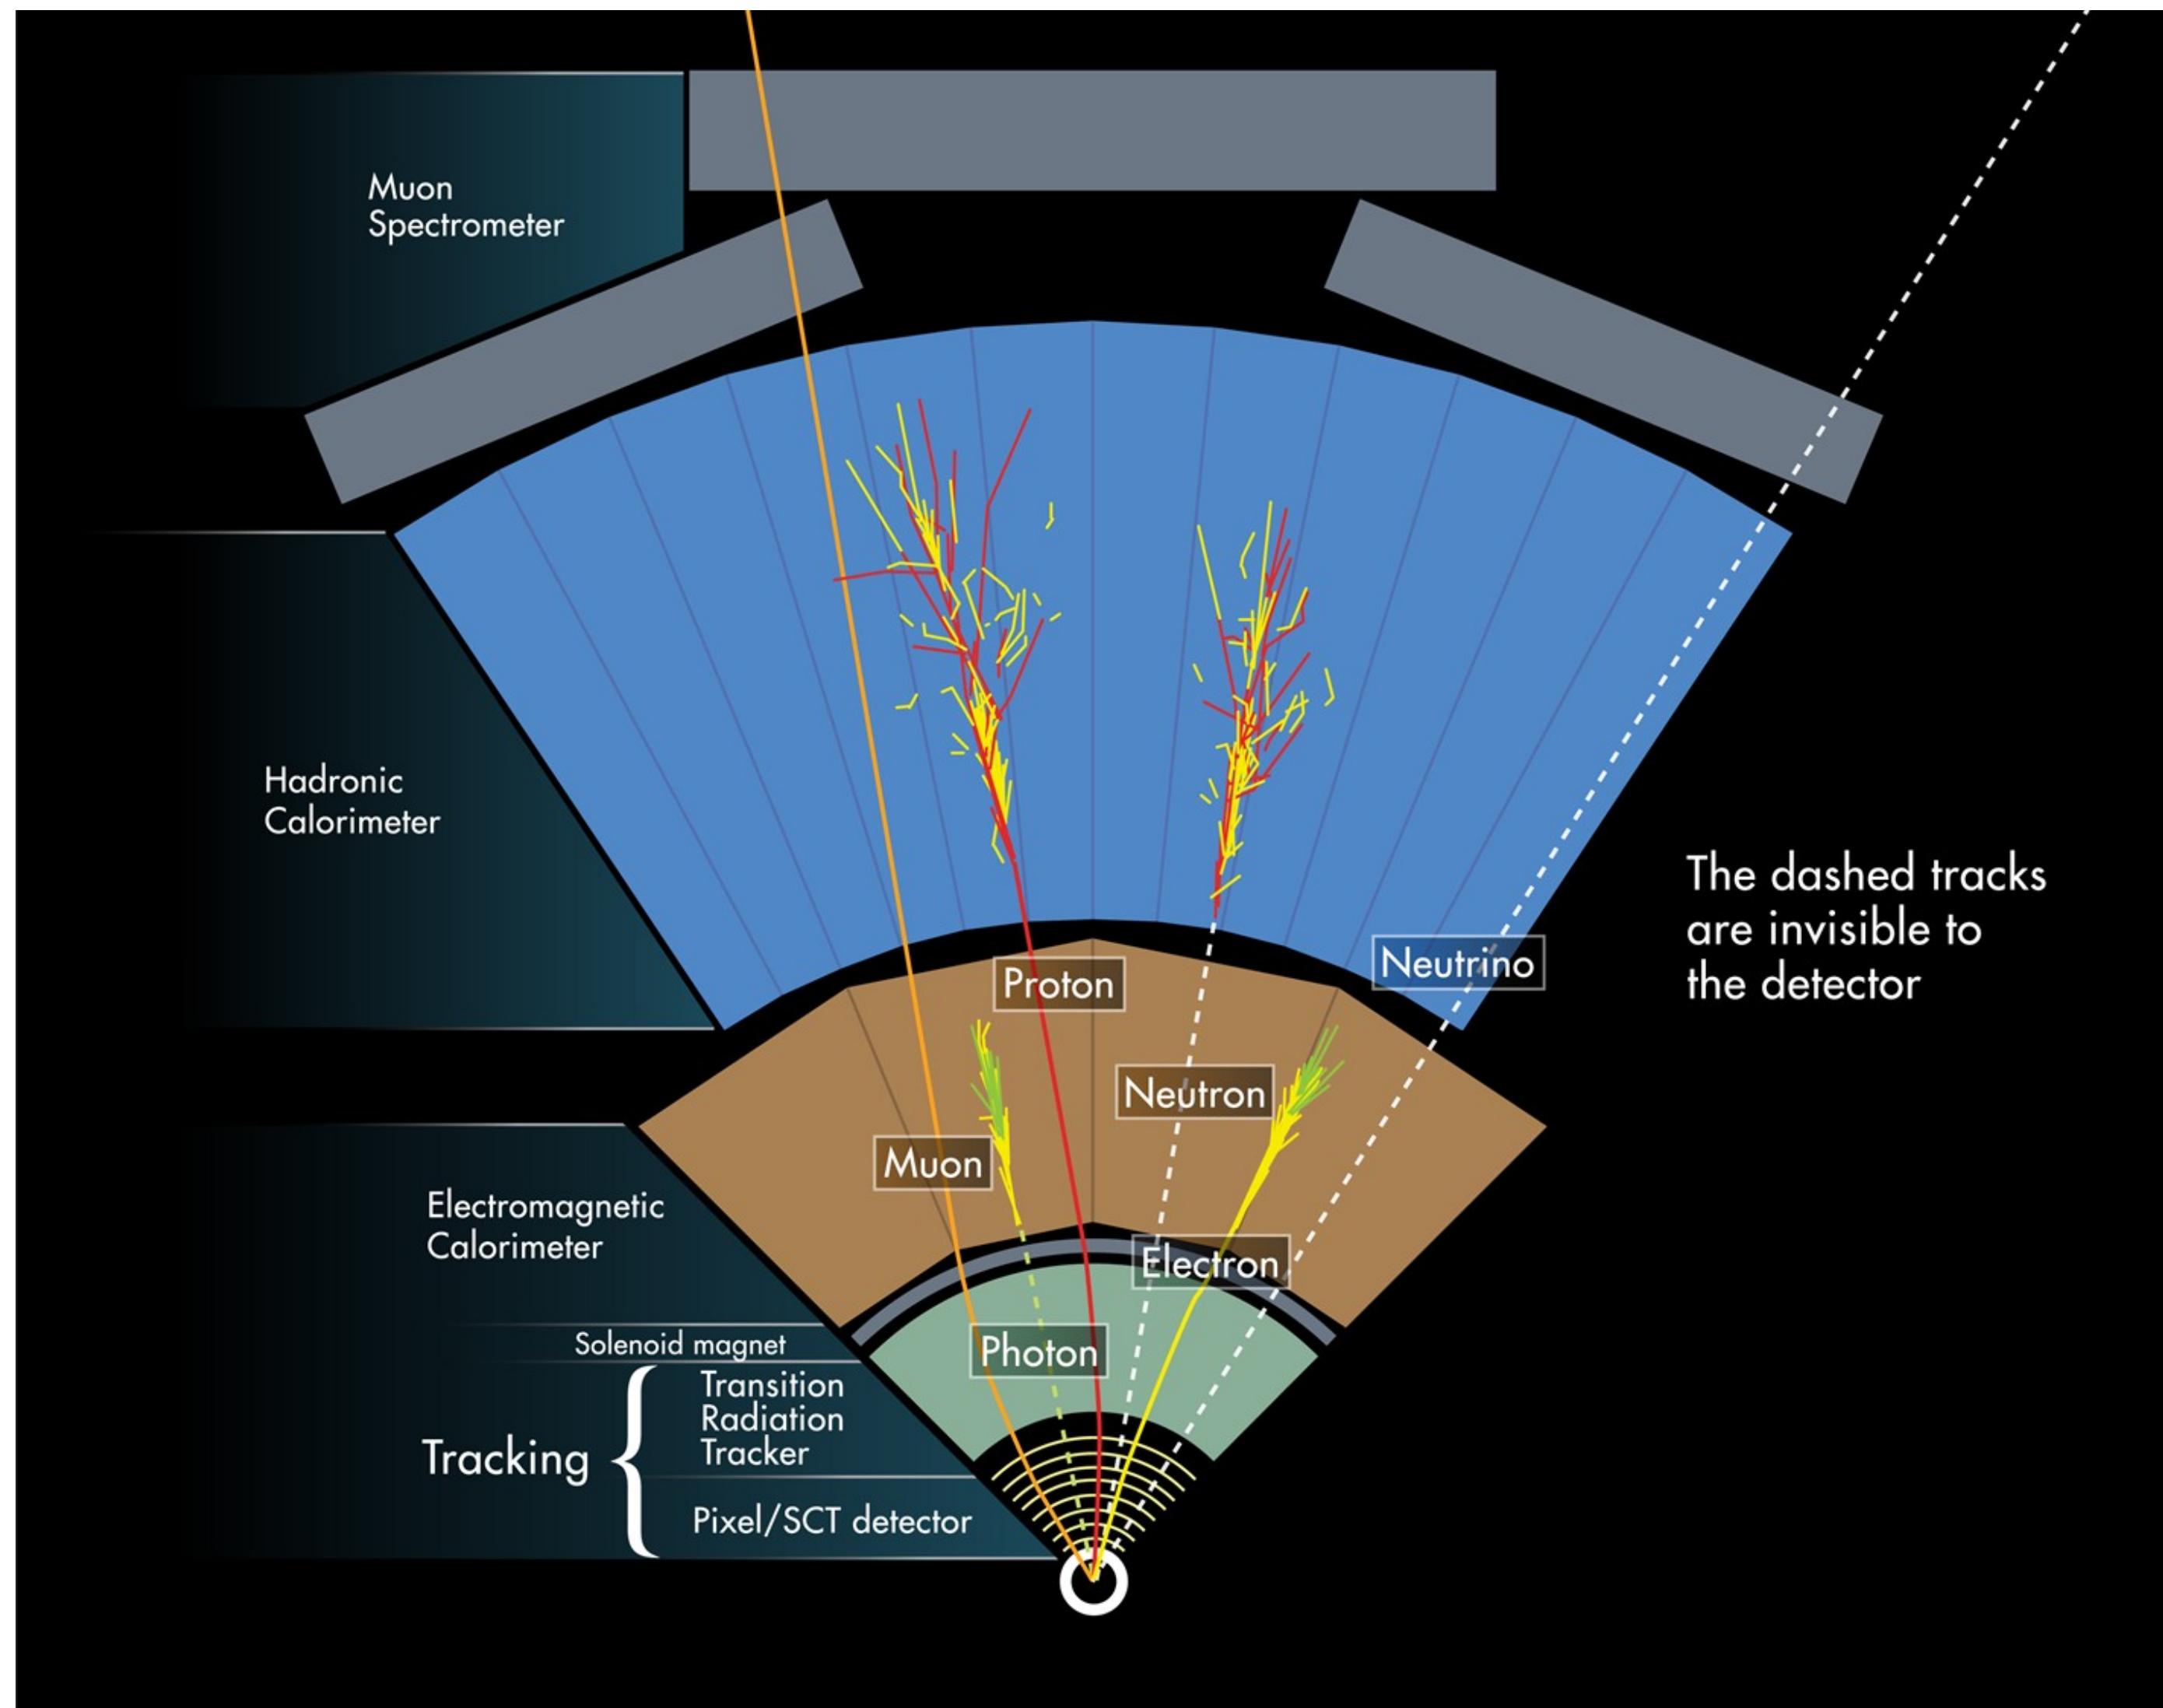
- Liquid Argon (LAr) calorimeter

- Surrounds the ID, measures the energy deposits of objects that interact via the electromagnetic force.

- Tile calorimeter

- It's a hadronic calorimeter, so it interacts with particles via the strong nuclear force.
- When a particle hits the steel, a cascade of secondary protons, neutrons and other hadrons (quark bound states) are produced.

Jets are produced here, along with reconstructed tracks in the ID, two or more tracks will be used to identify a **vertex**.



The ATLAS Detector

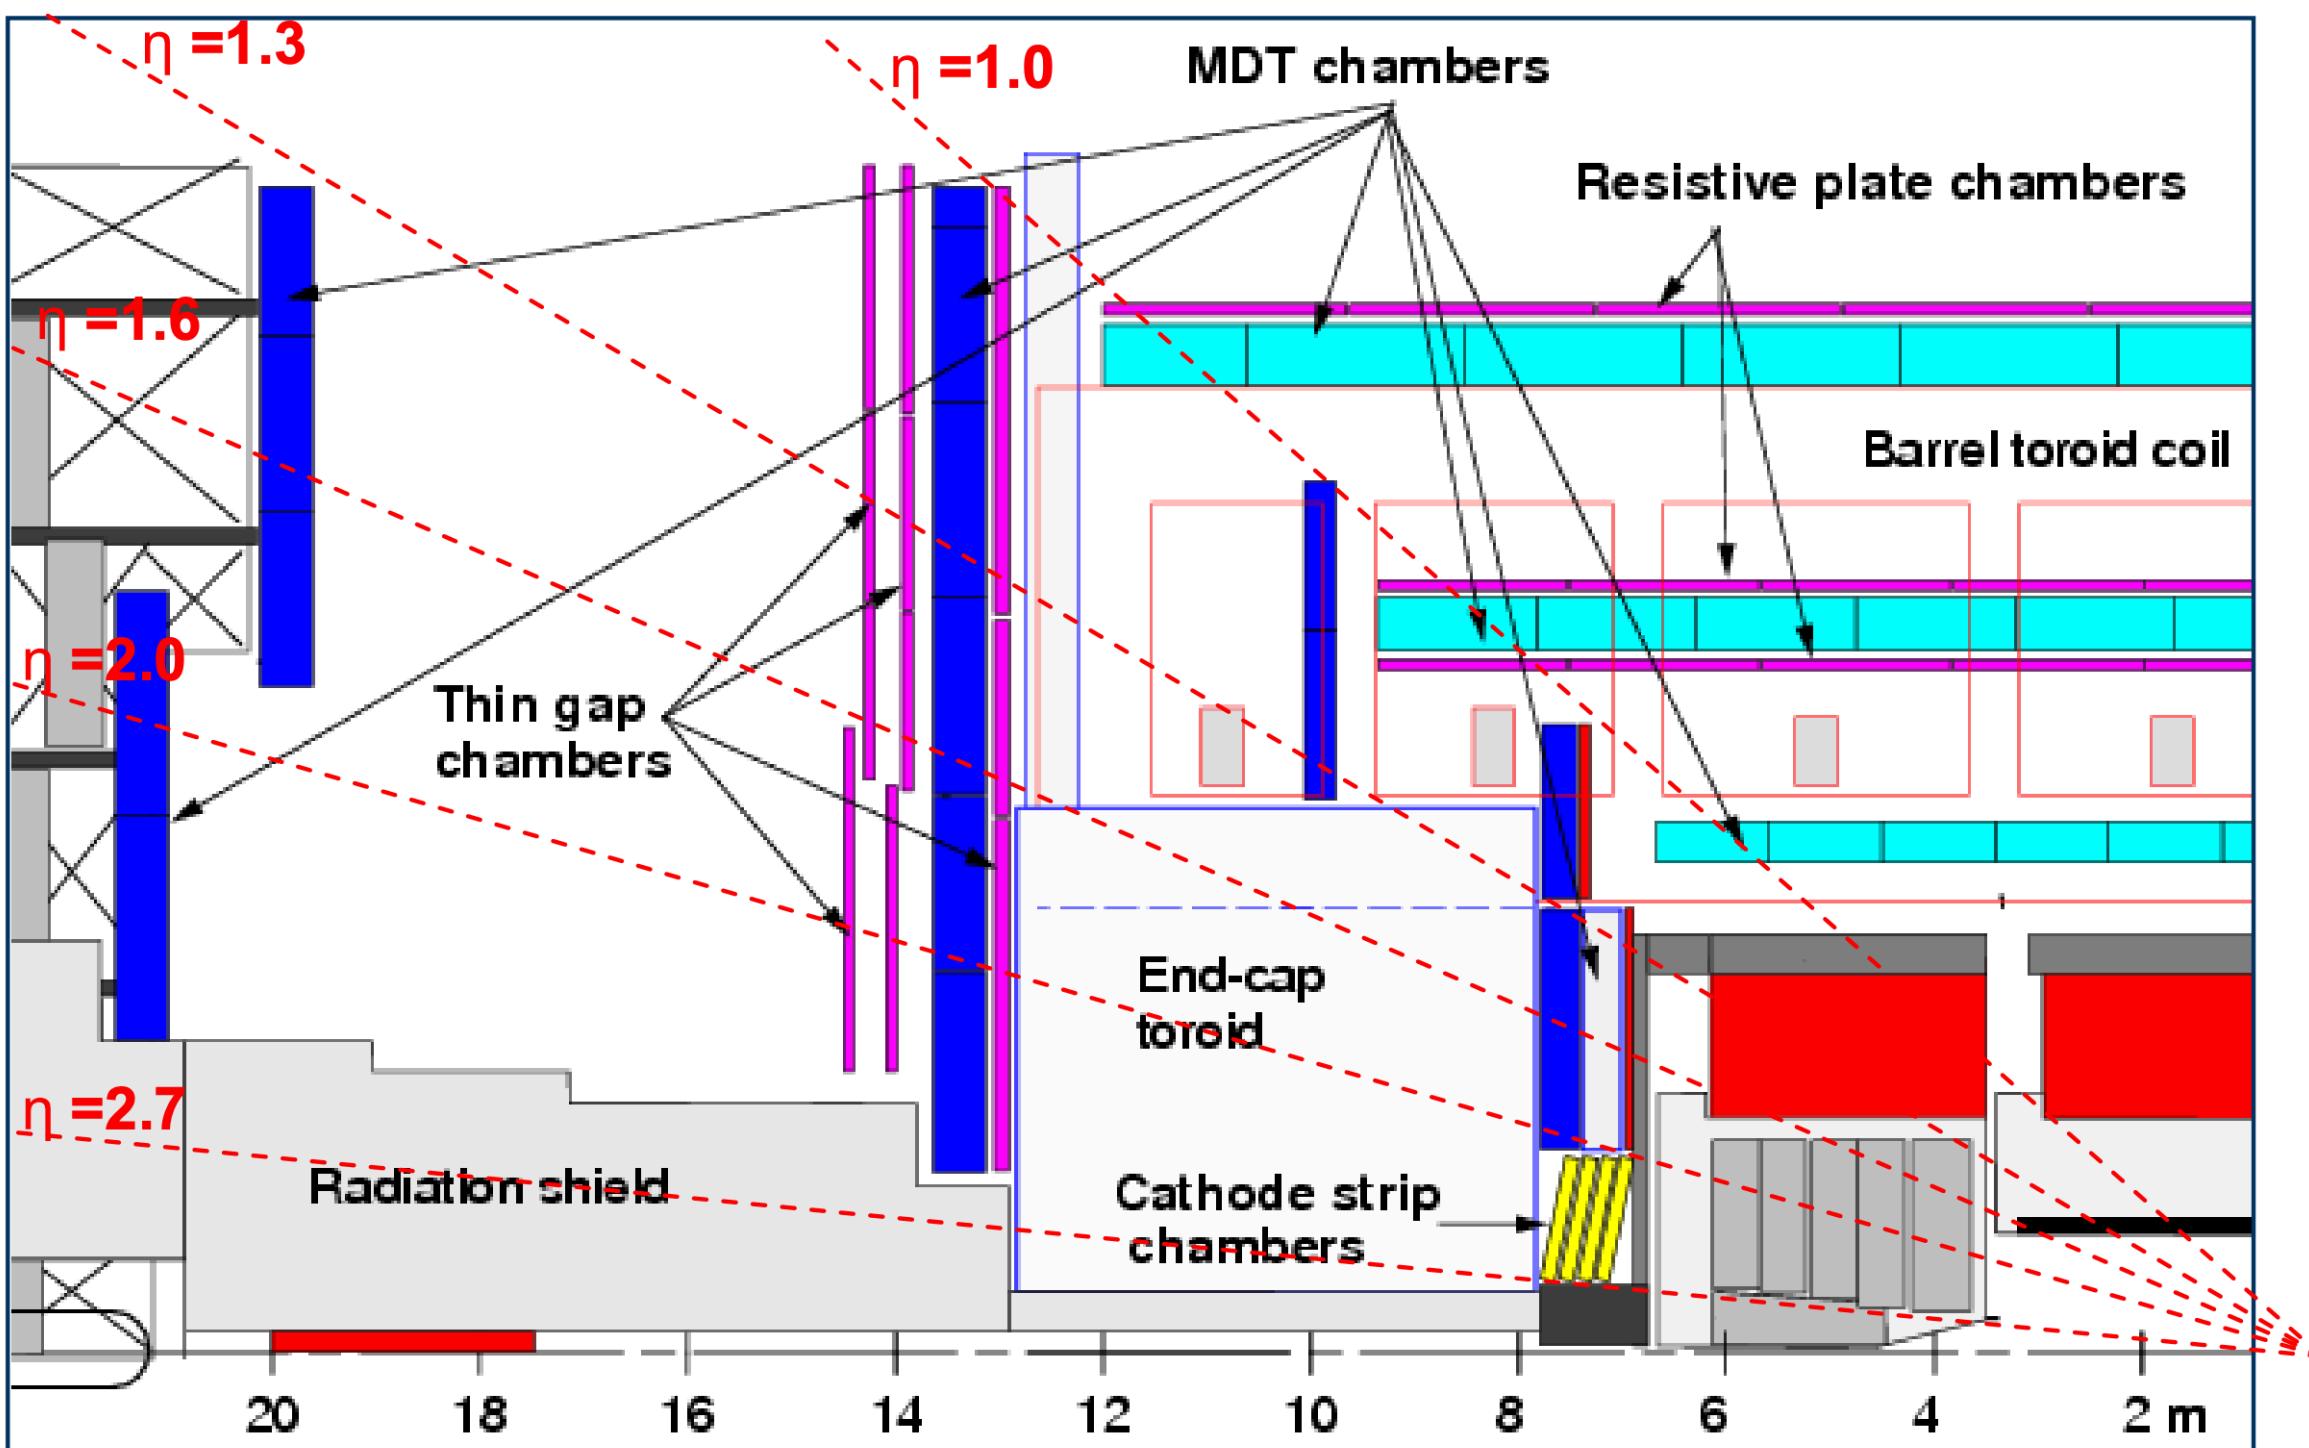
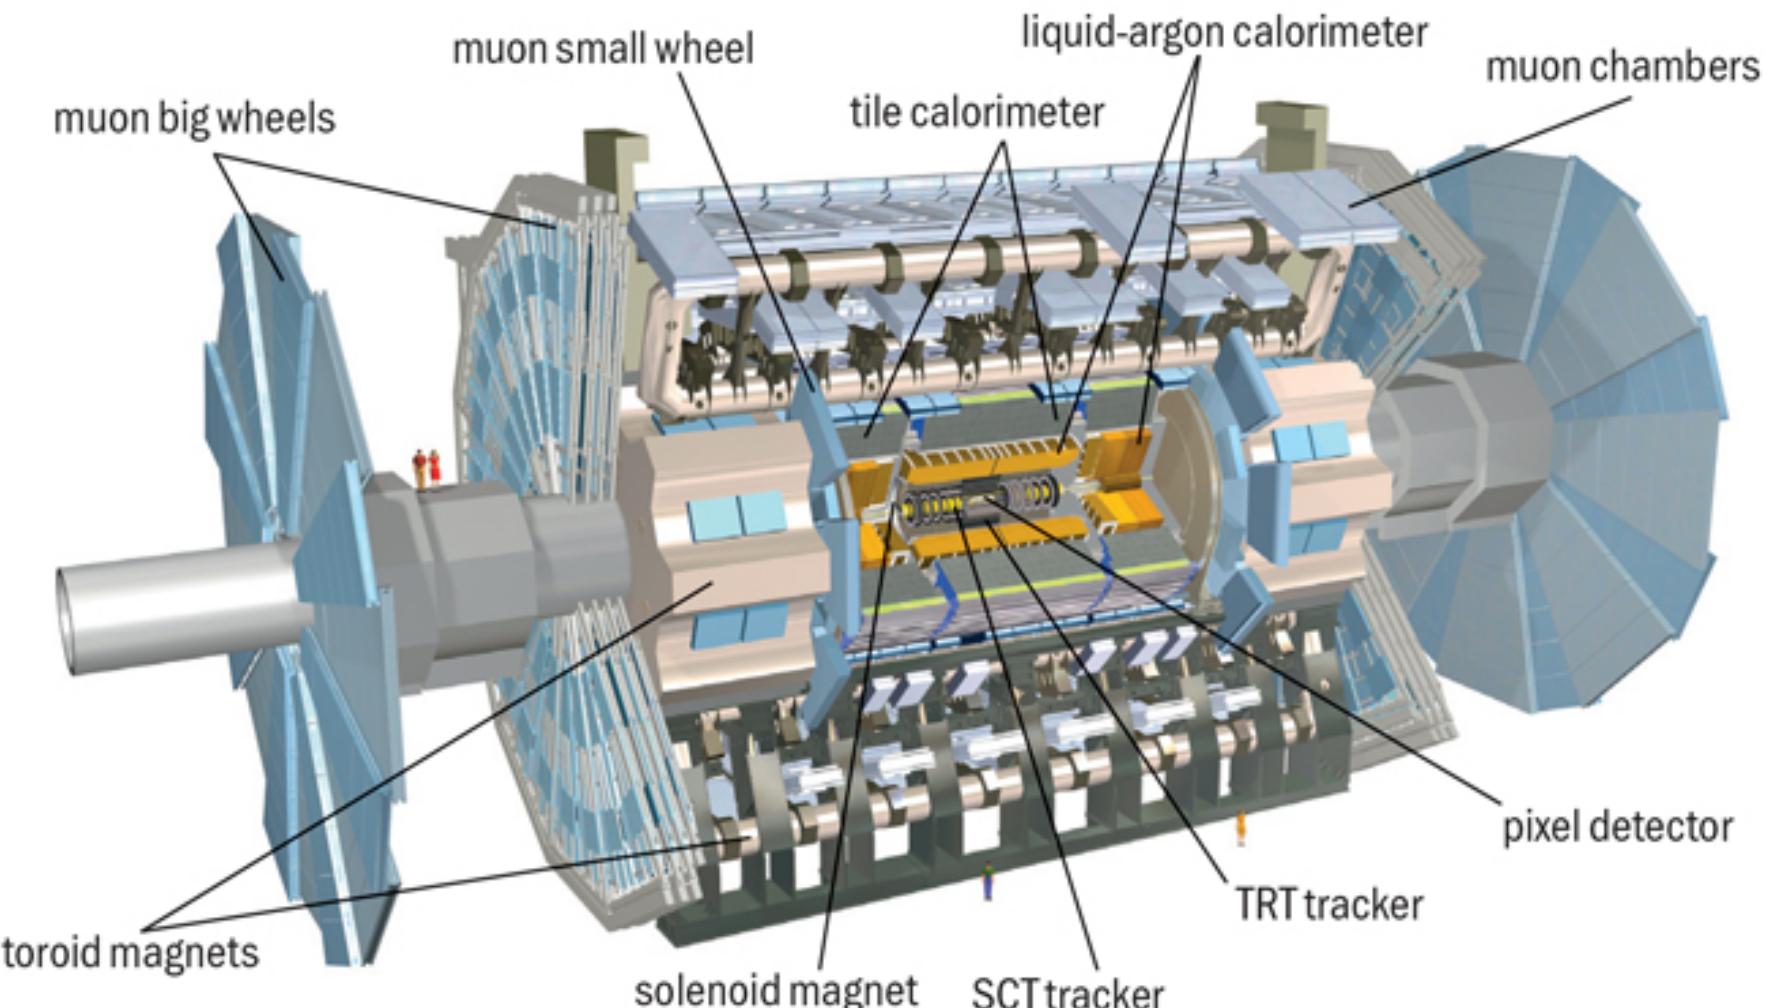
Muon Spectrometer (MS)

- Designed to identify muon tracks and momenta to high-res
- Monitored Drift Tubes (MDT) lie on the end caps and barrel regions, covering pseudo rapidity regions $0 < |\eta| < 2.7$

(Note: η can be expected to be stored as a number between 0 and infinity, but the area with highest flux is between $2 < |\eta| < 2.7$)

Also, η is chosen as a coordinate because it's difference in rapidity is Lorentz invariant under boosts along the beam axis, making it easier to ID tracks due to symmetry of the collision.

$$\eta \equiv -\ln \left[\tan \left(\frac{\theta}{2} \right) \right]$$



ATLAS Trigger/Data Acquisition

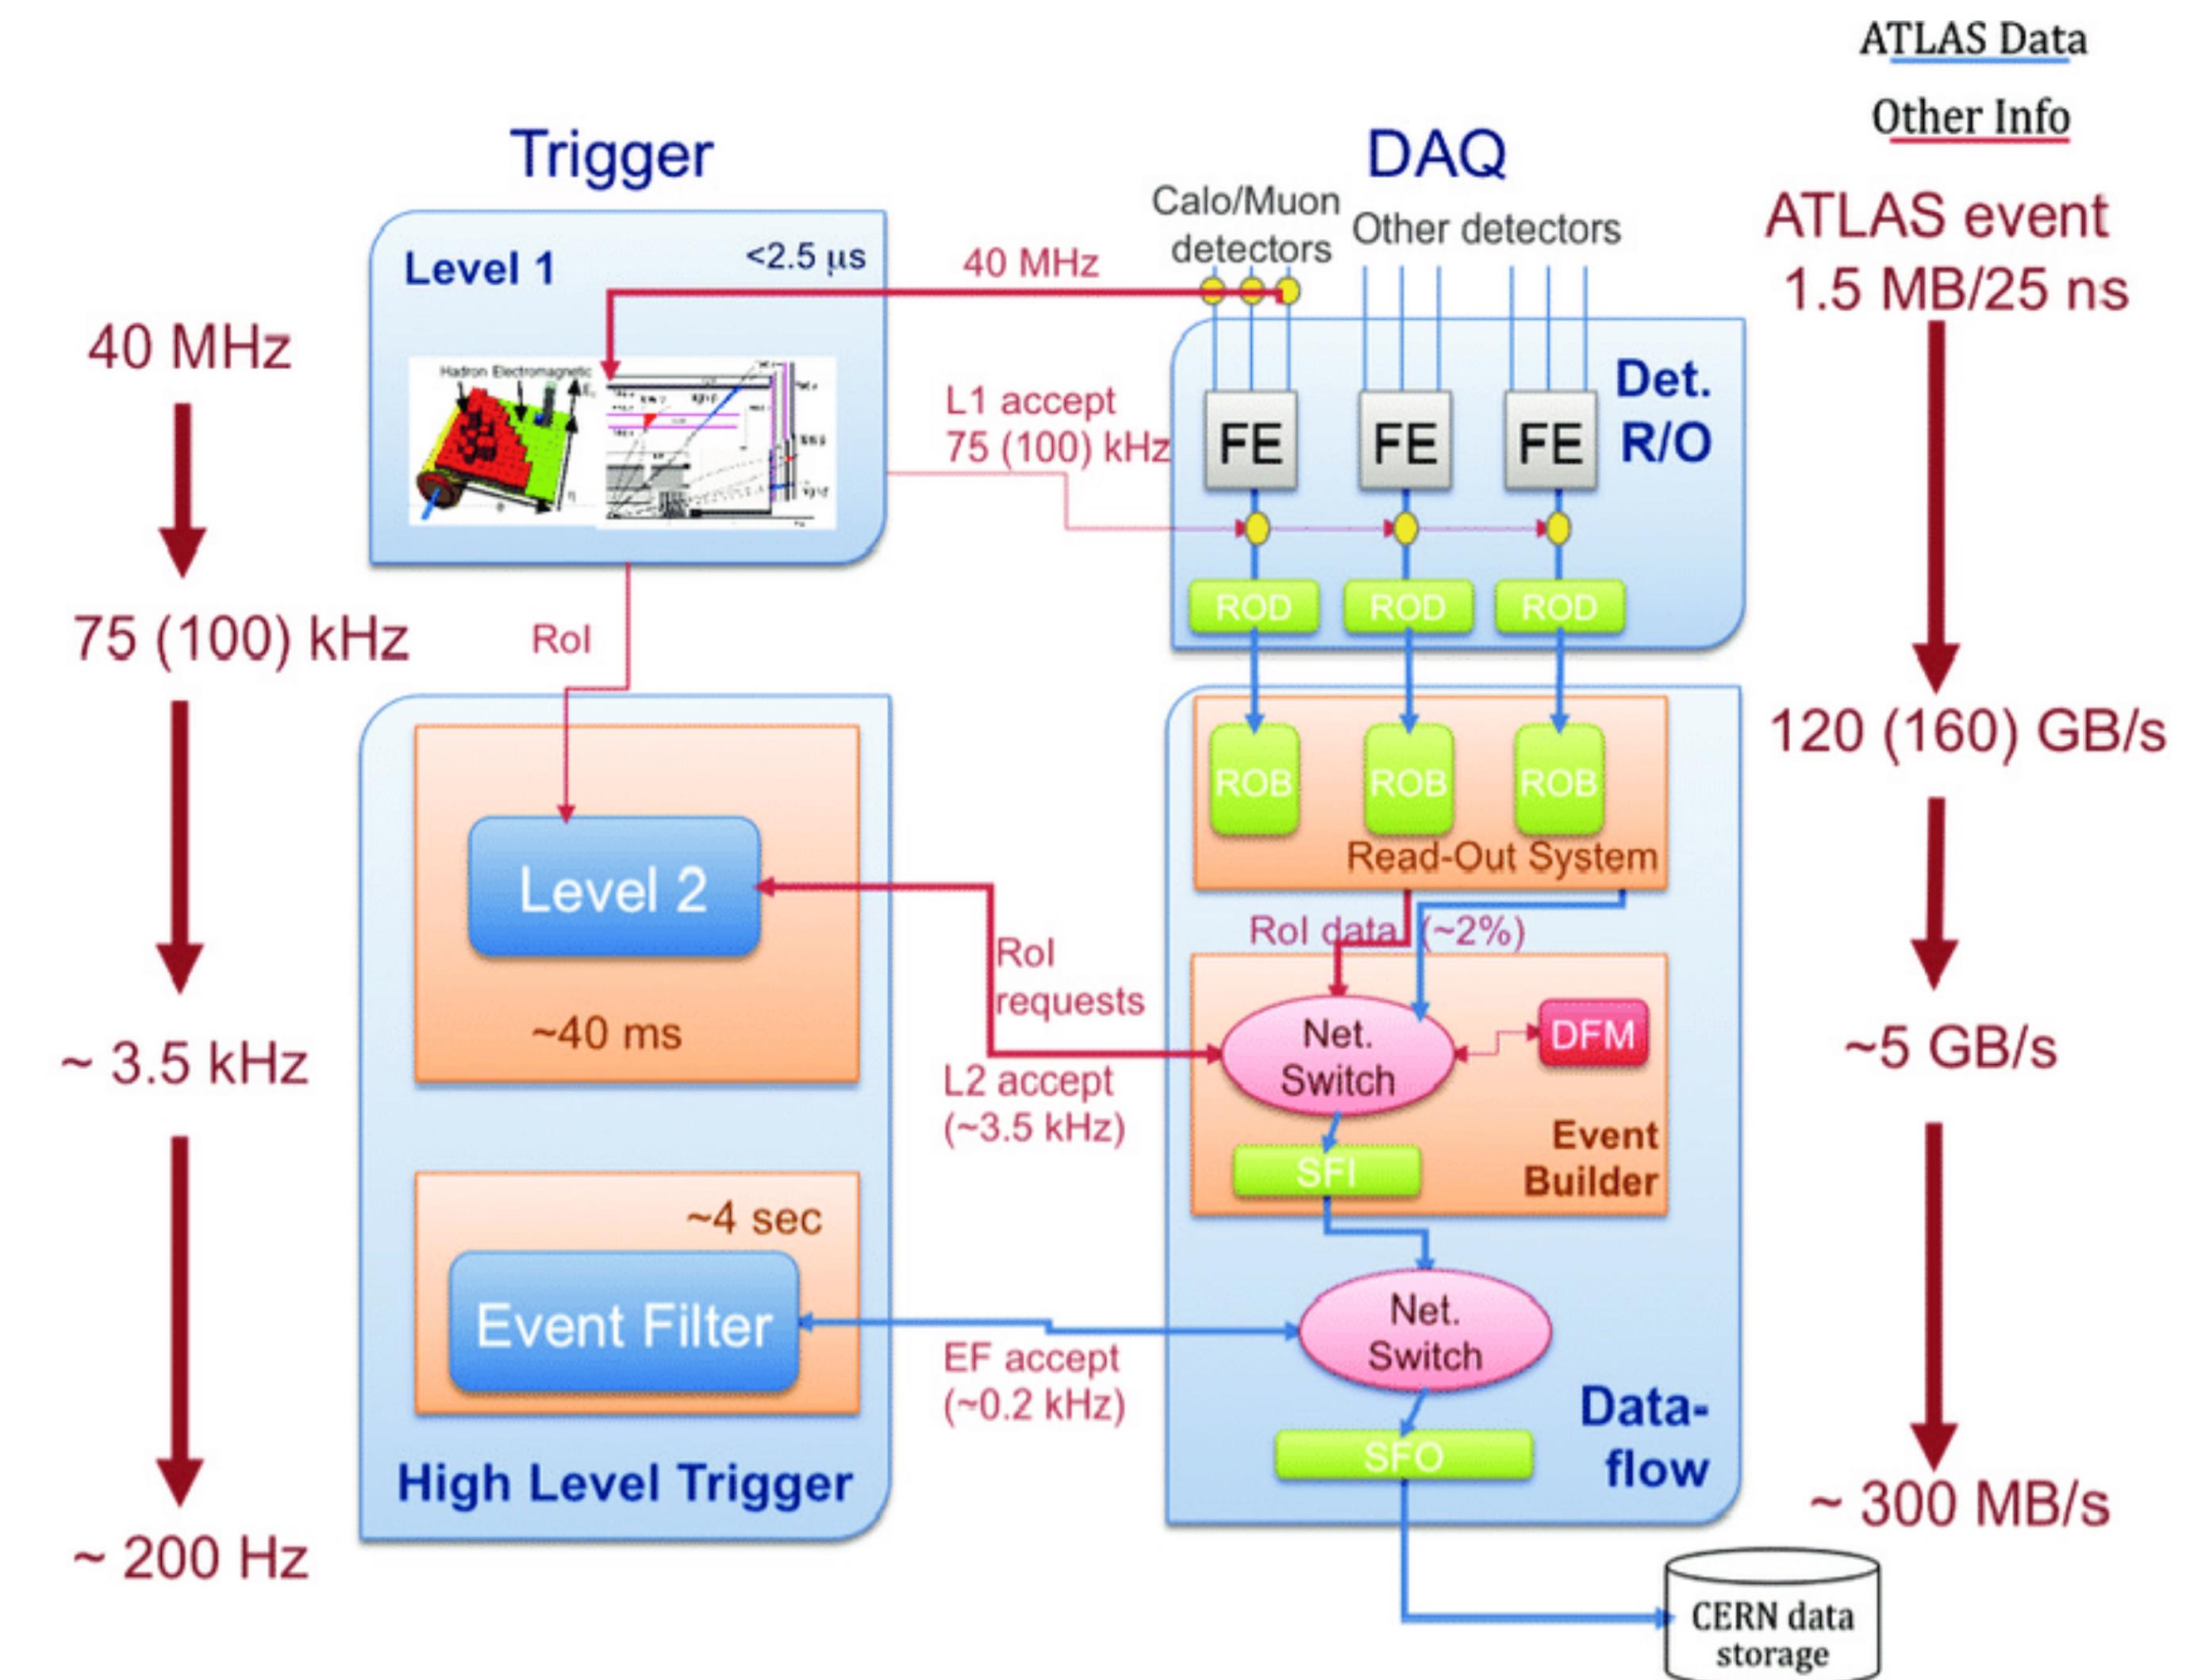
~40 Million (40 MHz) pp-collision per second

Each collision is an “event.”

- One bunch has $O(10^{11})$ protons, with ~2800 bunches per proton beam, each bunch is separated 25 ns apart from each other:
- 20 collisions happen at one bunch crossing

Data: ~ 25 MB / raw event, or about ~ 1 PB of raw data per second

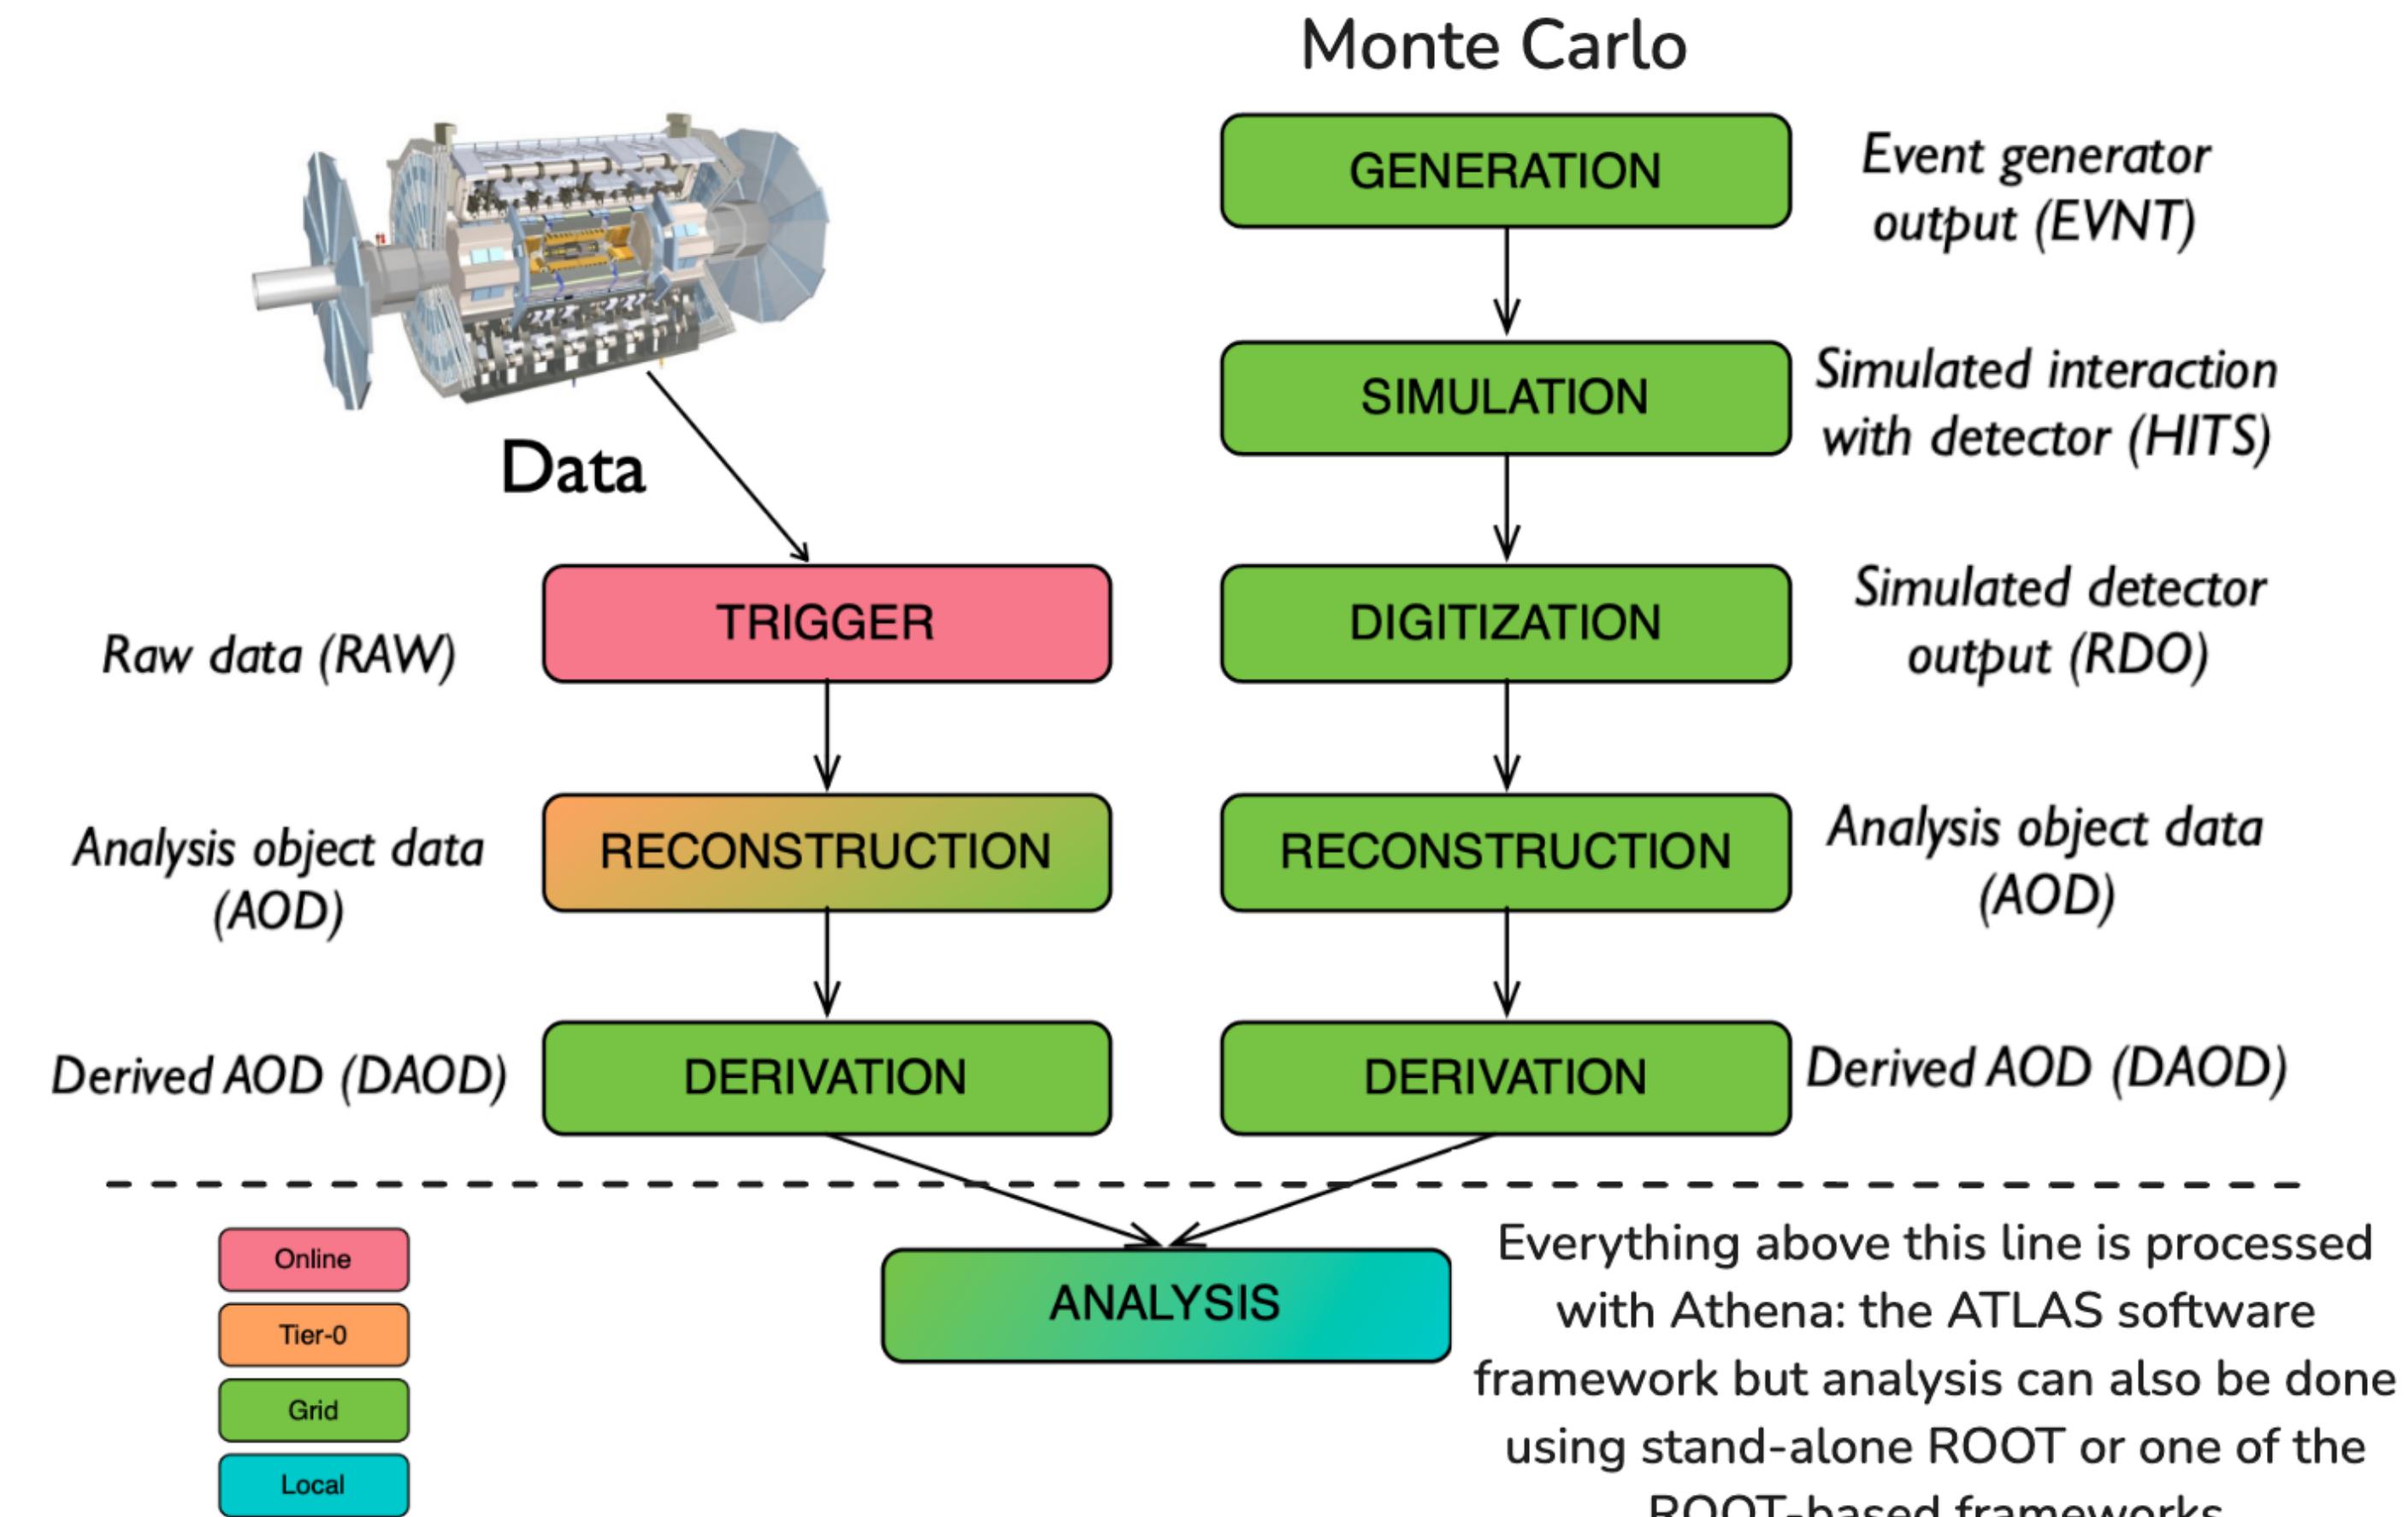
“Pileup” is the result of multiple collisions occurring from one bunch crossing.



ATLAS Analysis Pipeline

Analysis requires both real data and Monte Carlo simulated data (MC)

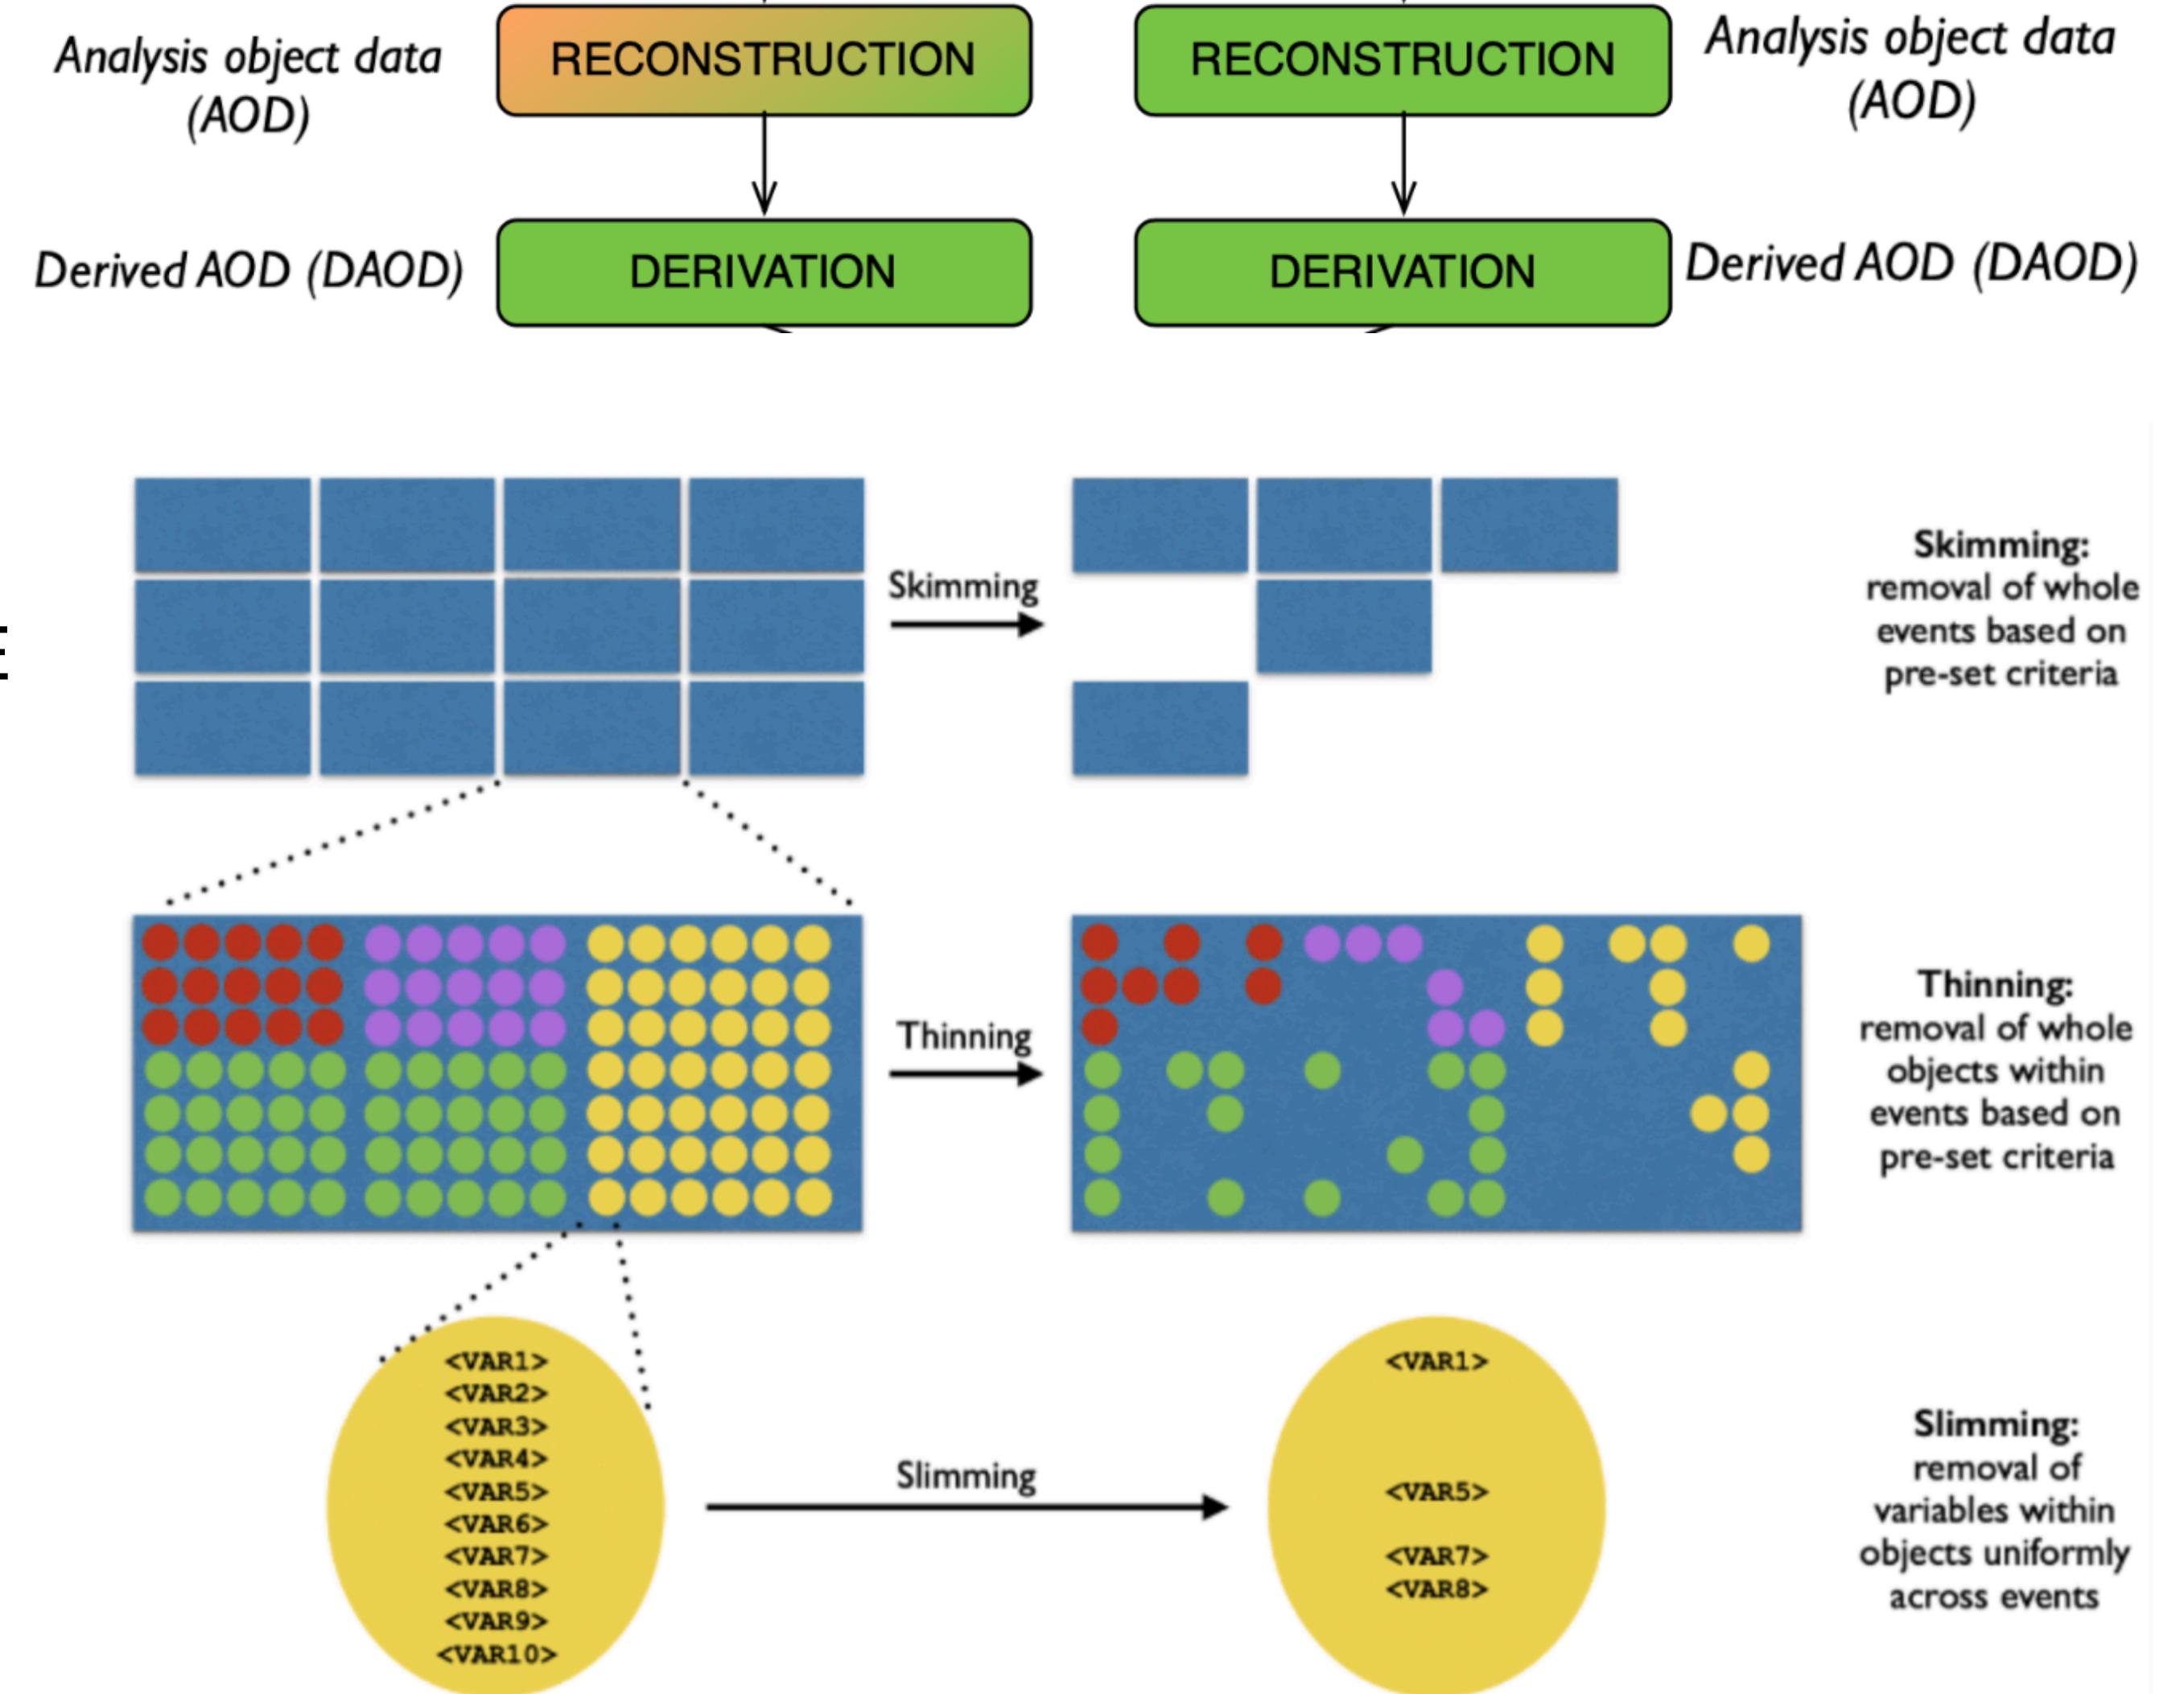
- Data:
 - Comes from the detector, trigger decides how interesting an event it comes from
- MC:
 - Files include more information per AOD, i.e. the full Trigger Menu, as the event generation might need to be selected for various detector configurations



ATLAS Analysis Pipeline: Derivation



- Derived Analysis Object Data (or DAOD), or a derivation, is a ROOT file used for analysis
 - From: AOD (**O(1 MB) / event**)
 - To: DAOD (**O(10 kB) / event**)
 - Main formats for ATLAS: PHYS/PHYSLITE
- Grid is used to create these files in a job, these jobs can be resource intensive (primarily memory/disk)
- Grid jobs have a 2 GB / core budget.
- Persistent storage should be minimized whenever possible.



Athena and LCG

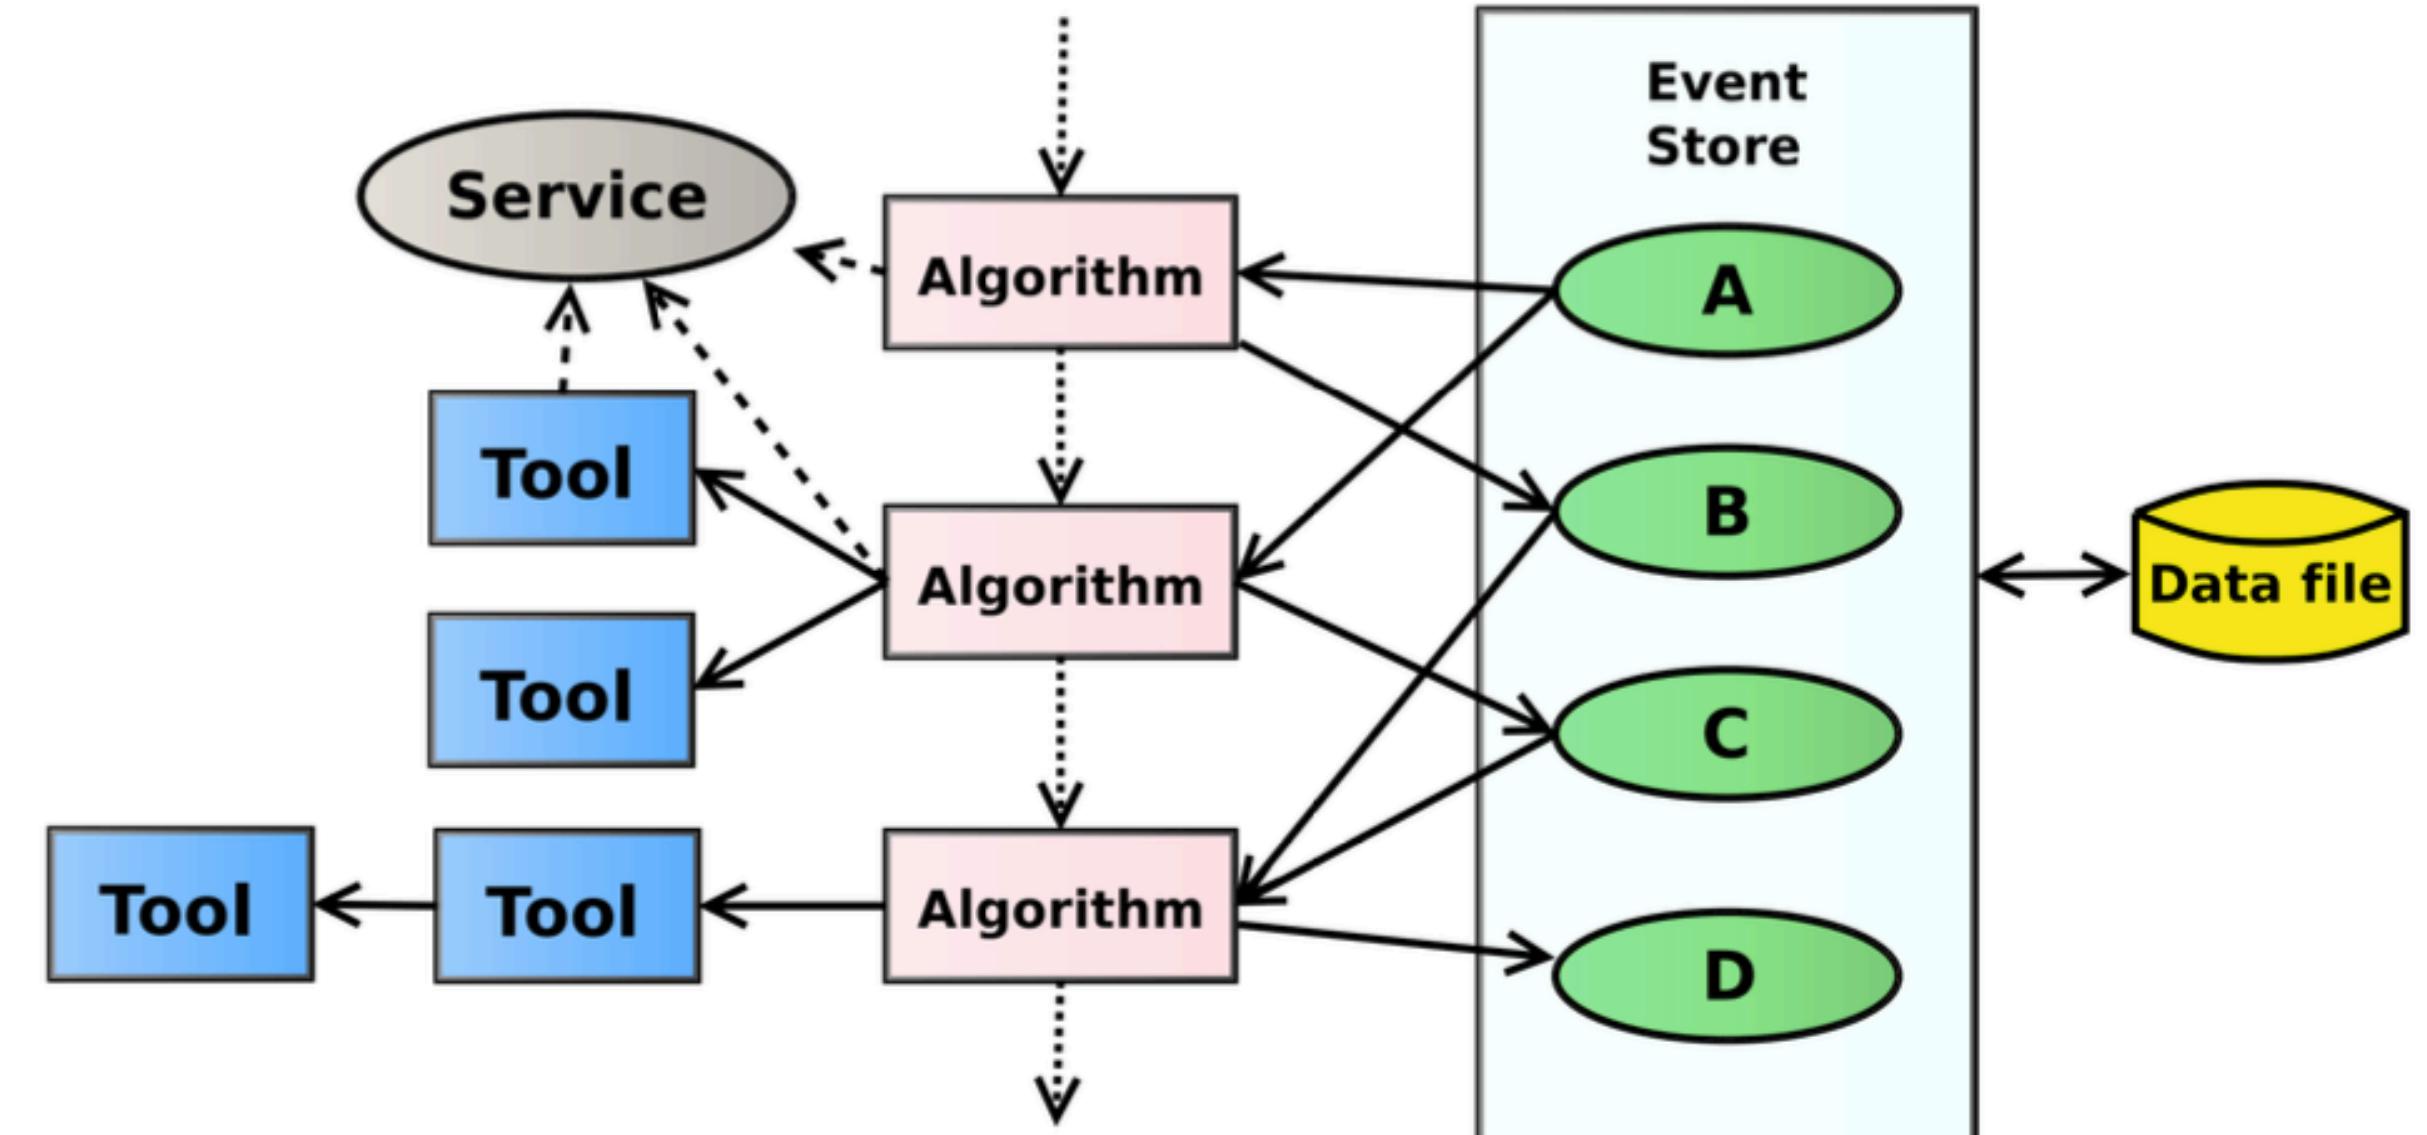
Athena is the open-source software framework for the ATLAS experiment.

- Based off of the Gaudi project, it uses ROOT and other software included in the LHC Computing Grid (LCG) software stack.
 - LCG stack has a set of software frameworks that provide general solutions to LHC experiments computing needs.
- ~ 500 packages:
- binary builders, compilers, language libraries, dependencies, simulation and analysis software, and more.

The screenshot shows the GitHub repository for the ATLAS Experiment's main offline software repository, "athena". The repository has 146,044 commits, 32 branches, 3,401 tags, and 536 releases. The "Code" tab is selected, displaying a table of files with their last commit details. The table includes columns for Name, Last commit, and Last update. Notable commits include "Merge branch 'gtest' into 'main'" by Frank Winklmeier and "CaloRec: Avoid some spurious F..." by CaloRec. The repository was created on December 06, 2018.

Name	Last commit	Last update
.devcontainer	Projects: enable flake8-bugbear ...	9 months ago
.vscode	Projects: enable flake8-bugbear ...	9 months ago
AsgExternal/Asg_Test	Added new ASG Run 3 test files ...	4 months ago
AtlasGeometryCommon	Added missing validity checks fo...	5 months ago
AtlasTest	Merge branch 'gtest' into 'main'	1 hour ago
Build	prepare_release_notes: minor fix...	3 months ago
Calorimeter	CaloRec: Avoid some spurious F...	18 hours ago

<https://gitlab.cern.ch/atlas/athena>



Athena Applications

An Athena application relies on components such as Algorithms, Tools, Services, and Properties.

Python is used for the applications job steering, a high-level configuration for the direction a user may want to take the job.

The **Event Store** is a Service

Tools can be anything Algorithms need, i.e. random number generators

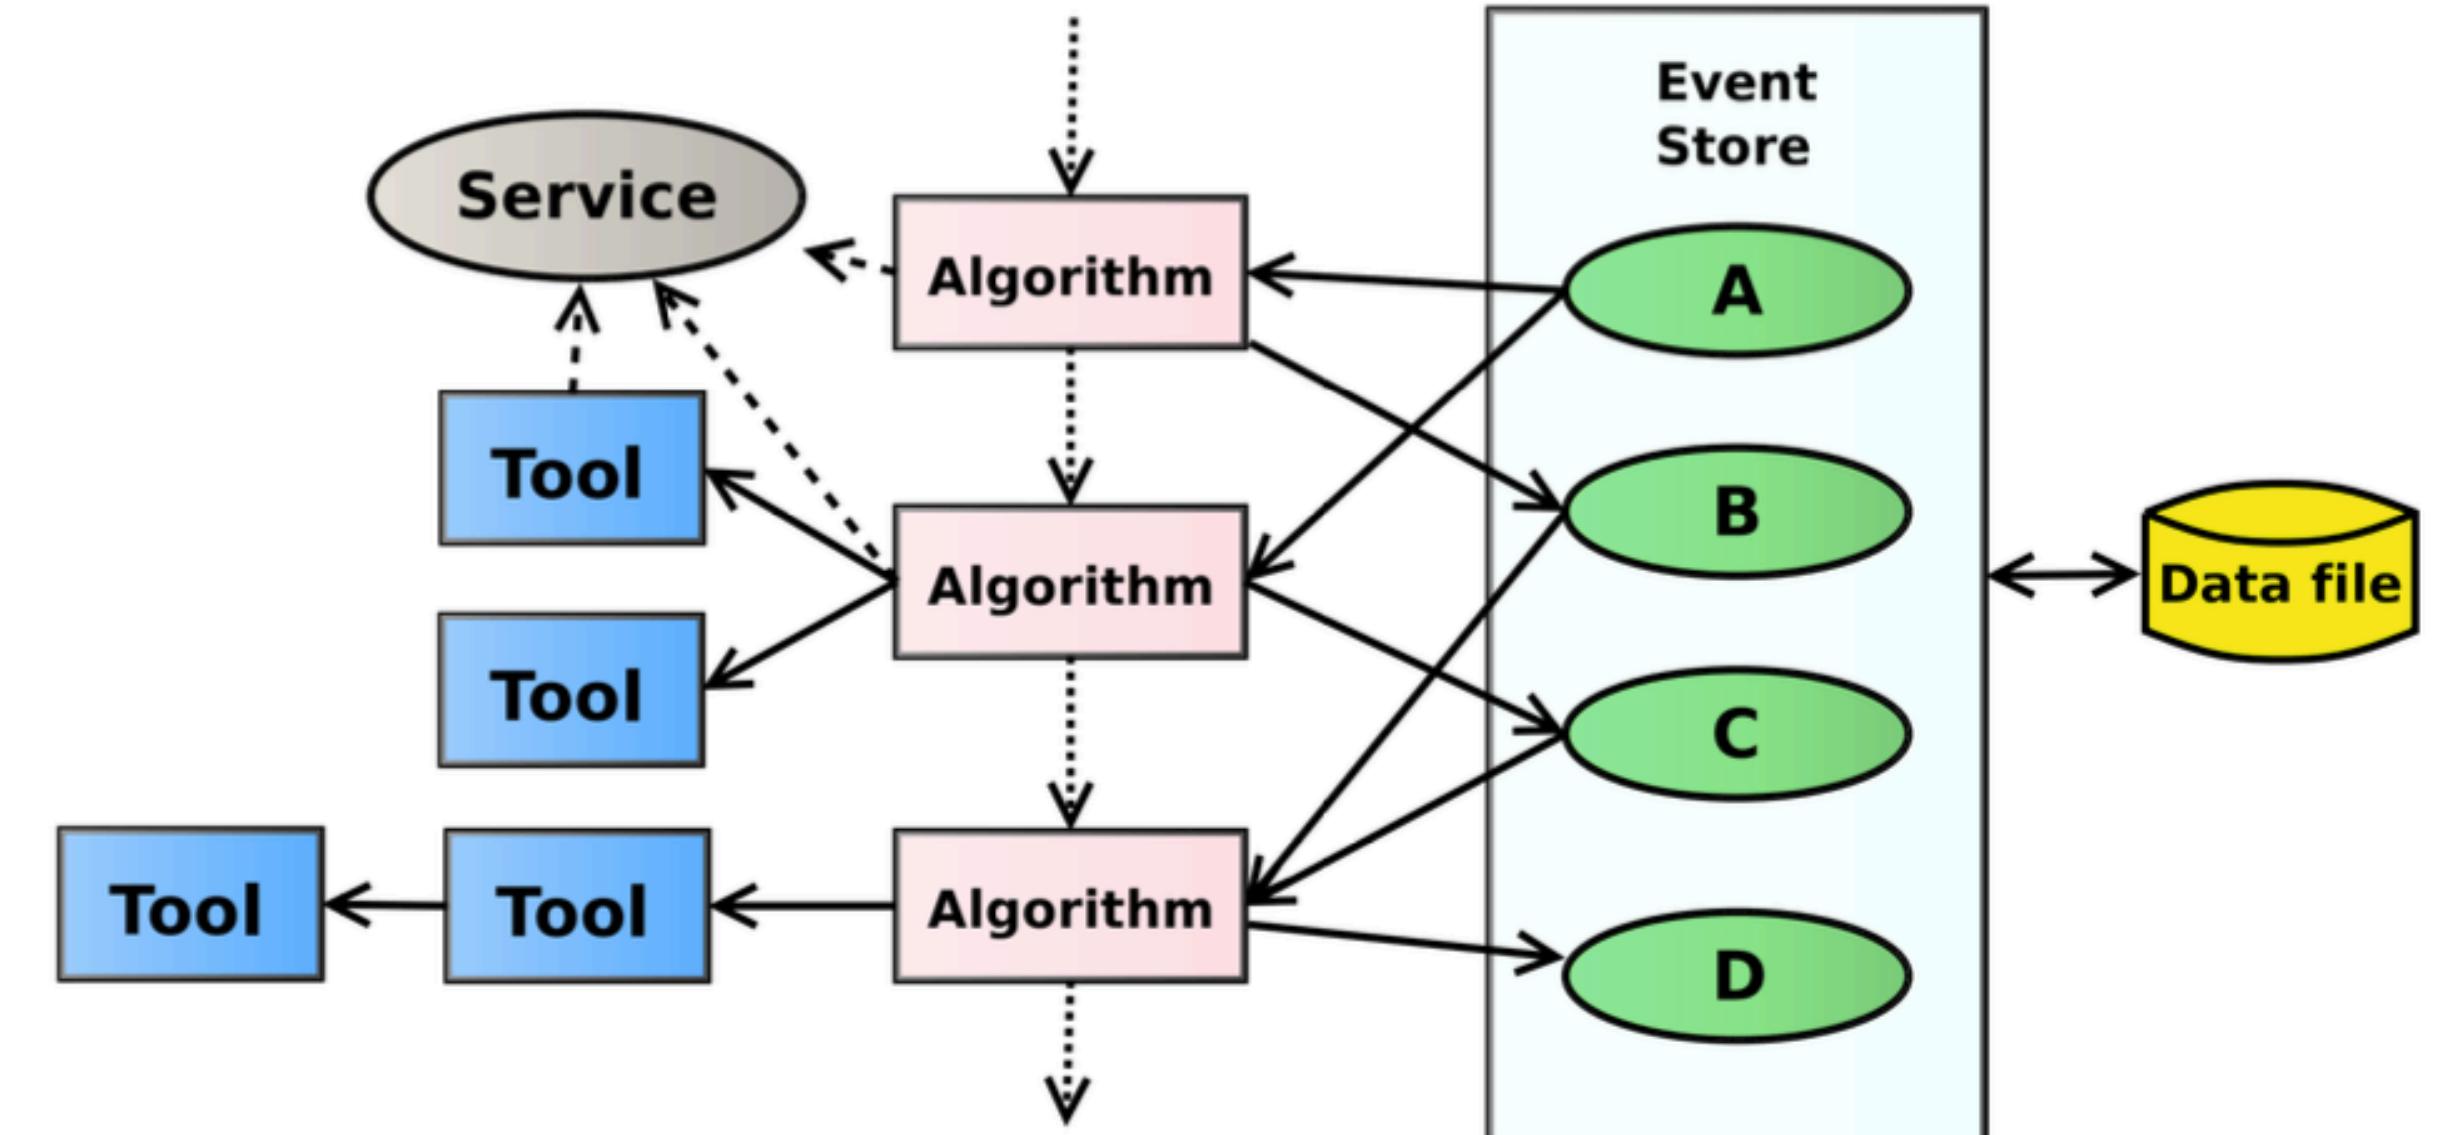
All of this sits on top of ROOT, using its own classes, Tools and Services, though it primarily acts on ROOT files.

The screenshot shows the GitHub repository for 'athena' under the 'ATLAS EXPERIMENT'. It displays a list of recent commits:

Name	Last commit	Last update
.devcontainer	Projects: enable flake8-bugbear ...	9 months ago
.vscode	Projects: enable flake8-bugbear ...	9 months ago
AsgExternal/Asg_Test	Added new ASG Run 3 test files ...	4 months ago
AtlasGeometryCommon	Added missing validity checks fo...	5 months ago
AtlasTest	Merge branch 'gtest' into 'main'	1 hour ago
Build	prepare_release_notes: minor fix...	3 months ago
Calorimeter	CaloRec: Avoid some spurious F...	18 hours ago

On the right side, there is a sidebar titled 'Project information' with details like 146,044 commits, 32 branches, 3,401 tags, 536 releases, and links to README and LICENSE.

<https://gitlab.cern.ch/atlas/athena>



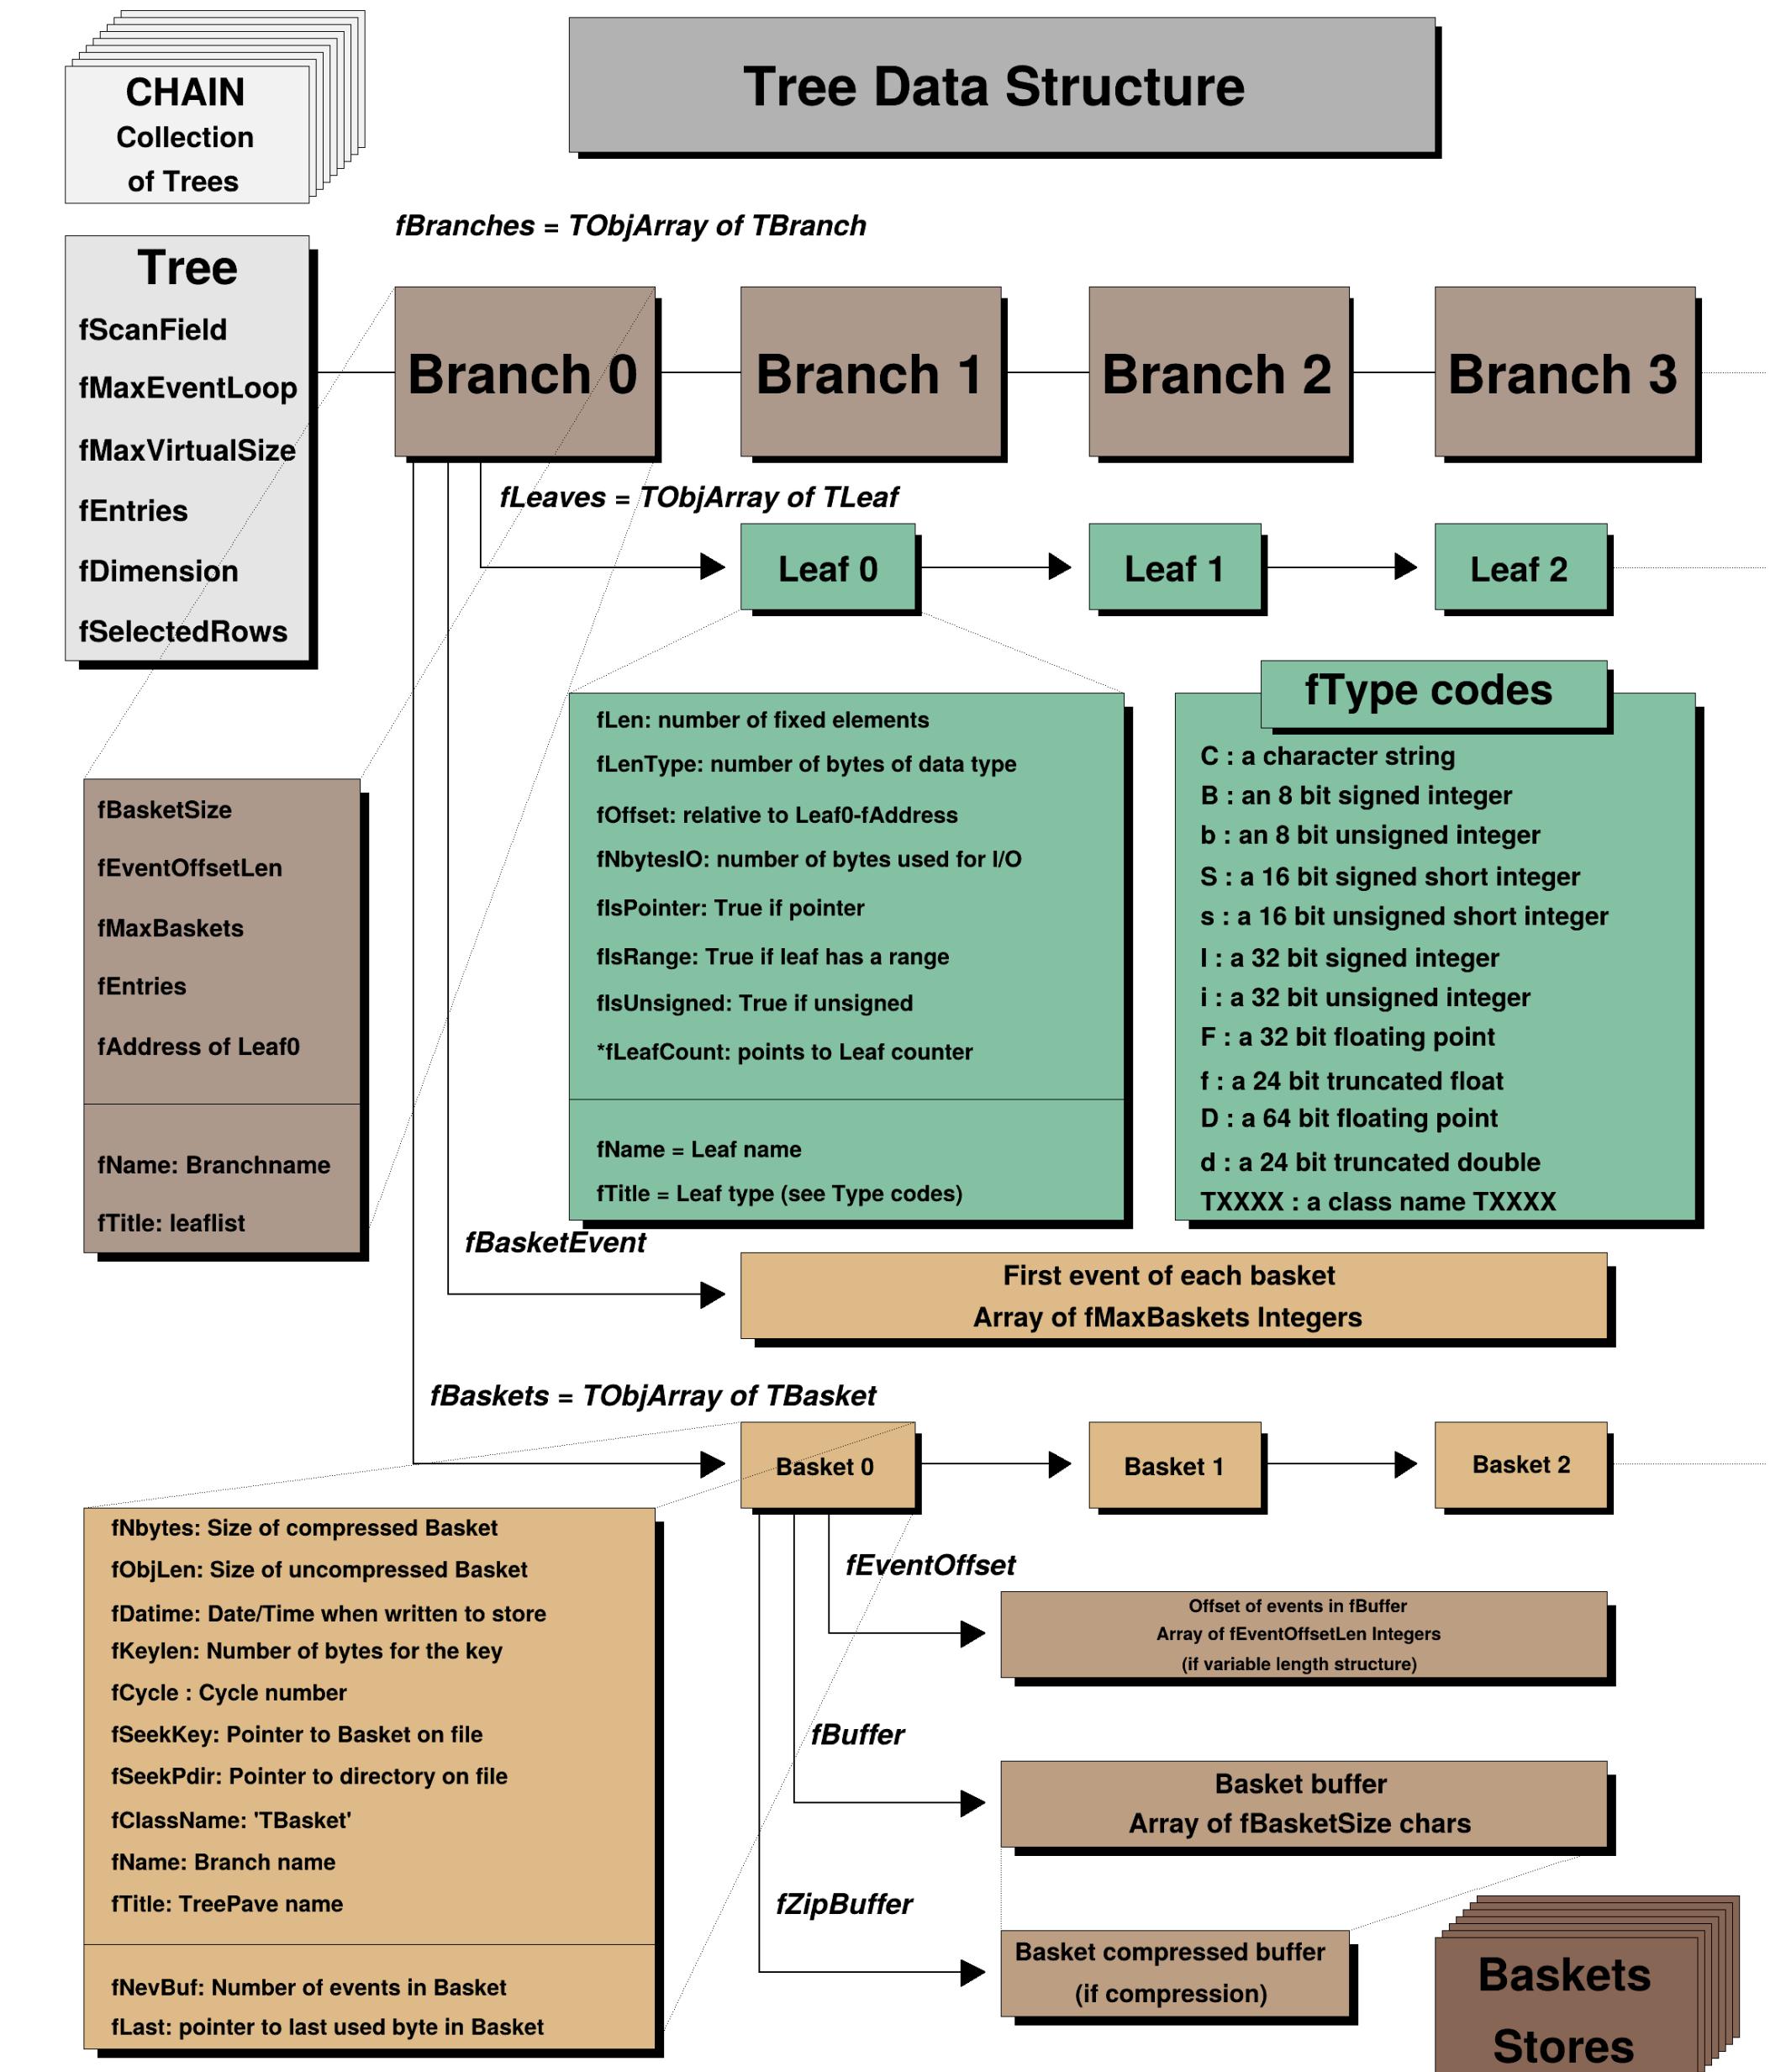
ROOT



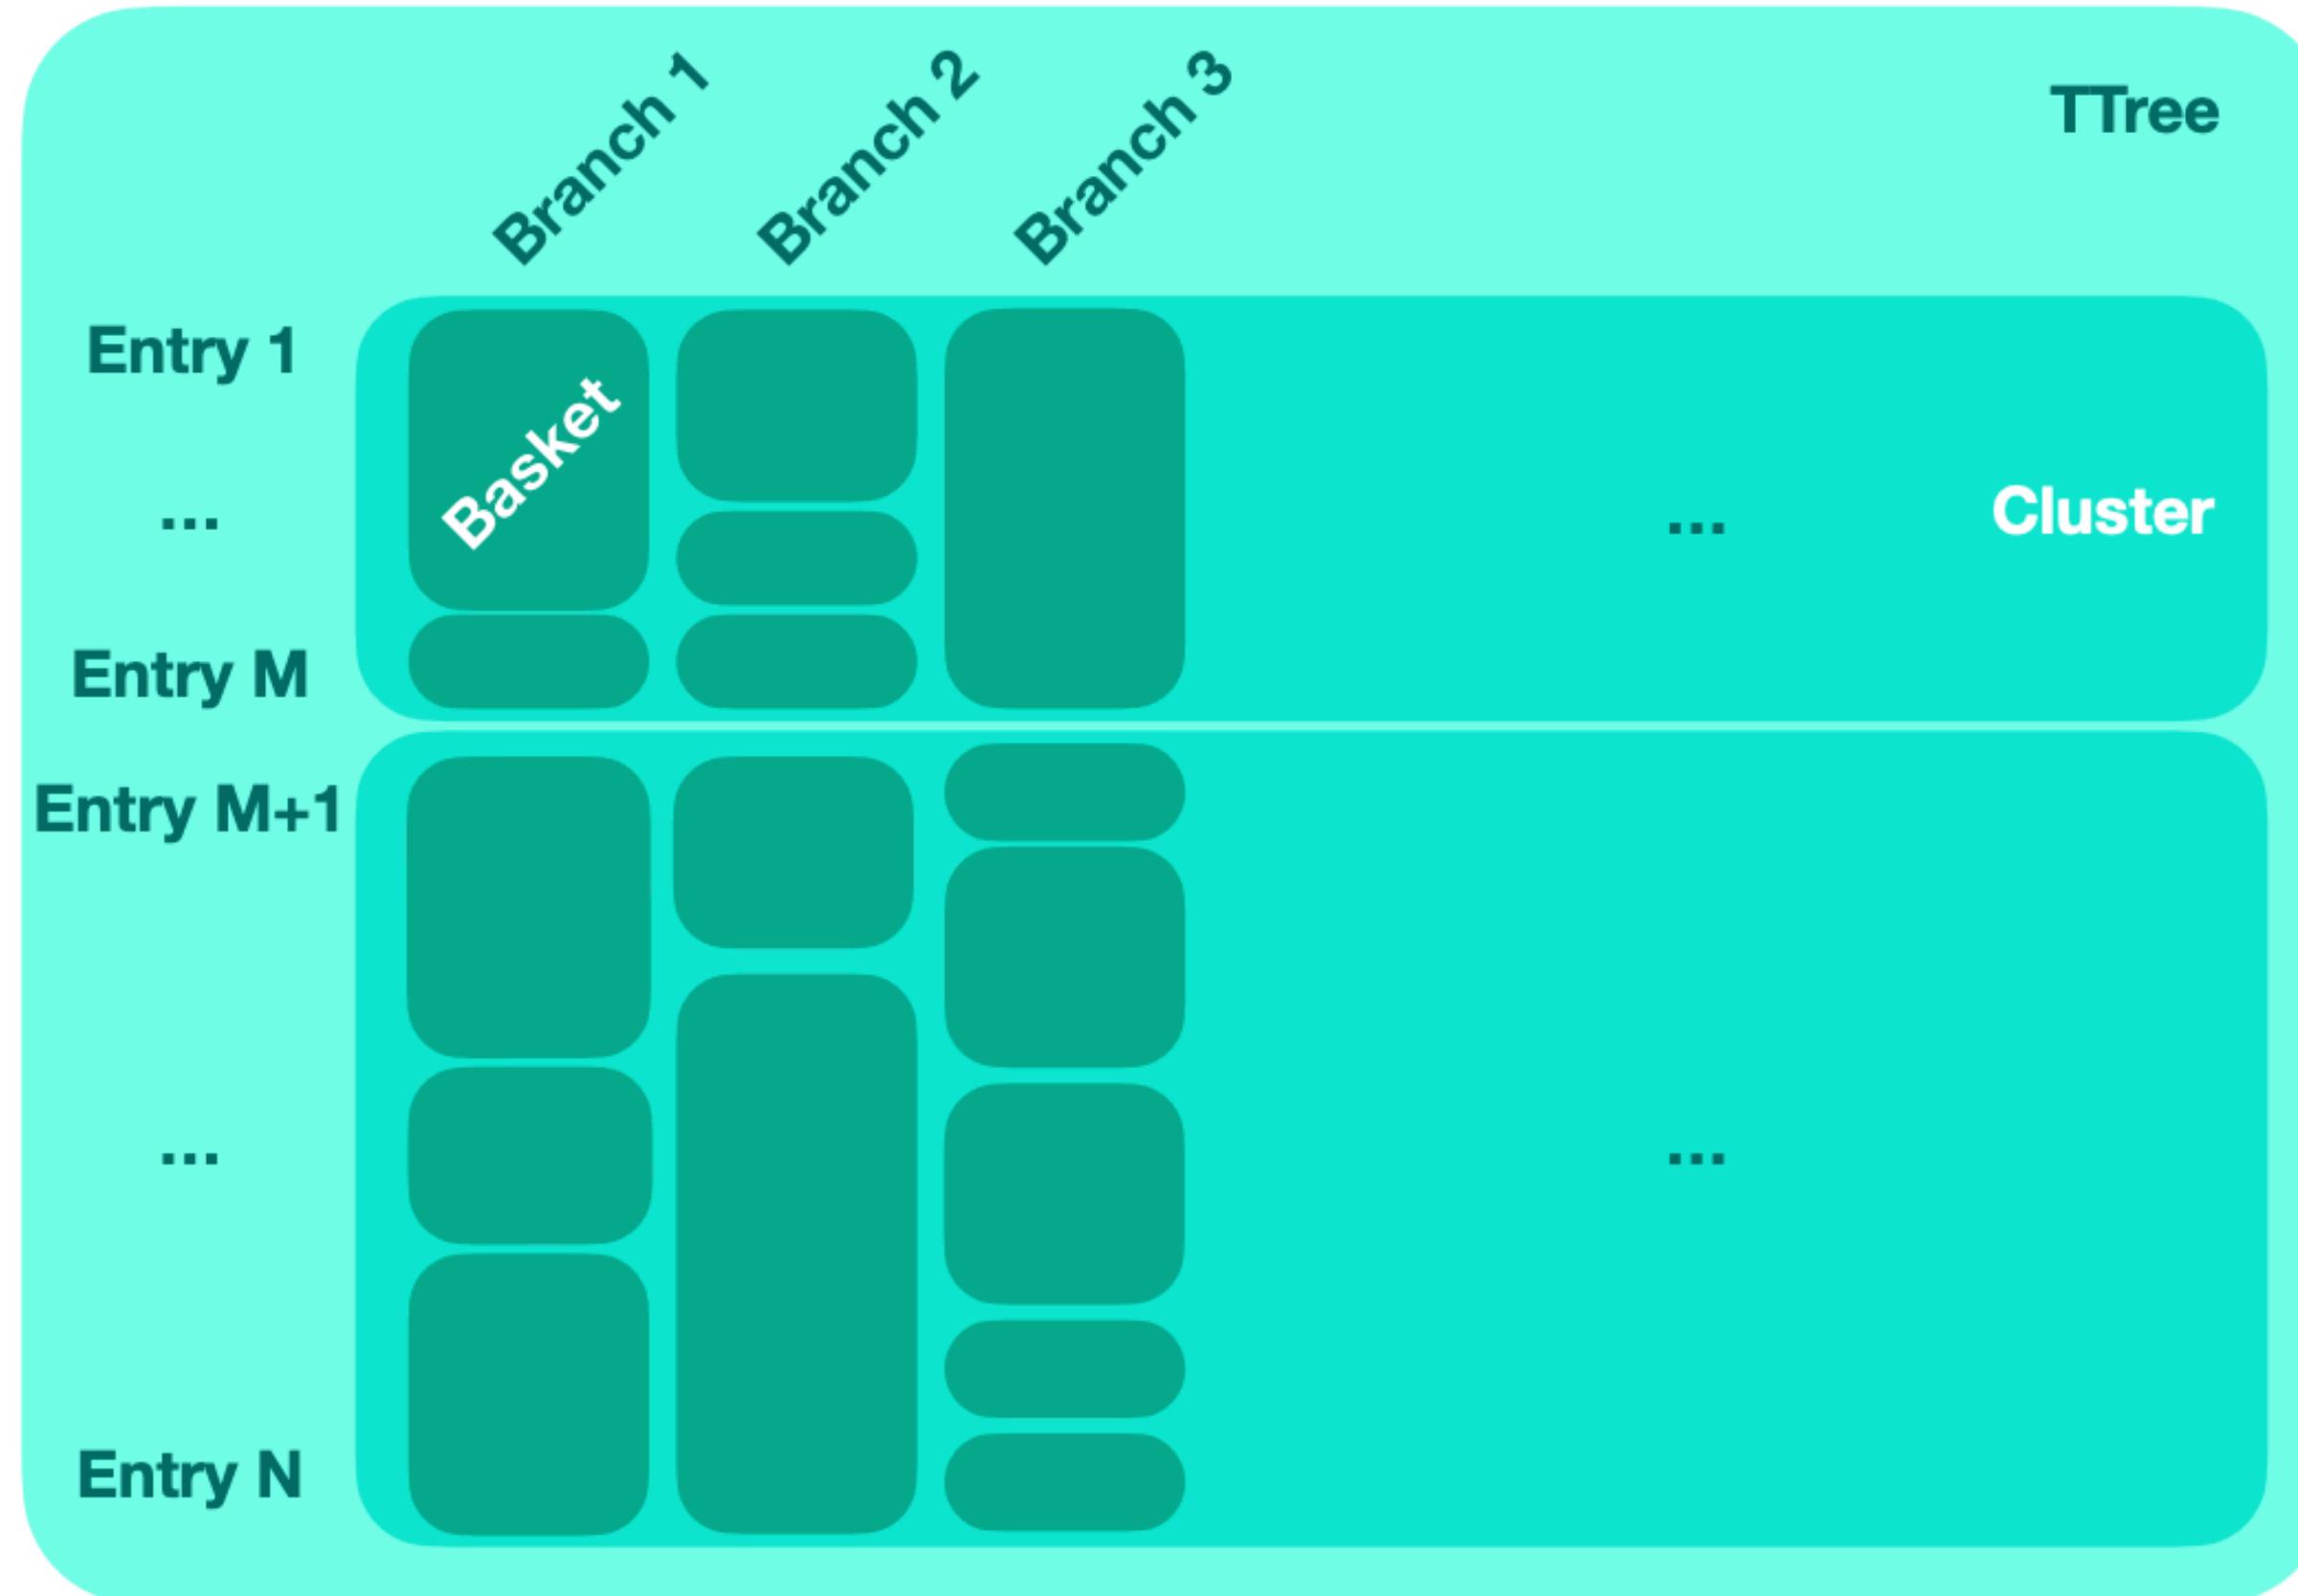
ROOT is one of the primary open-source software framework for high-energy physics analysis.

The primary structure ROOT uses is TTree, which includes a series of nested objects of varying size.

TTree is still used and will continue to be used for a while, but RNTuple is around the corner and has shown a lot of promise.



How are ROOT files organized?



DAODs are ROOT files, so they contain TTrees

TTrees have branches which can correspond to physics objects (i.e. Electrons, Muons, etc.)

Each branch stores entries into “baskets” for better I/O handling

Baskets are dedicated memory buffers that contain data which are compressed and then stored to disk

How can we reduce the load (memory/disk) to the Grid?

Currently, **ROOT during runtime**:

- Optimizes baskets after tree->Fill() is called, which assigns a basket buffer size optimized for the tree at that time. Maximum of 1.3 MB

Athena before runtime:

- Has several settings, but the two we're looking at are:
 - Maximum Basket Buffer Size: 128 kB
 - Minimum Number of Entries to the Buffer: 10 entries

For both ROOT and Athena: What happens to performance when we fix the maximum buffer limit prior to runtime memory allocation for the job occurs?

Now we start:

Toy Model Investigation into TBranch basket buffer sizes

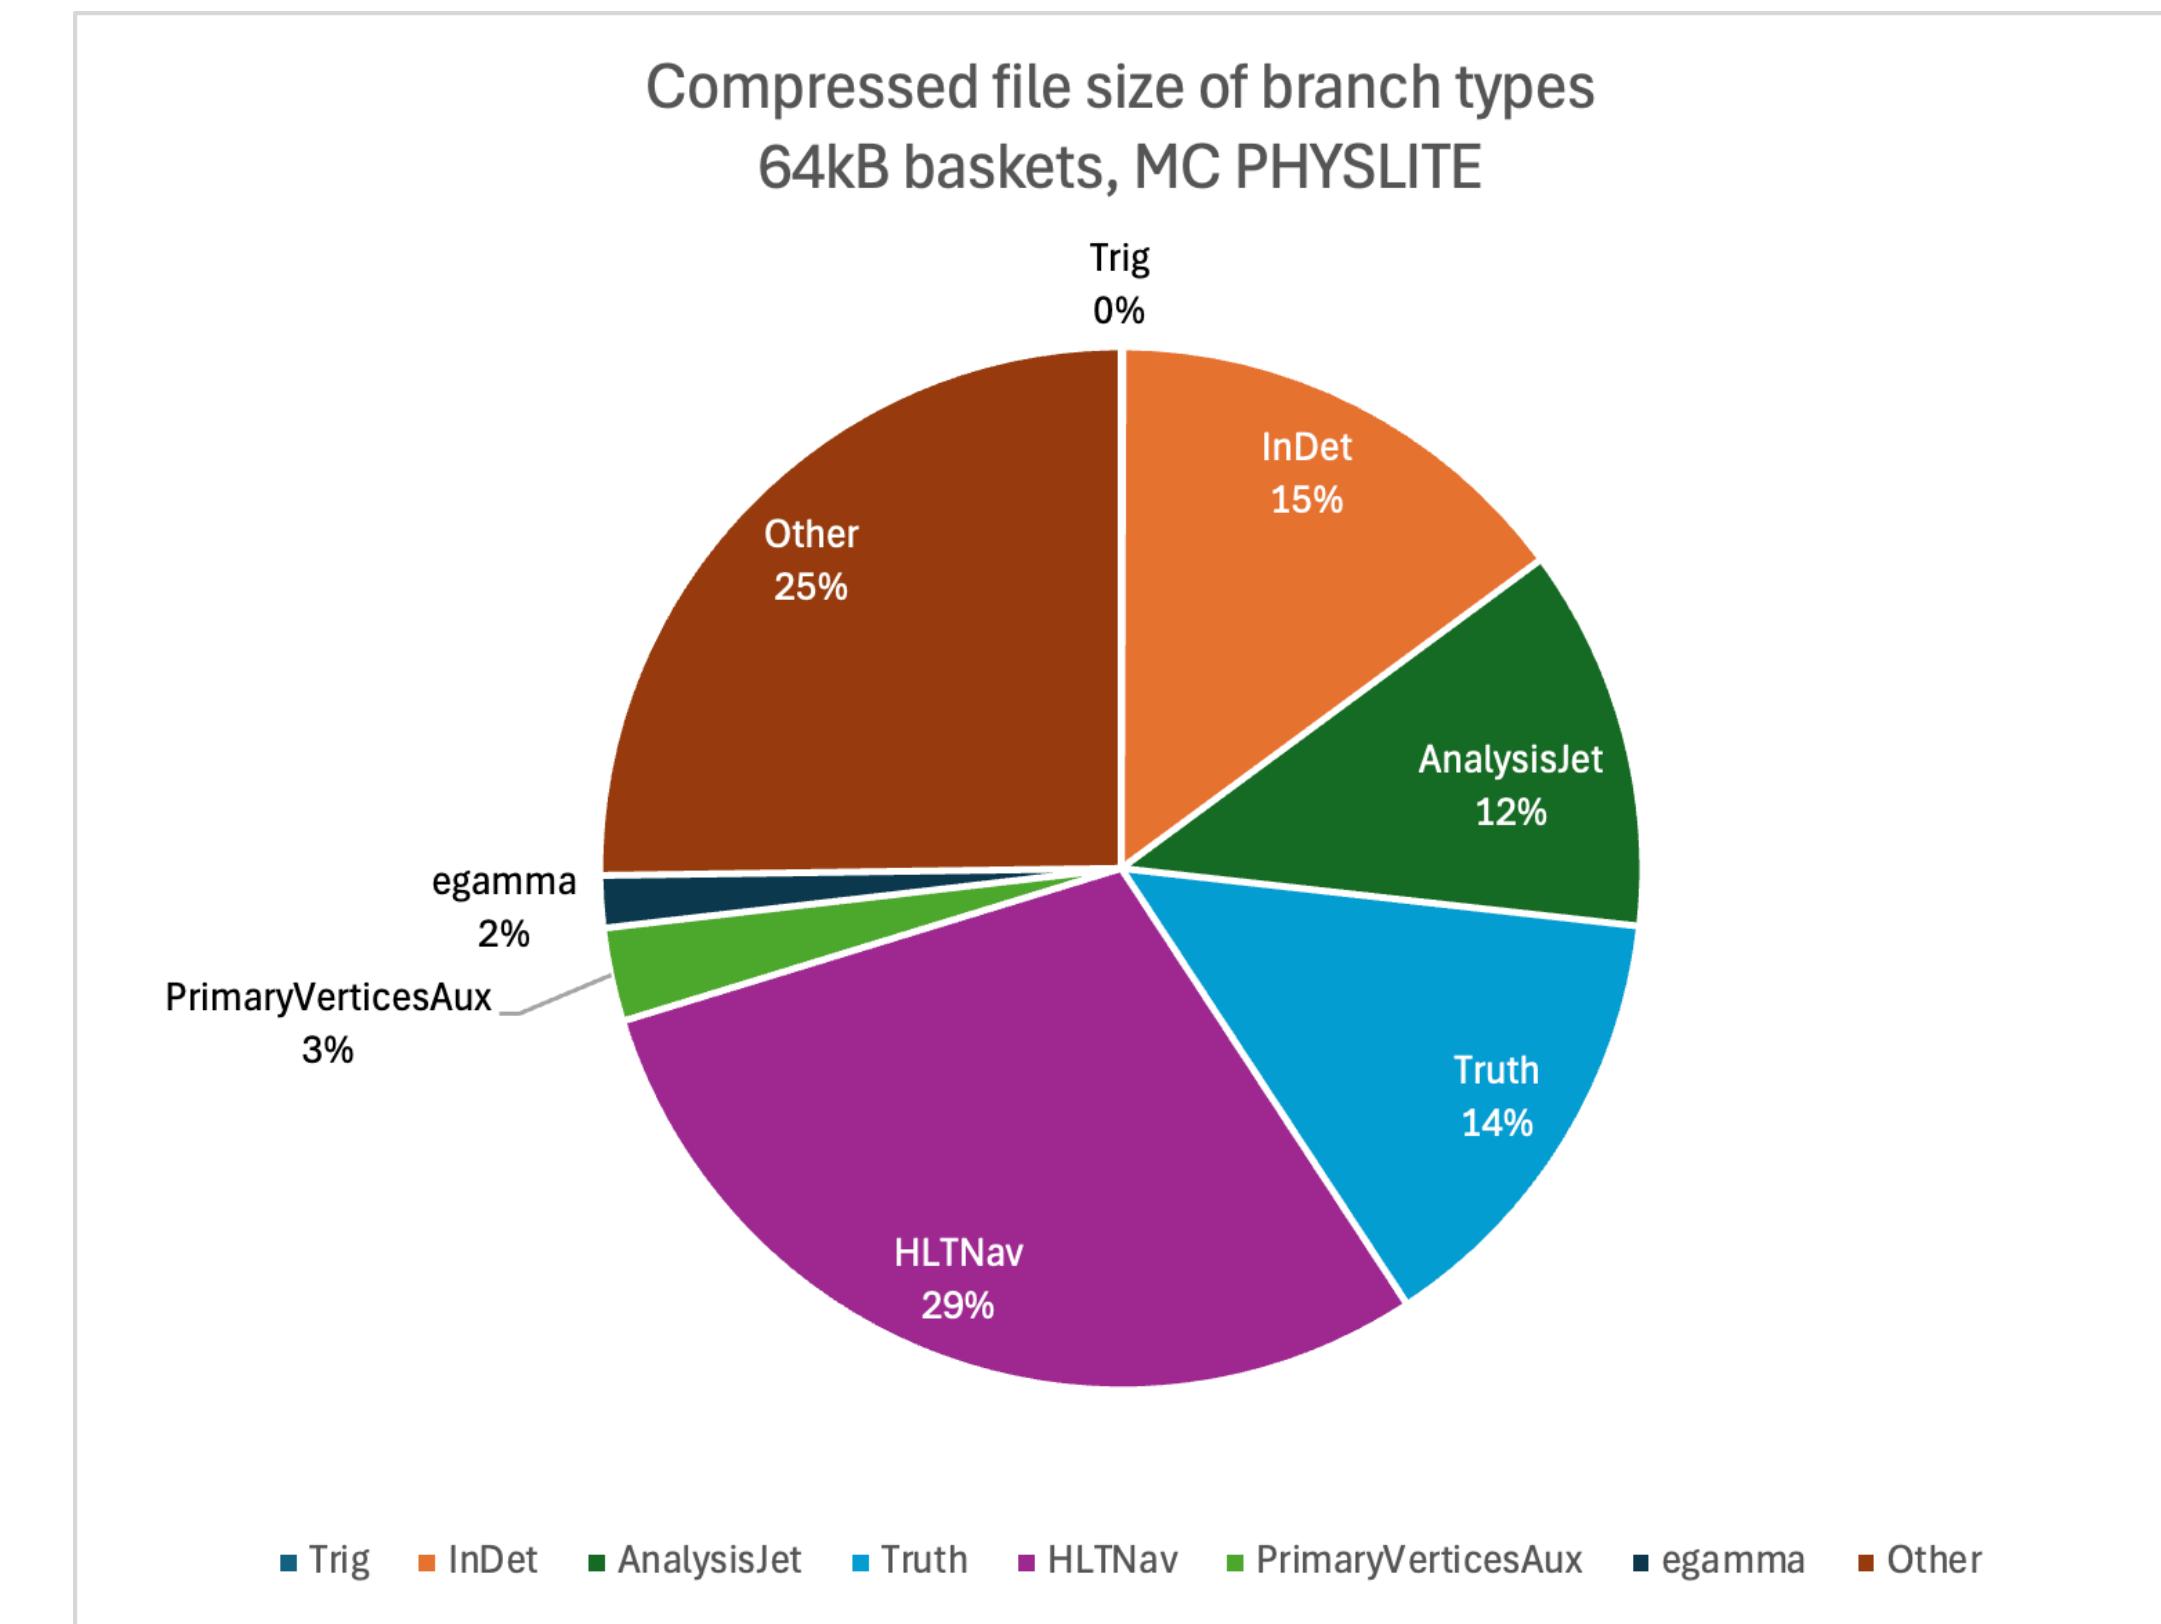
Toy Model I/O Tests: What to look out for

Can we mimic what will be put into branches?

What we know:

- MC AODs are typically larger than Data AODs (i.e. Full TriggerMenu for MC)
- Branch data consists of things like:
 - Charges (Electrons: $q = -1$, Muons: $q=-1$, etc)
 - pT (best represented by floating point numbers)
 - η values (range from 0 \rightarrow Infinity, but majority are in that $2 < |\eta| < 2.7$ range)
 - ϕ values

Given an example MC derivation we might expect to see an average compression ratio of about ~ 5.5



Tag	Compressed File Size (MB) (out of 2.5 GB)
Trig	1.13
InDet	336.94
AnalysisJet	269.24
Truth	317.48
HLTNavy	669.33
PrimaryVerticesAux	65.81
egamma	36.68
Other	572.08

Toy Model I/O Tests: What to look out for

branch-name	tag	num-of-entries	basket-size	total-file-size	compressed-file-size	compression-ratio
891	TruthForwardProtonsAuxDyn.classifierParticleOutCome	Truth	140000	10752	2034540	403401
892	TruthForwardProtonsAuxDyn.polarizationPhi	Truth	140000	10752	2029090	397102
893	TruthForwardProtonsAuxDyn.polarizationTheta	Truth	140000	10752	2030180	400159
894						
895					AVG:	5.77
...						

Just one example

Tag	Compressed File Size (MB)
Trig	1.13
InDet	336.94
AnalysisJet	269.24
Truth	317.48
HLTNav	669.33
PrimaryVerticesAux	65.81
egamma	36.68
Other	572.08

Toy Model I/O Tests

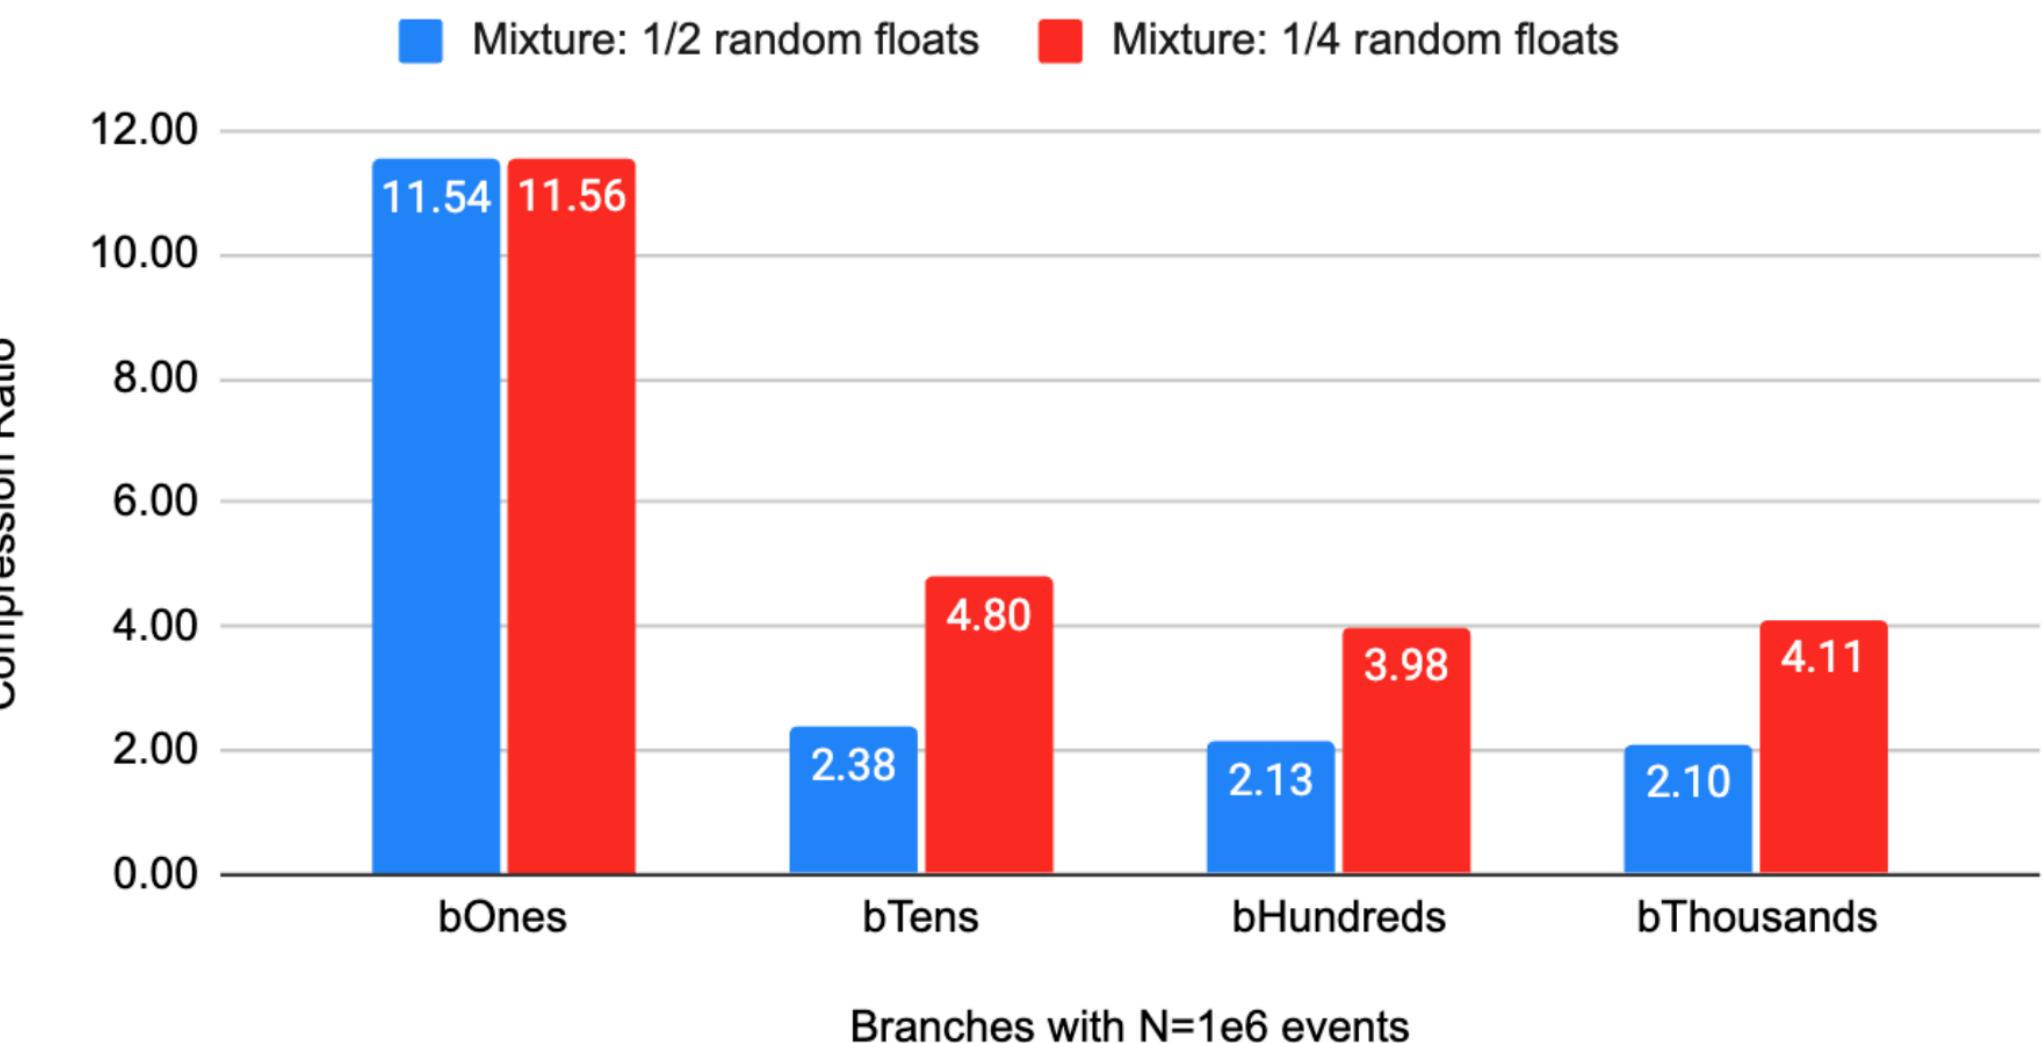
Early toy model derivations test the role of basket sizes with compression factors

For any toy model where branches are:

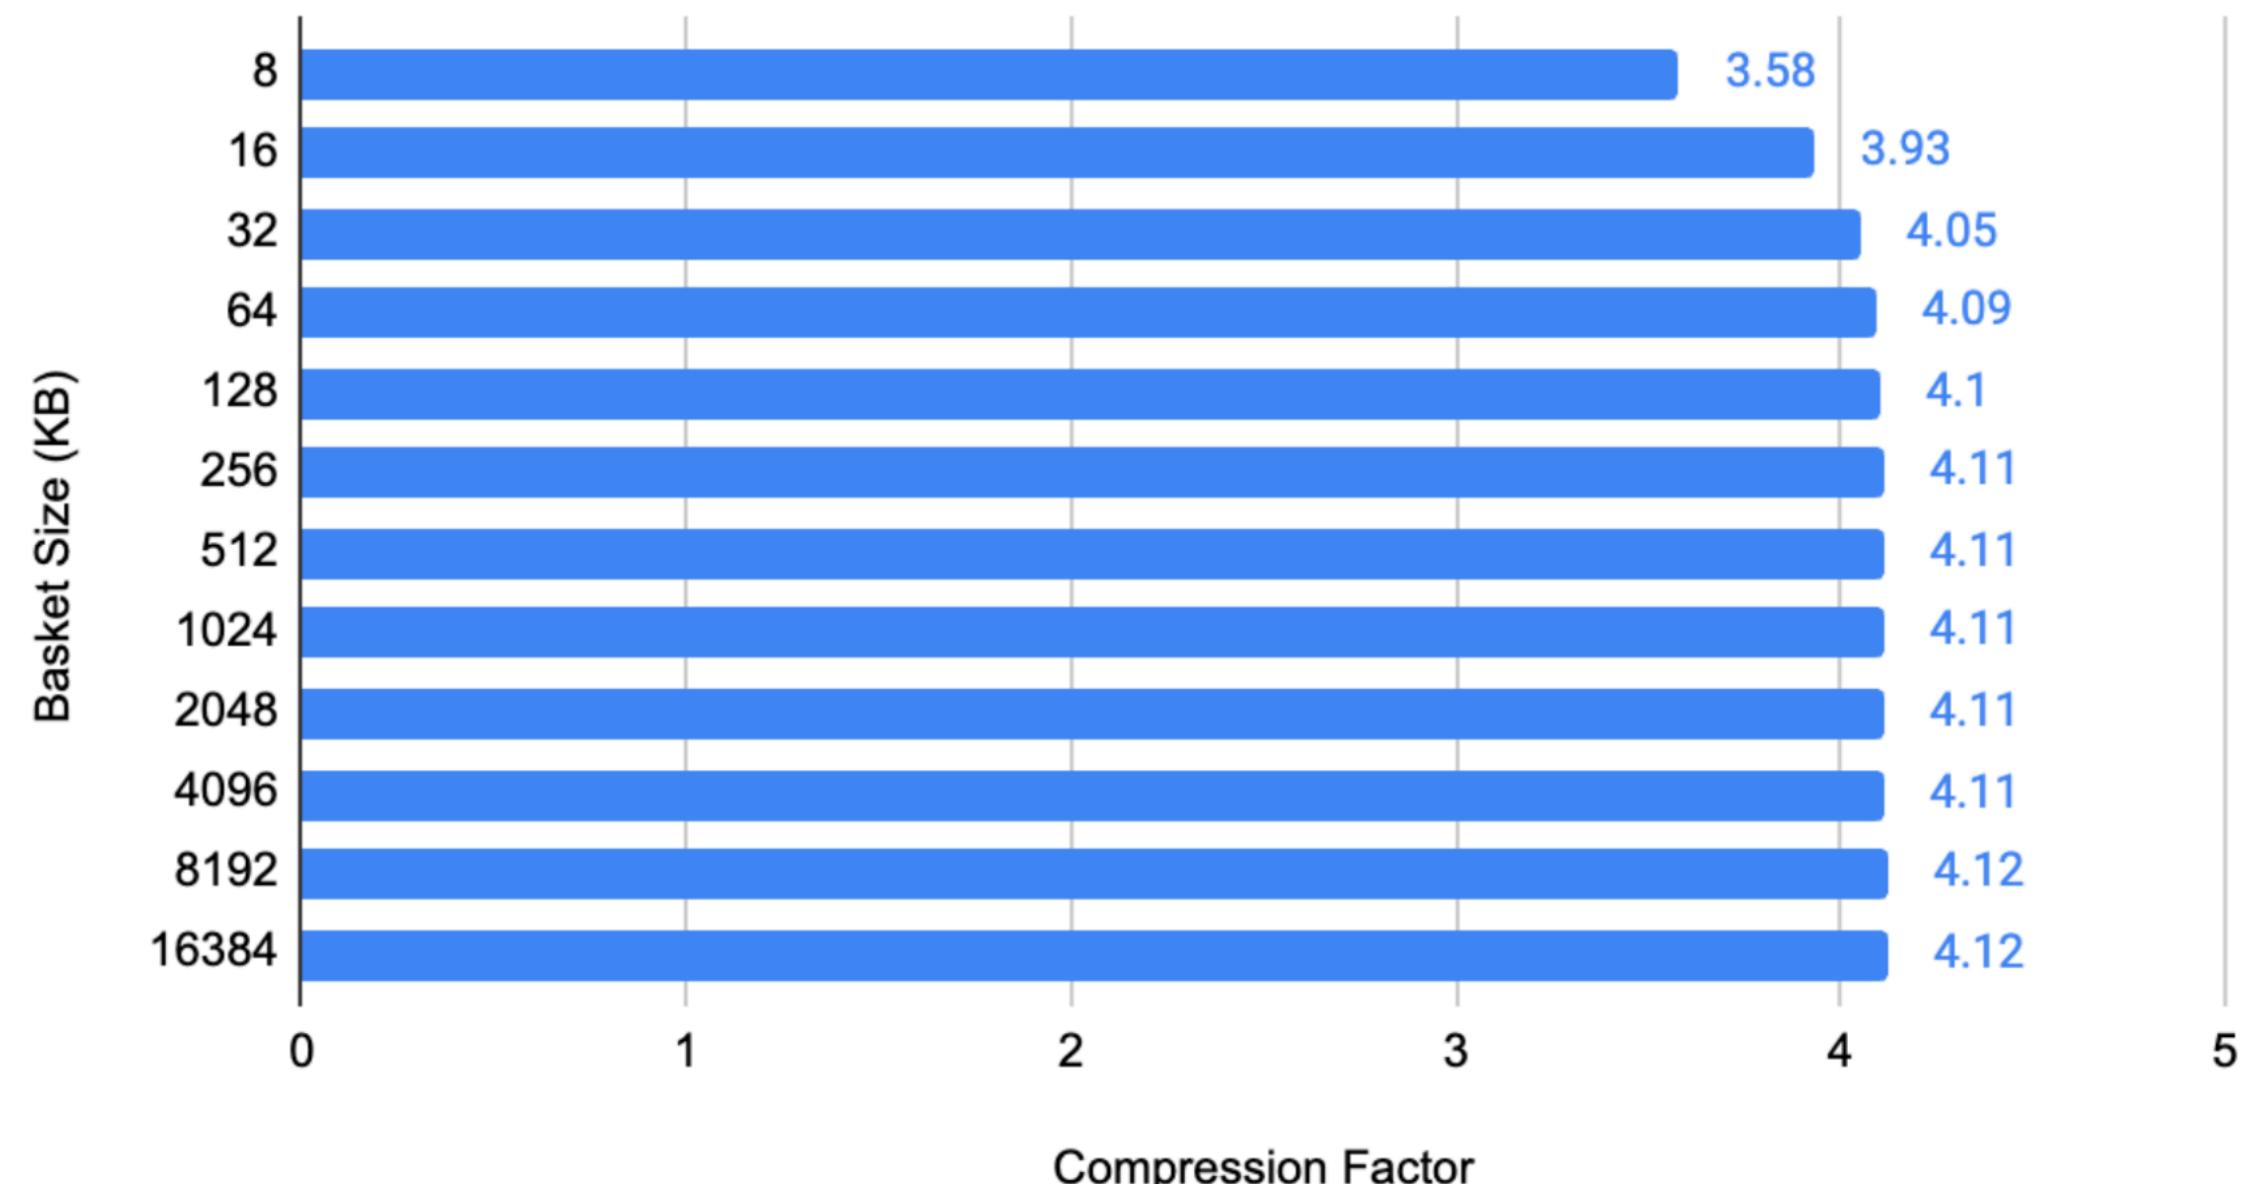
- Filled with some fraction of (random floats & the same constant integer)
- Mimic what data looks like (mixed bag of easy/hard compressible data)

We do see that increasing basket size does not scale the compression linearly

Compression Ratios for (1/2 random) and (1/4 random) branches at (N=1,000,000 events)



1/4 Random - Compression Factor vs. Basket Size (KB)



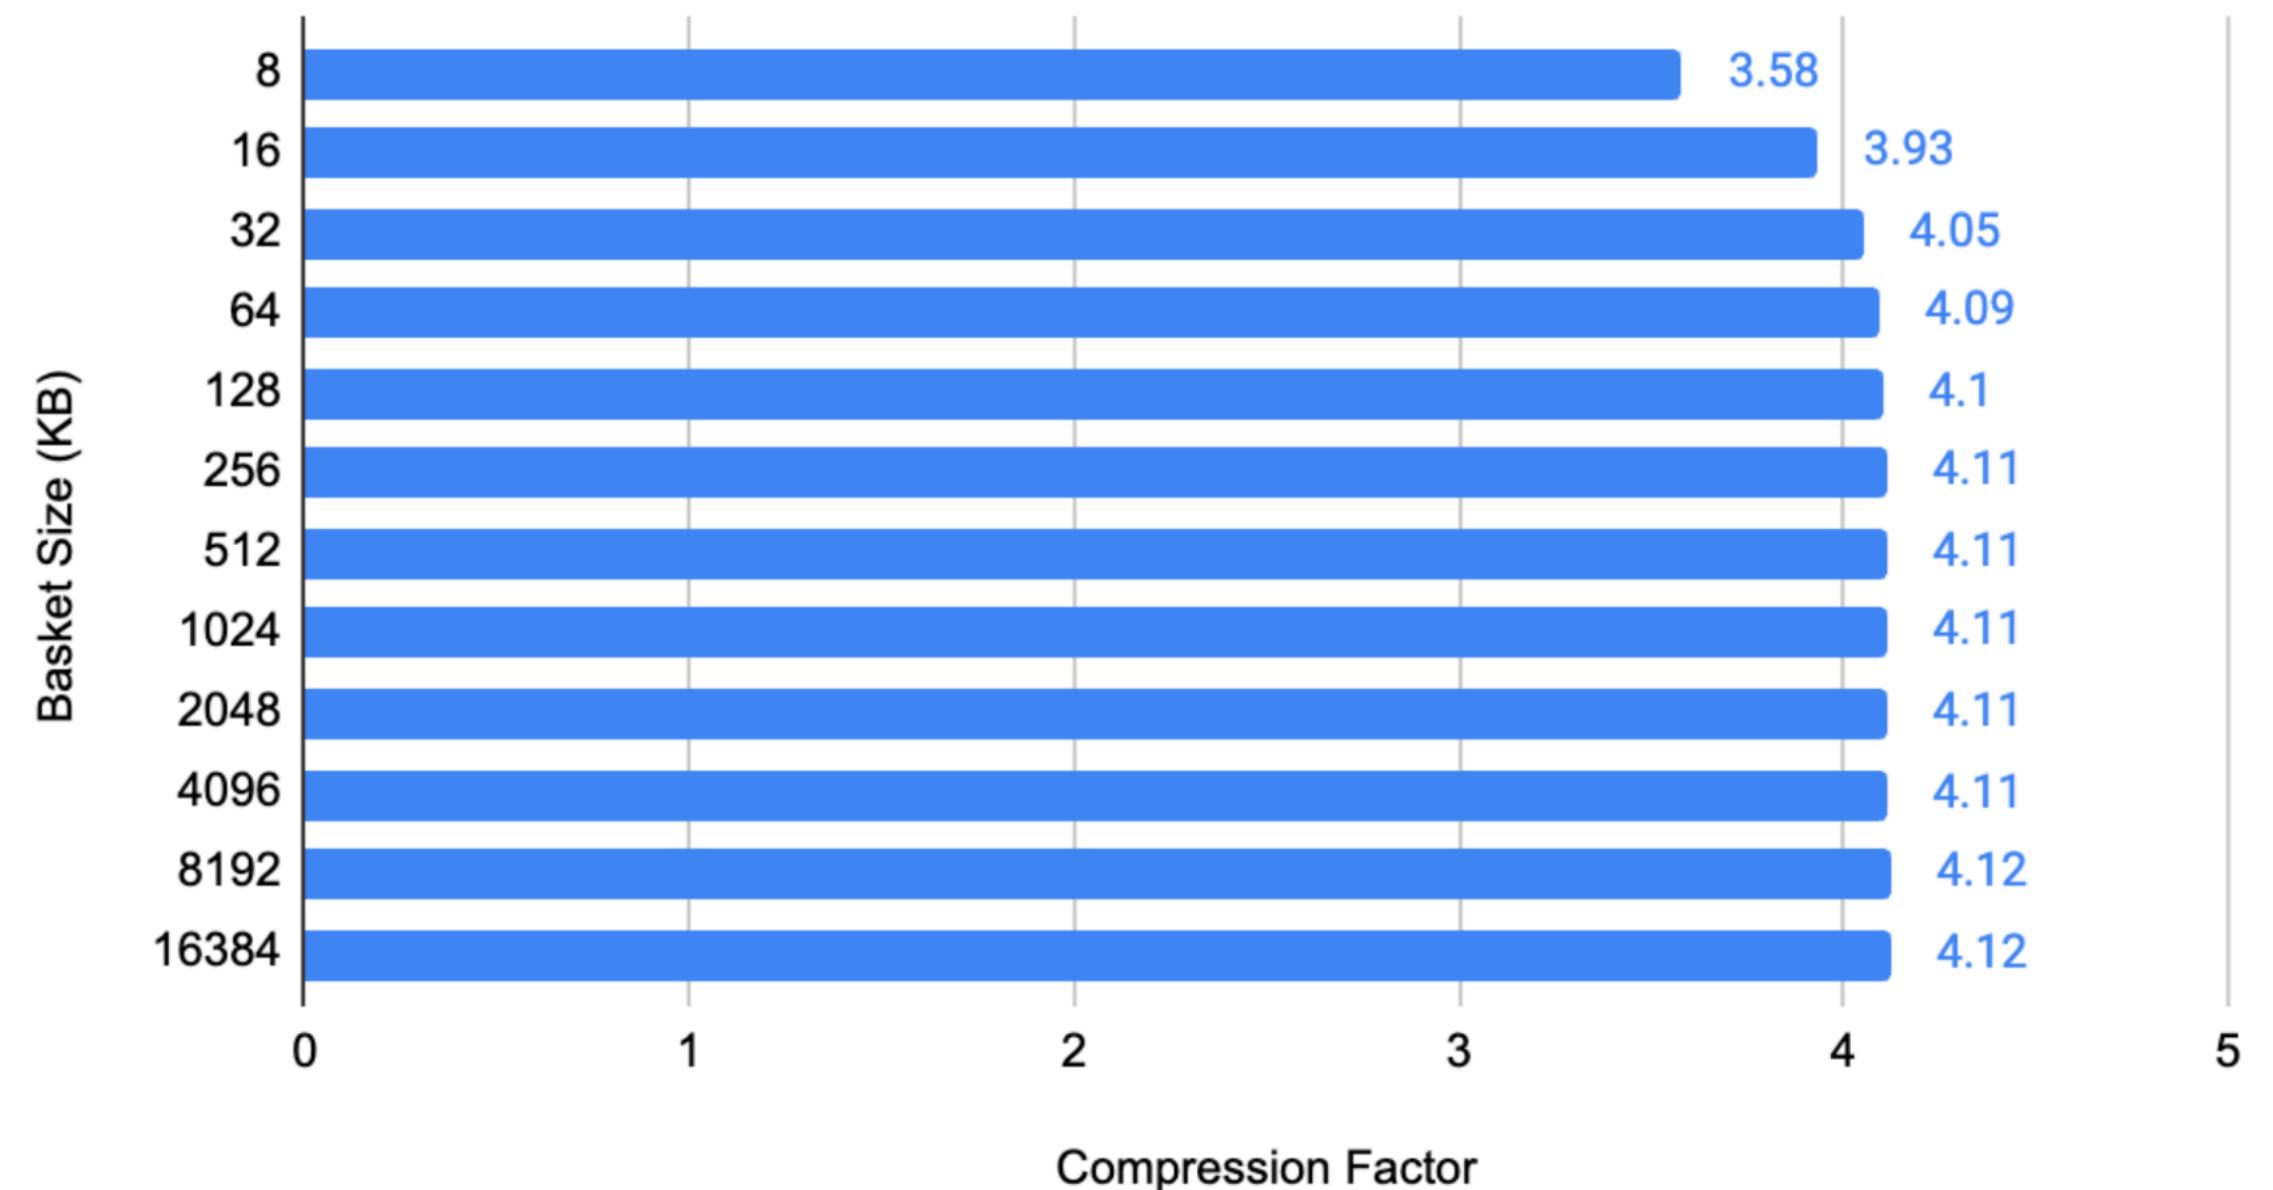
Toy Model I/O Tests

During the writing of these branches, the setting was done on within the script that filled the TTree full of values with:

```
tree->SetBasketSize(*, 8000) // 8 kB
```

But this is done during runtime... What if we wanted to set or remove the basket buffer size beforehand?

1/4 Random - Compression Factor vs. Basket Size (KB)



How can we reduce the load (memory/disk) to the Grid?

Currently, **ROOT** during runtime:

- Optimizes baskets after tree->Fill() is called, which assigns a basket buffer size optimized for the tree at that time. Maximum of 1.3 MB

Athena before runtime:

- Has several settings, but the two we're looking at are:
 - Maximum Basket Buffer Size: 128 kB
 - Minimum Number of Entries to the Buffer: 10 entries

For both ROOT and Athena: What happens to performance when we fix the maximum buffer limit prior to runtime memory allocation for the job occurs?

For this we'll need to run real derivation jobs...

Then we can get to:

**Applying changes to derivation of real
AODs**

After running derivation jobs: 128k-Basket / min-entries

Key | = 128kB_basket present [not present]
| = min_entries=10 present [not present]

Confias	Athena v22.0.16 Data Derivations		Output File Size (GB) [$\Delta\%$]	PHYSLITE
	Max PSS (GB) [$\Delta\%$ default]	PHYS + PHYSLITE		
	27.11		3.216	1.034
	27.81 [+ 2.5]		3.222 [+ 0.2]	1.036 [+ 0.2]
	27.81 [+ 2.5]		3.216 [- 0.0]	1.030 [- 0.4]
	27.30 [+ 0.7]		3.221 [+ 0.2]	1.042 [+ 0.7]

Configs	Athena v22.0.16 MC Derivations		Output File Size (GB) [$\Delta\%$]	PHYSLITE
	Max PSS (GB) [$\Delta\%$ default]	PHYS + PHYSLITE		
	14.1		5.8	2.59
	16.1 [+ 12.1]		6.0 [+ 2.9]	2.7 [+ 5.1]
	16.0 [+ 11.5]		5.7 [- 2.8]	2.5 [- 5.6]
	14.2 [+ 0.4]		6.2 [+ 5.4]	2.9 [+ 9.9]

Default configuration

Testing: 128k-Basket / min-entries

Key = 128kB_basket present [not present]
 = min_entries=10 present [not present]

Athena v22.0.16 Data Derivations				
Confias	Max PSS (GB) [$\Delta\%$ default]		Output File Size (GB) [$\Delta\%$]	
	PHYS + PHYSLITE	PHYS	PHYSLITE	
	27.11	3.216	1.034	
	27.81 [+ 2.5]	3.222 [+ 0.2]	1.036 [+ 0.2]	
	27.81 [+ 2.5]	3.216 [- 0.0]	1.030 [- 0.4]	
	27.30 [+ 0.7]	3.221 [+ 0.2]	1.042 [+ 0.7]	

- Data derivation jobs see little effect. Except a bit of memory increase when removing the basket limit.

Athena v22.0.16 MC Derivations				
Configs	Max PSS (GB) [$\Delta\%$ default]		Output File Size (GB) [$\Delta\%$]	
	PHYS + PHYSLITE	PHYS	PHYSLITE	
	14.1	5.8	2.59	
	16.1 [+ 12.1]	6.0 [+ 2.9]	2.7 [+ 5.1]	
	16.0 [+ 11.5]	5.7 [- 2.8]	2.5 [- 5.6]	
	14.2 [+ 0.4]	6.2 [+ 5.4]	2.9 [+ 9.9]	

Testing: 128k-Basket / min-entries

Key = 128kB_basket present [not present]
 = min_entries=10 present [not present]

Configs	Athena v22.0.16 Data Derivations			
	Max PSS (GB) [$\Delta\%$ default]		Output File Size (GB) [$\Delta\%$]	
	<u>PHYS + PHYSLITE</u>	<u>PHYS</u>	<u>PHYSLITE</u>	
	27.11	3.216	1.034	
	27.81 [+ 2.5]	3.222 [+ 0.2]	1.036 [+ 0.2]	
	27.81 [+ 2.5]	3.216 [- 0.0]	1.030 [- 0.4]	
	27.30 [+ 0.7]	3.221 [+ 0.2]	1.042 [+ 0.7]	

Configs	Athena v22.0.16 MC Derivations			
	Max PSS (GB) [$\Delta\%$ default]		Output File Size (GB) [$\Delta\%$]	
	<u>PHYS + PHYSLITE</u>	<u>PHYS</u>	<u>PHYSLITE</u>	
	14.1	5.8	2.59	
	16.1 [+ 12.1]	6.0 [+ 2.9]	2.7 [+ 5.1]	
	16.0 [+ 11.5]	5.7 [- 2.8]	2.5 [- 5.6]	
	14.2 [+ 0.4]	6.2 [+ 5.4]	2.9 [+ 9.9]	

- MC derivation sees a much larger effect

Reduction in the output file-size by (~5.6%) but at the price of much higher memory usage (> 11% for MC derivation)

Removing the basket cap in ROOT doesn't make sense, cannot assume every ROOT user can afford the memory increase this large. Ignore ROOT

Testing: 128k-Basket / min-entries

Key = 128kB_basket present [not present]
 = min_entries=10 present [not present]

Configs	Athena v22.0.16 Data Derivations	
	Max PSS (GB) [$\Delta\%$ default]	Output File Size (GB) [$\Delta\%$]
	<u>PHYS + PHYSLITE</u>	<u>PHYS</u> <u>PHYSLITE</u>
	27.11	3.216 1.034
	27.81 [+ 2.5]	3.222 [+ 0.2] 1.036 [+ 0.2]
	27.81 [+ 2.5]	3.216 [- 0.0] 1.030 [- 0.4]
	27.30 [+ 0.7]	3.221 [+ 0.2] 1.042 [+ 0.7]

Configs	Athena v22.0.16 MC Derivations		
	Max PSS (GB) [$\Delta\%$ default]	Output File Size (GB) [$\Delta\%$]	
	<u>PHYS + PHYSLITE</u>	<u>PHYS</u> <u>PHYSLITE</u>	
	14.1	5.8 2.59	
	16.1 [+ 12.1]	6.0 [+ 2.9] 2.7 [+ 5.1]	
	16.0 [+ 11.5]	5.7 [- 2.8] 2.5 [- 5.6]	
	14.2 [+ 0.4]	6.2 [+ 5.4] 2.9 [+ 9.9]	

- MC derivation sees a much larger effect

When both features were first implemented, it just so happened to have the opposite effect—**hit a happy medium! Keep Both!**

When only changing Athena: just how good is choosing 128 kB as a basket limit?

MC Derivation by Altered Basket Sizes			
Basket sizes	MaxPSS (GB) [$\Delta\%$ default]	PHYS Size (GB) [$\Delta\%$]	PHYSLITE Size (GB) [$\Delta\%$]
(default) 128k_basket	15.0	5.88	2.59
512k_basket	16.4 [+ 8.6]	5.74 [- 2.5]	2.46 [- 5.1]
no_limit	16.9 [+ 11.3]	5.72 [- 2.8]	2.45 [- 5.6]

Pretty good!

Any larger and we immediately start heading toward high-memory usage territory. Need to stay under budget for these jobs.

Interesting branches [MC PHYSLITE DAOD]

MC PHYS-LITE branch: HLTPNav_Summary_DAODSlimmedAuxDyn.decisions					
Configs	Baskets	Basket Size (KB)	Total Size (MB)	File Size (MB)	Compression
128kB_basket	11211	128	1862.8	417.0	4.47
no_limit	5223	1293.5	1861.9	342.4	5.44

Again: Basket size beyond 128kB doesn't massively improve compression [See [Toy Model](#)]

Interesting Branches	Basket size (kB)	Total-file-size (MB)	Compressed-file-size (MB)	Compression
PrimaryVerticesAuxDyn.trackParticleLinks	1293.5	2145.5	22.9	93.5
HLTPNav_Summary_DAODSlimmedAuxDyn.decisions	1293.5	1861.8	342.4	5.44
HardScatterVerticesAuxDyn.incomingParticleLinks	693.0	118.5	1.3	90.1
HardScatterVerticesAuxDyn.outgoingParticleLinks	635.5	108.5	1.5	74.0
AnalysisTauJetsAuxDyn.tauTrackLinks	447.0	74.9	1.9	39.2
Max = 1293.5 kB = 1.3 MB				

Interesting branches [MC PHYSLITE DAOD (no-lim)]

Contains mostly repeated numbers, highly compressible

Interesting Branches	Basket size (kB)	Total-file-size (MB)	Compressed-file-size (MB)	Compression
PrimaryVerticesAuxDyn.trackParticleLinks	1293.5	2145.5	22.9	93.5
HLTNav_Summary_DAODSlimmedAuxDyn.decisions	1293.5	1861.8	342.4	5.44
HardScatterVerticesAuxDyn.incomingParticleLinks	693.0	118.5	1.3	90.1
HardScatterVerticesAuxDyn.outgoingParticleLinks	635.5	108.5	1.5	74.0
AnalysisTauJetsAuxDyn.tauTrackLinks	447.0	74.9	1.9	39.2

Max = 1293.5 kB = 1.3 MB

Poorly compressed, Large file, reaches ROOT basket limit!

Other highly compressible branches

Conclusion to branch buffer size studies:

Removing the basket limit, a greater decrease in DAOD output file size is achieved at the expense of greater memory usage. However, what we're currently doing is good, which wasn't explored in detail before.

Finding these highly compressible branches leads us to want to look further into these branches to see what can be done to further optimize beyond derivation production.

Finally

**Creating new unit tests covering new I/O
core functionality**

Transient/Persistent EDM and the xAOD EDM

Transient/Persistent (T/P):

- Transient data is located in memory and is readily accessible to the user.
- While Persistent data is saved to disk and need to be converted into transient data before it can be accessible by the user.

Some problems:

- Transient data would often gain in complexity the more information attached to it, and ROOT had a hard time dealing with this complexity.
- Developers would have to convert data samples from the full EDM to simpler ones (more readily available to ROOT) which would lead to duplication of data.

Transient/Persistent EDM and the xAOD EDM

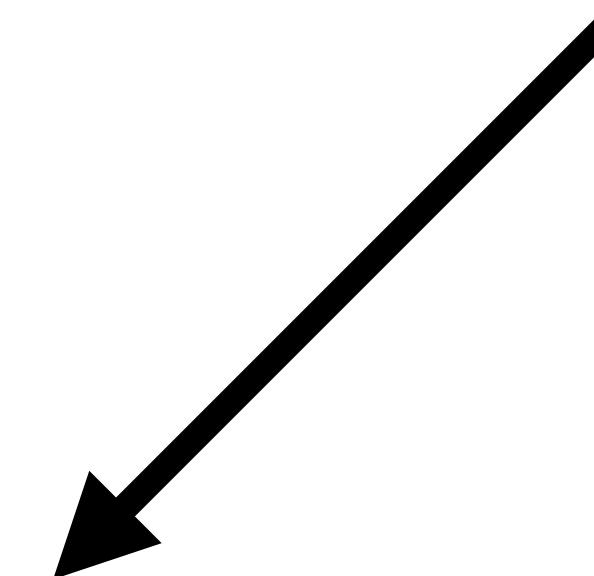
Transient/Persistent (T/P):

- Transient data is located in memory and is readily accessible to the user.
- While Persistent data is saved to disk and need to be converted into transient data before it can be accessible by the user.

Some problems:

- Transient data would often gain in complexity the more information attached to it, and ROOT had a hard time dealing with this complexity.
- Developers would have to convert data samples from the full EDM to simpler ones (more readily available to ROOT) which would lead to duplication of data.

(From what I could find, they put 'x' in-front of the AOD to distinguish it as the new AOD, though I can't find the source)



xAOD EDM:

- Seeks to simplify things further in to interface objects and payload objects.

Transient/Persistent EDM and the xAOD EDM

Transient/Persistent (T/P):

- Transient data is located in memory and is readily accessible to the user.
- While Persistent data is saved to disk and need to be converted into transient data before it can be accessible by the user.

Some problems:

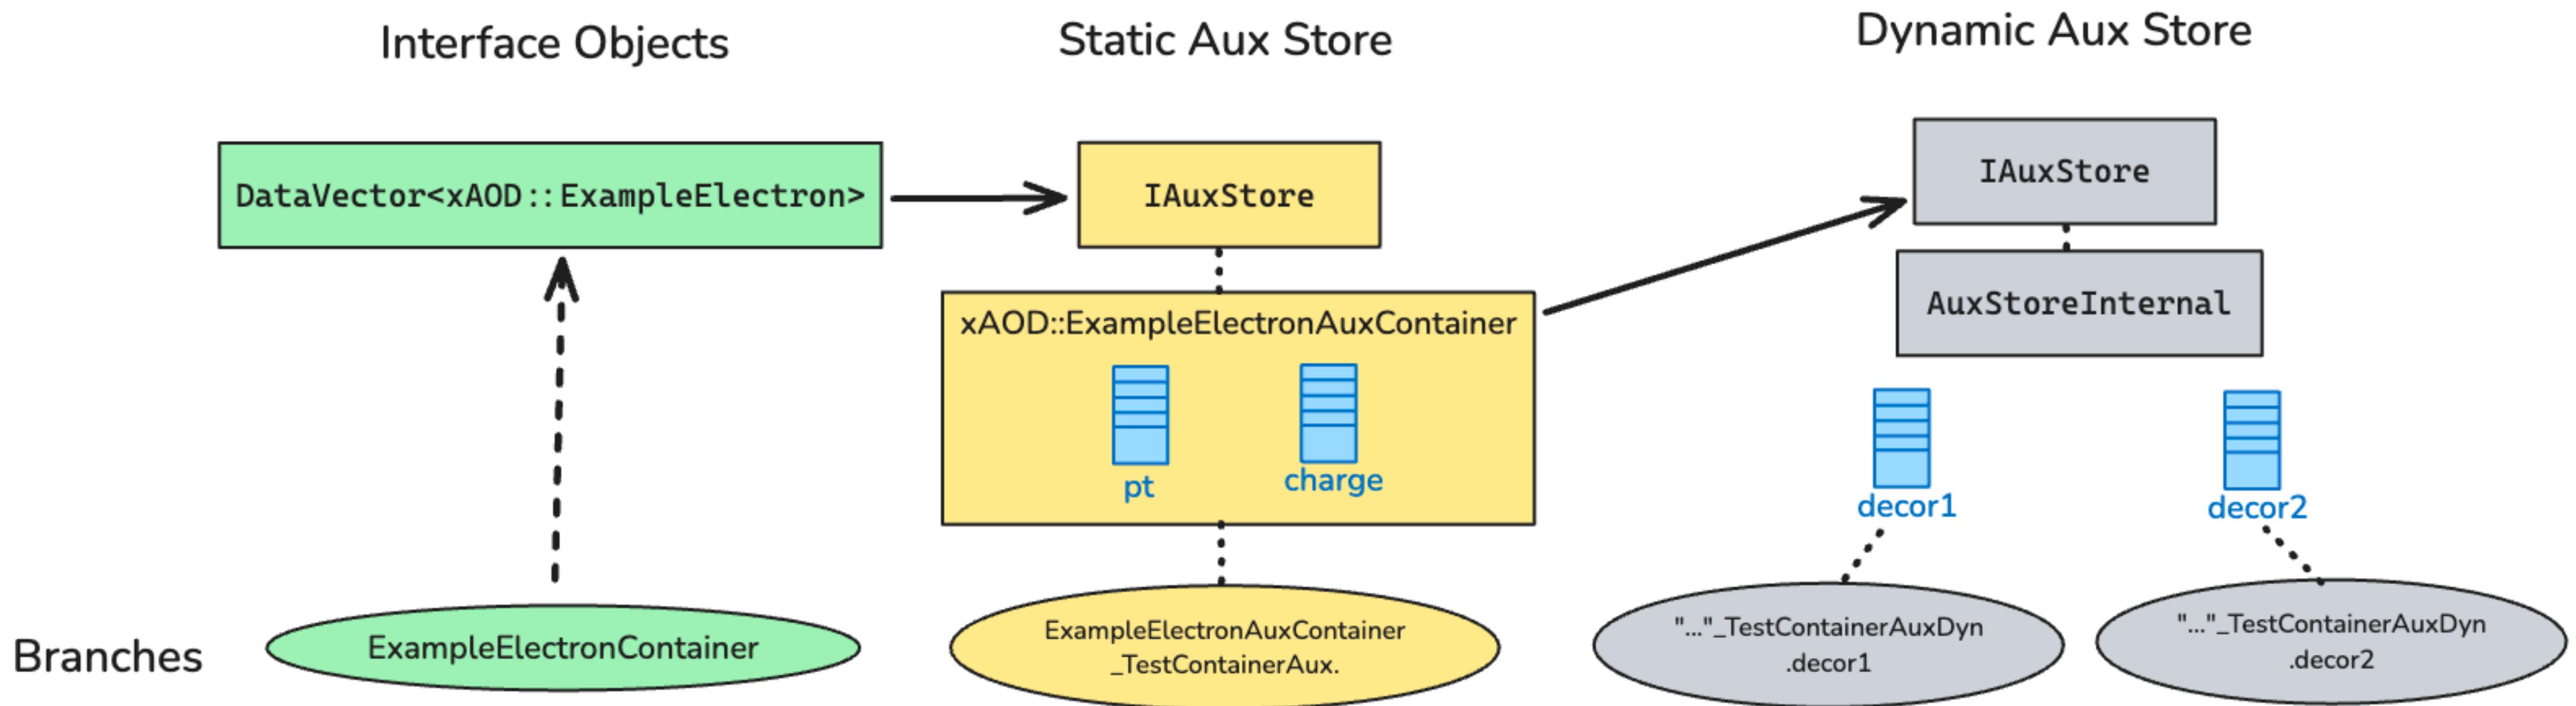
- Transient data would often gain in complexity the more information attached to it, and ROOT had a hard time dealing with this complexity.
- Developers would have to convert data samples from the full EDM to simpler ones (more readily available to ROOT) which would lead to duplication of data.

xAOD EDM:

- Seeks to simplify things further in to interface objects and payload objects.
- Interfaces allow the user to delay loading data into memory
- While the payload objects contain the data for the interface objects and allocate contiguous blocks of memory. Payload objects are synonymous with auxiliary storage.

Note: xAOD EDM was adopted during Run 2, so it's not *that* new—and T/P EDM is still used for certain things.

Transient/Persistent EDM and the xAOD EDM

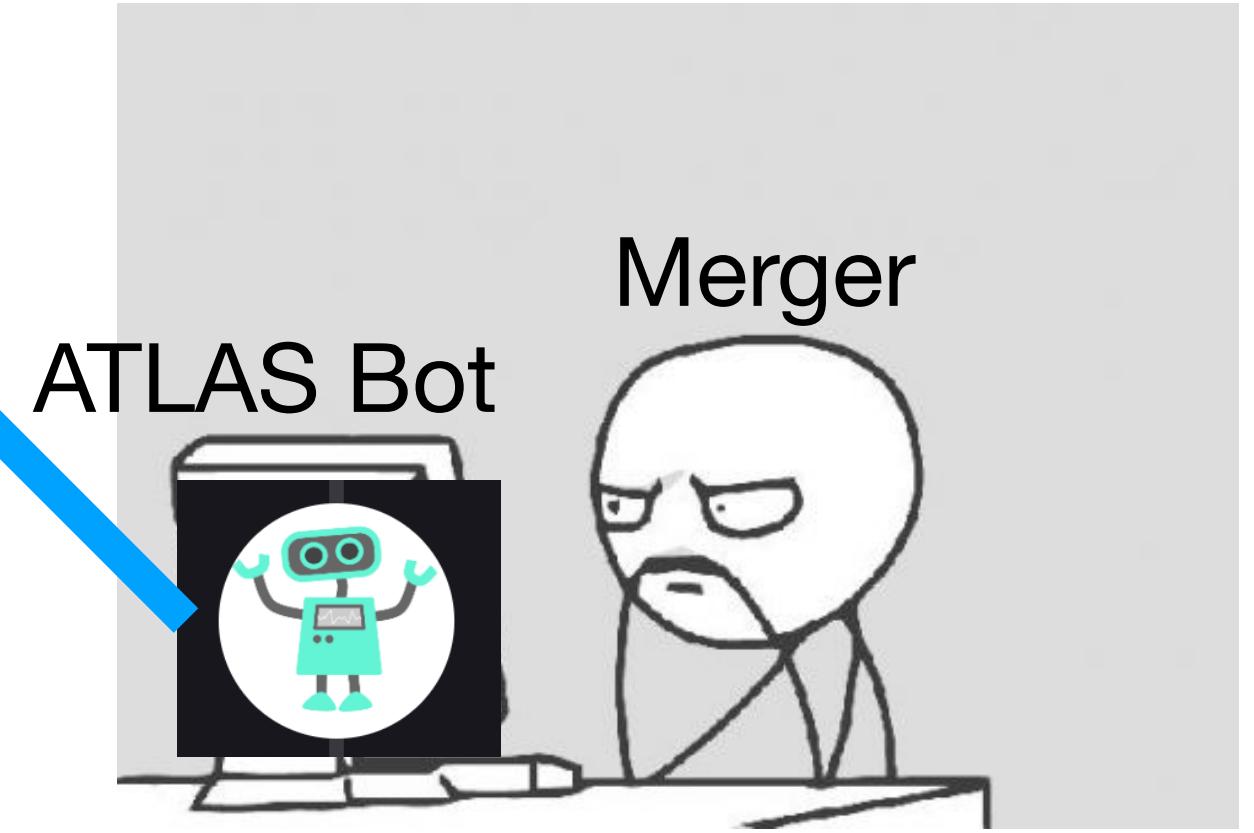
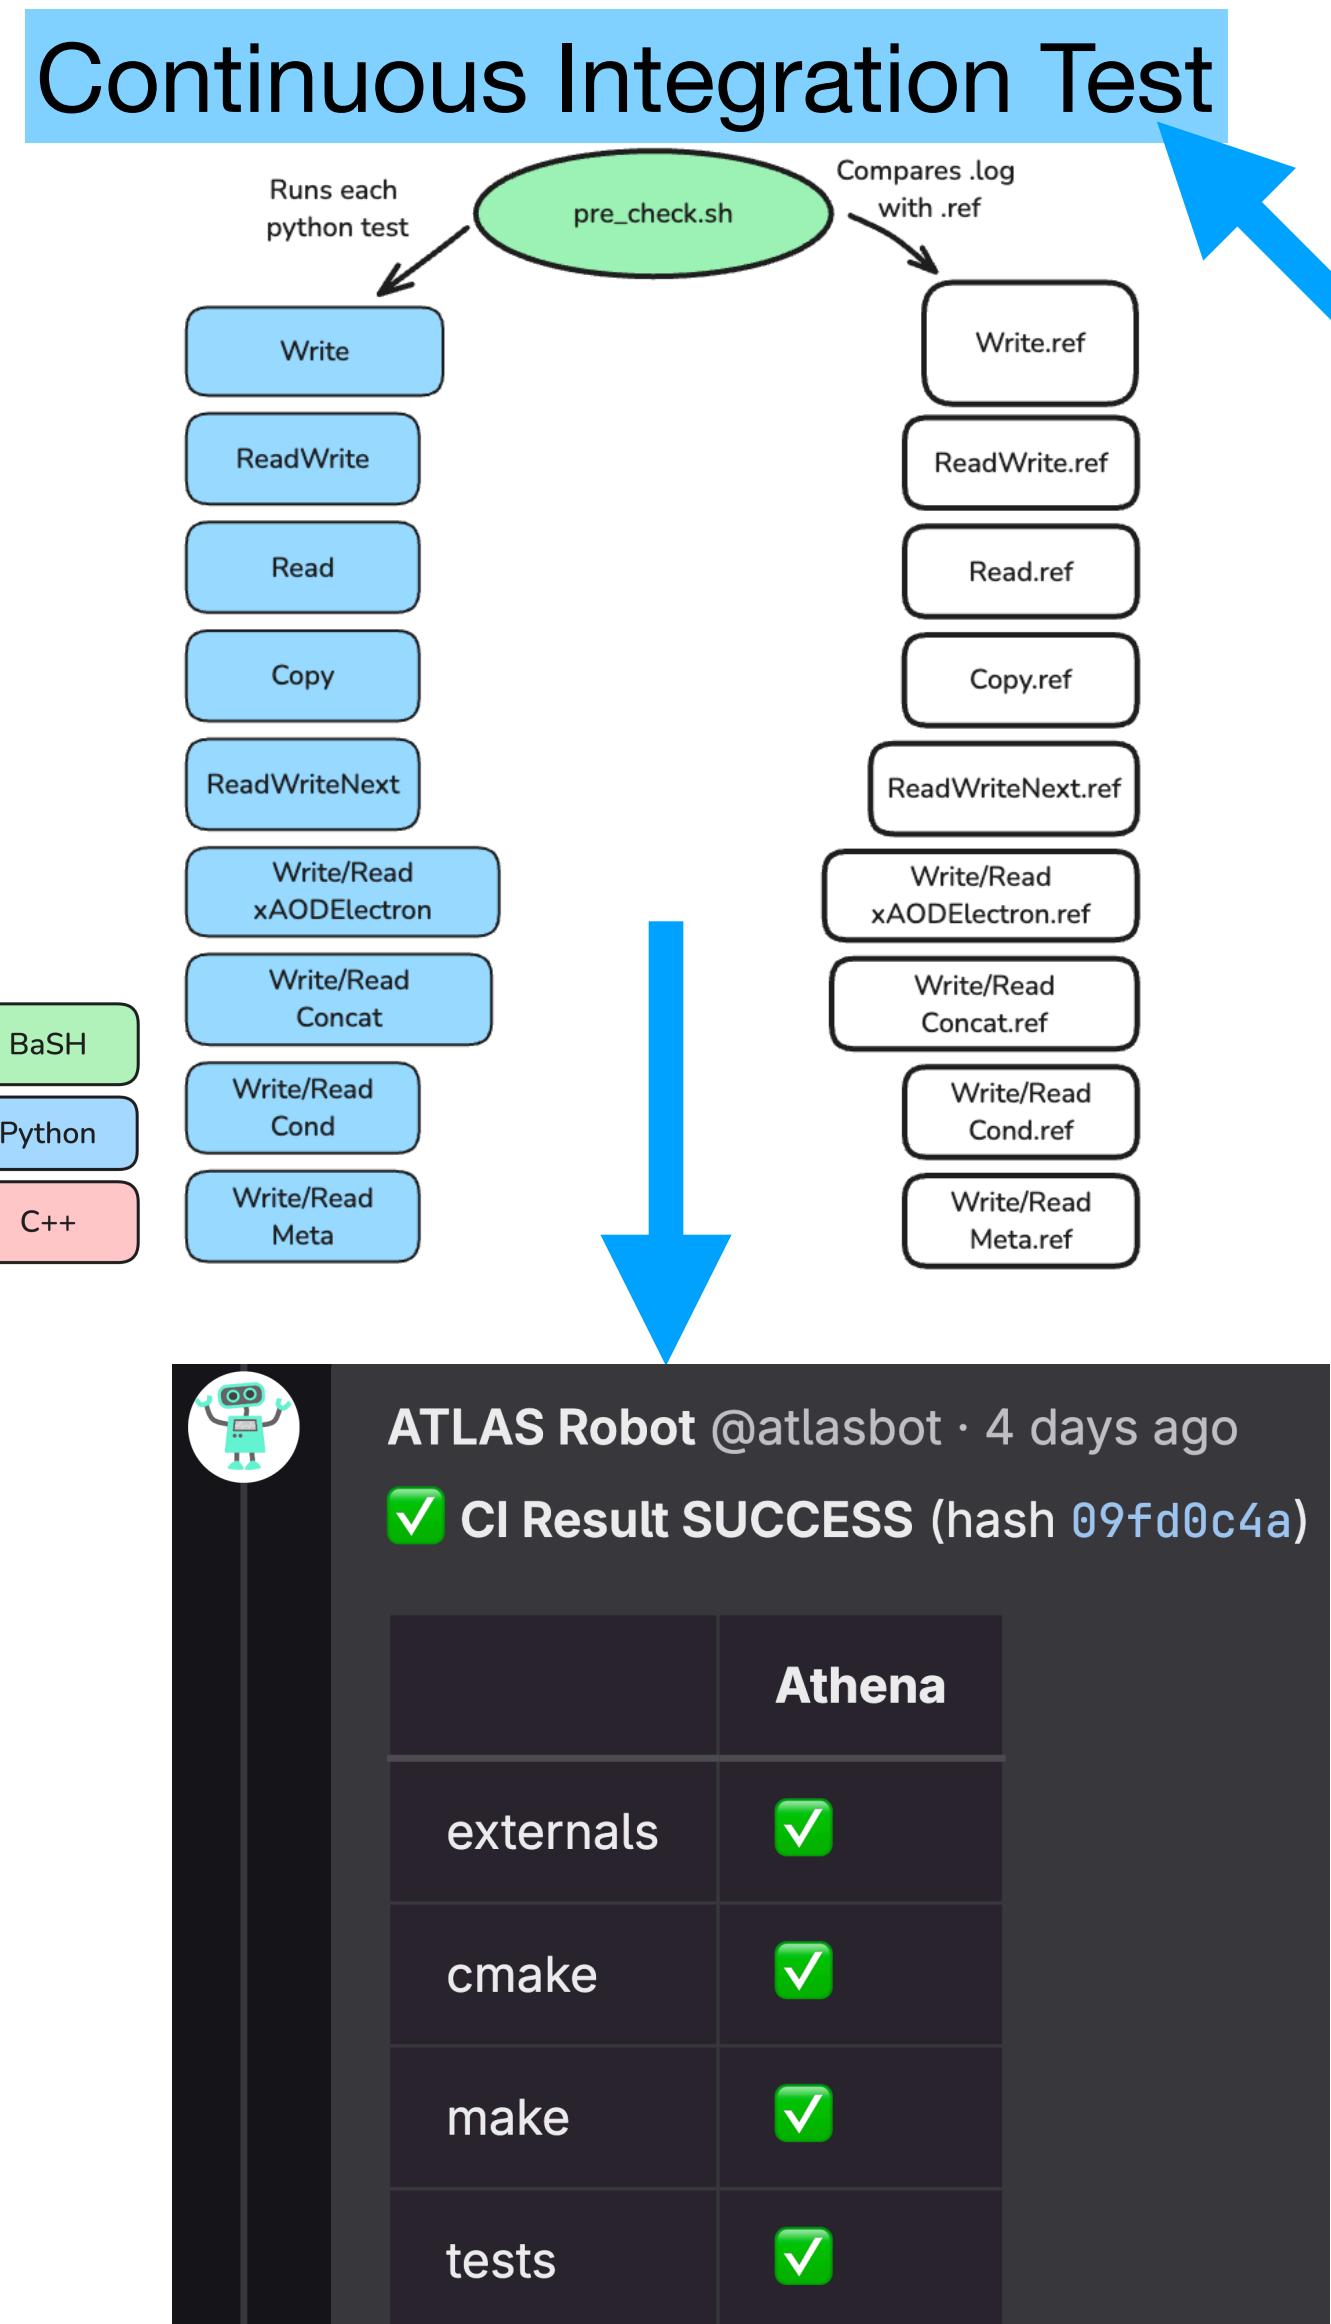


xAOD EDM:

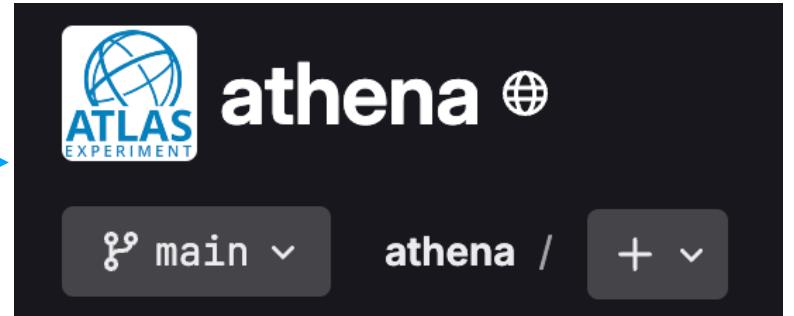
- Seeks to simplify things further in to interface objects and payload objects.
- Allows for decorations which can be readily accessed by ROOT and don't have as much of the complexity problem.

Continuous Integration Testing

- Athena is updated every night (a **nightly**) and built during continuous integration, but each time it's built it needs to reference unit tests to ensure safety of core features
- Unit tests can test core features across many areas, but the package we're focused on is core I/O functionality in AthenaPoolExample



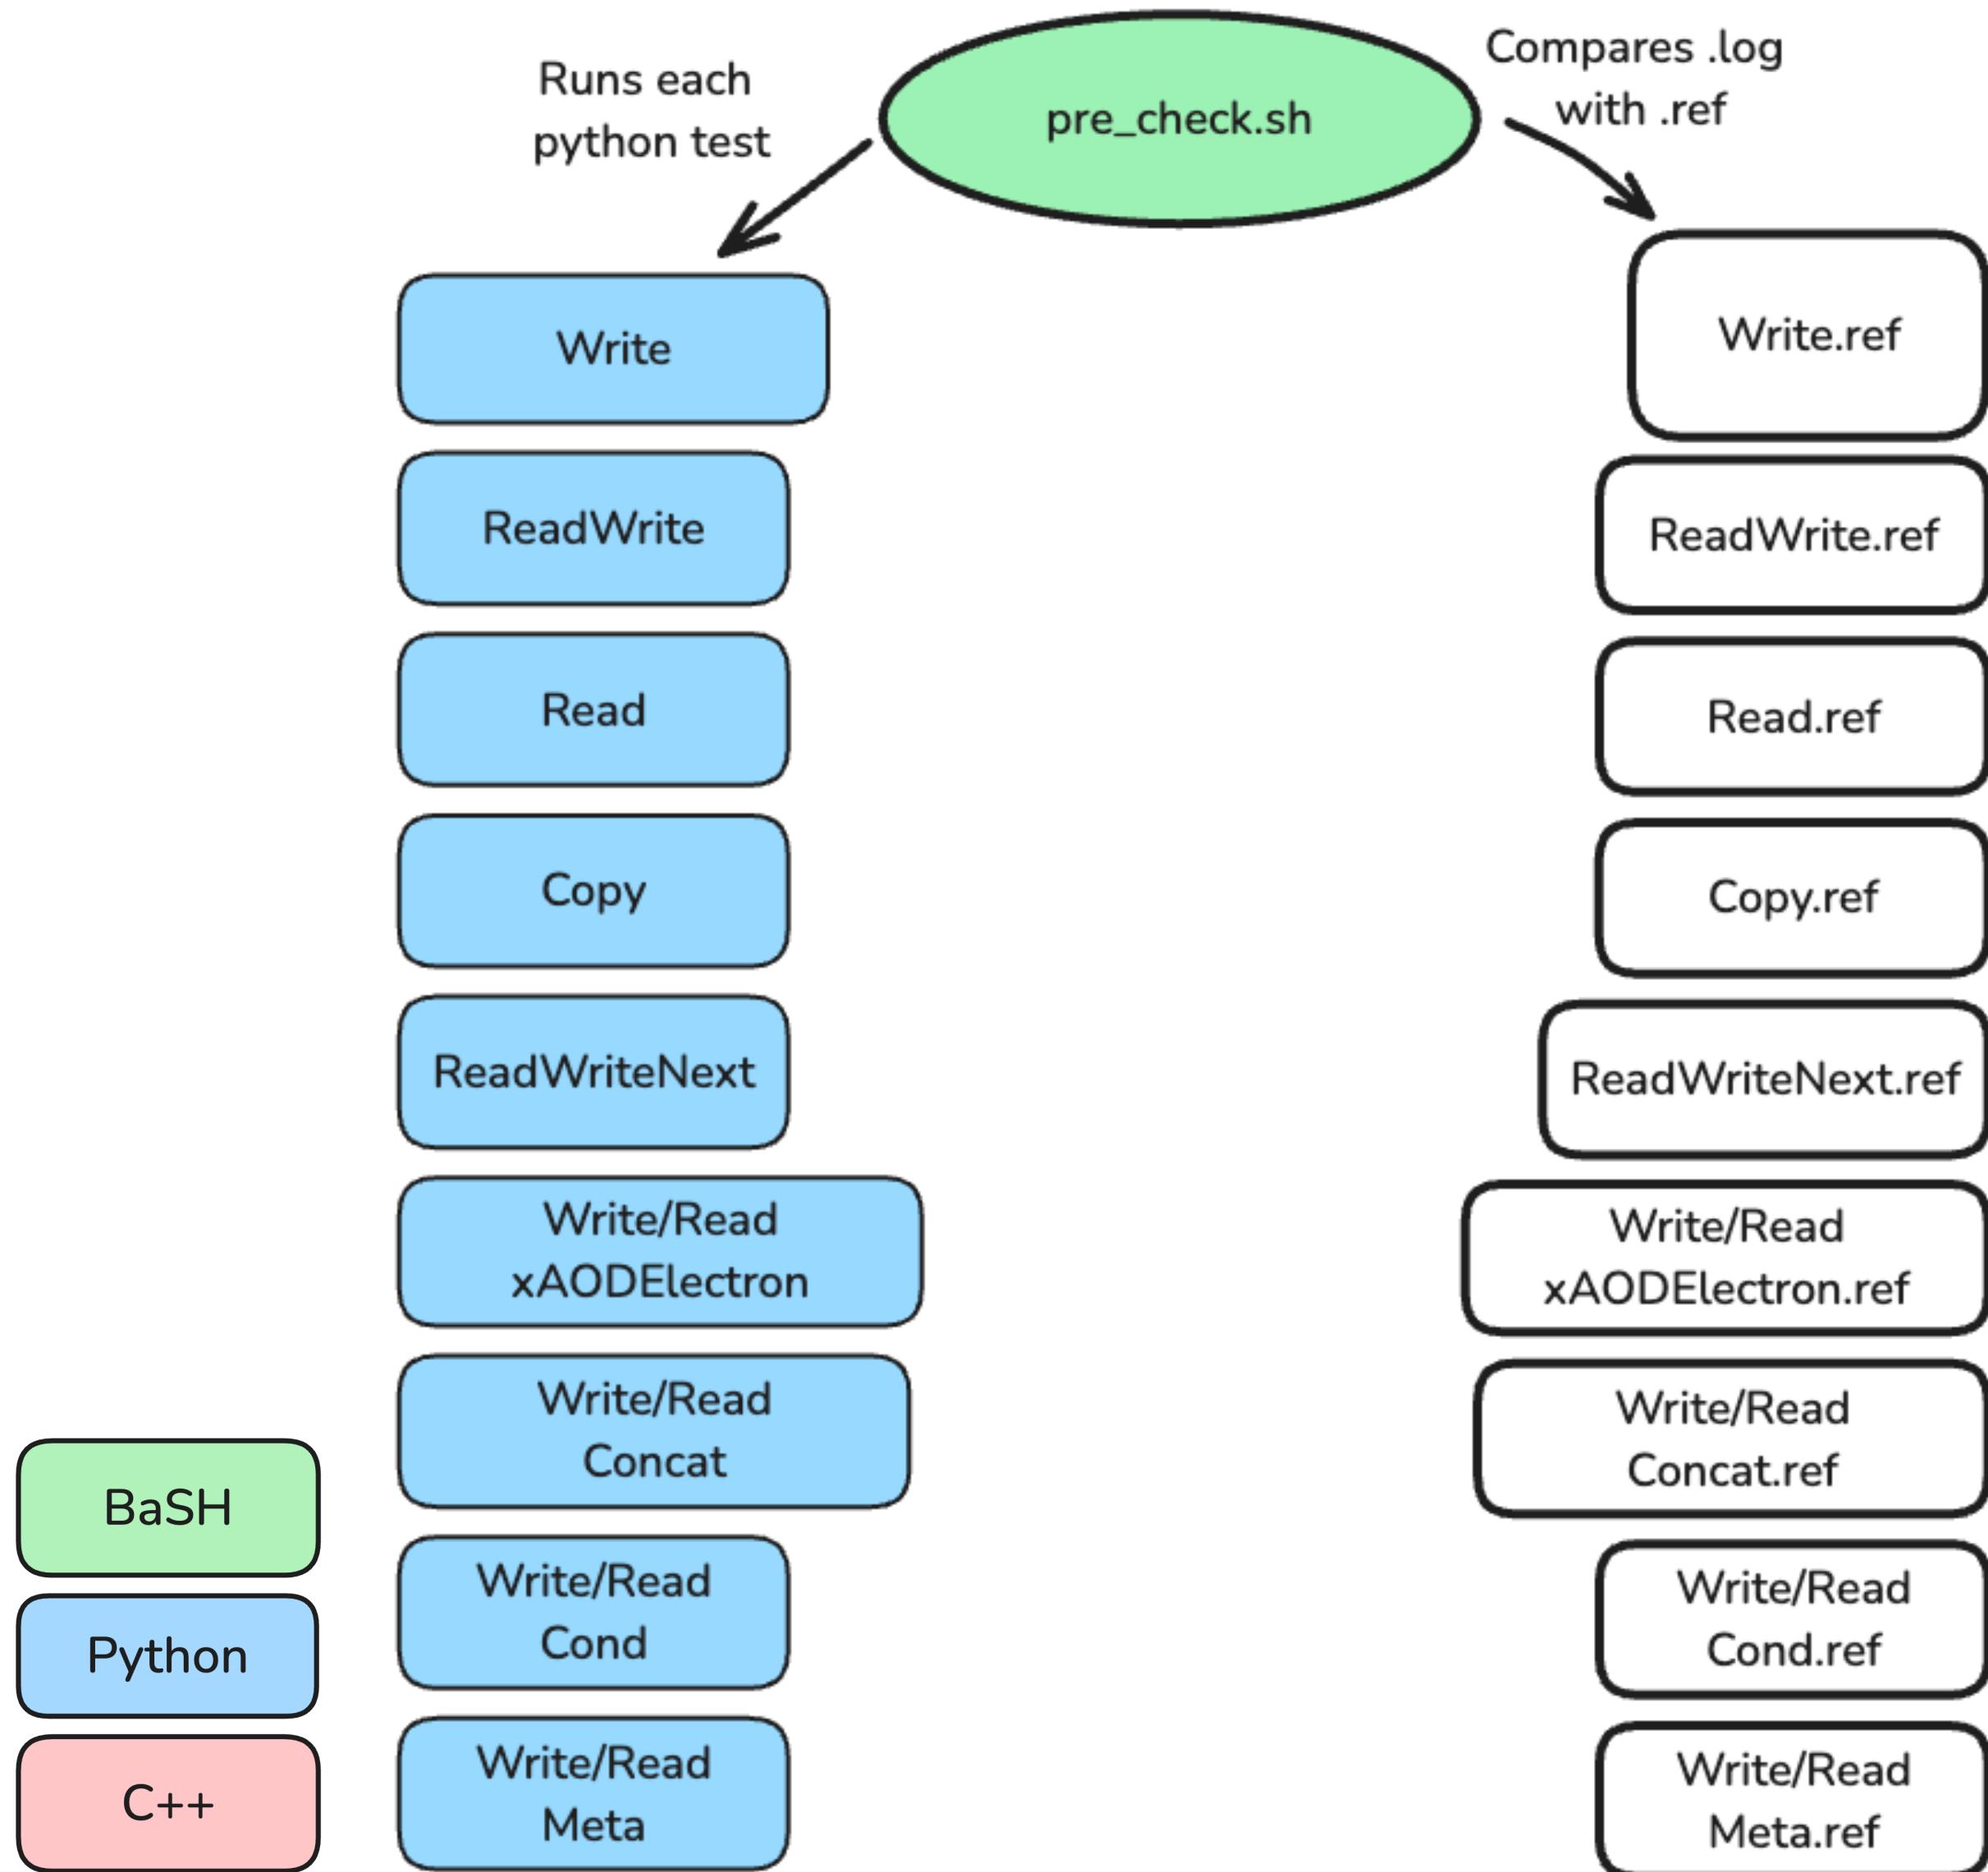
Merges with



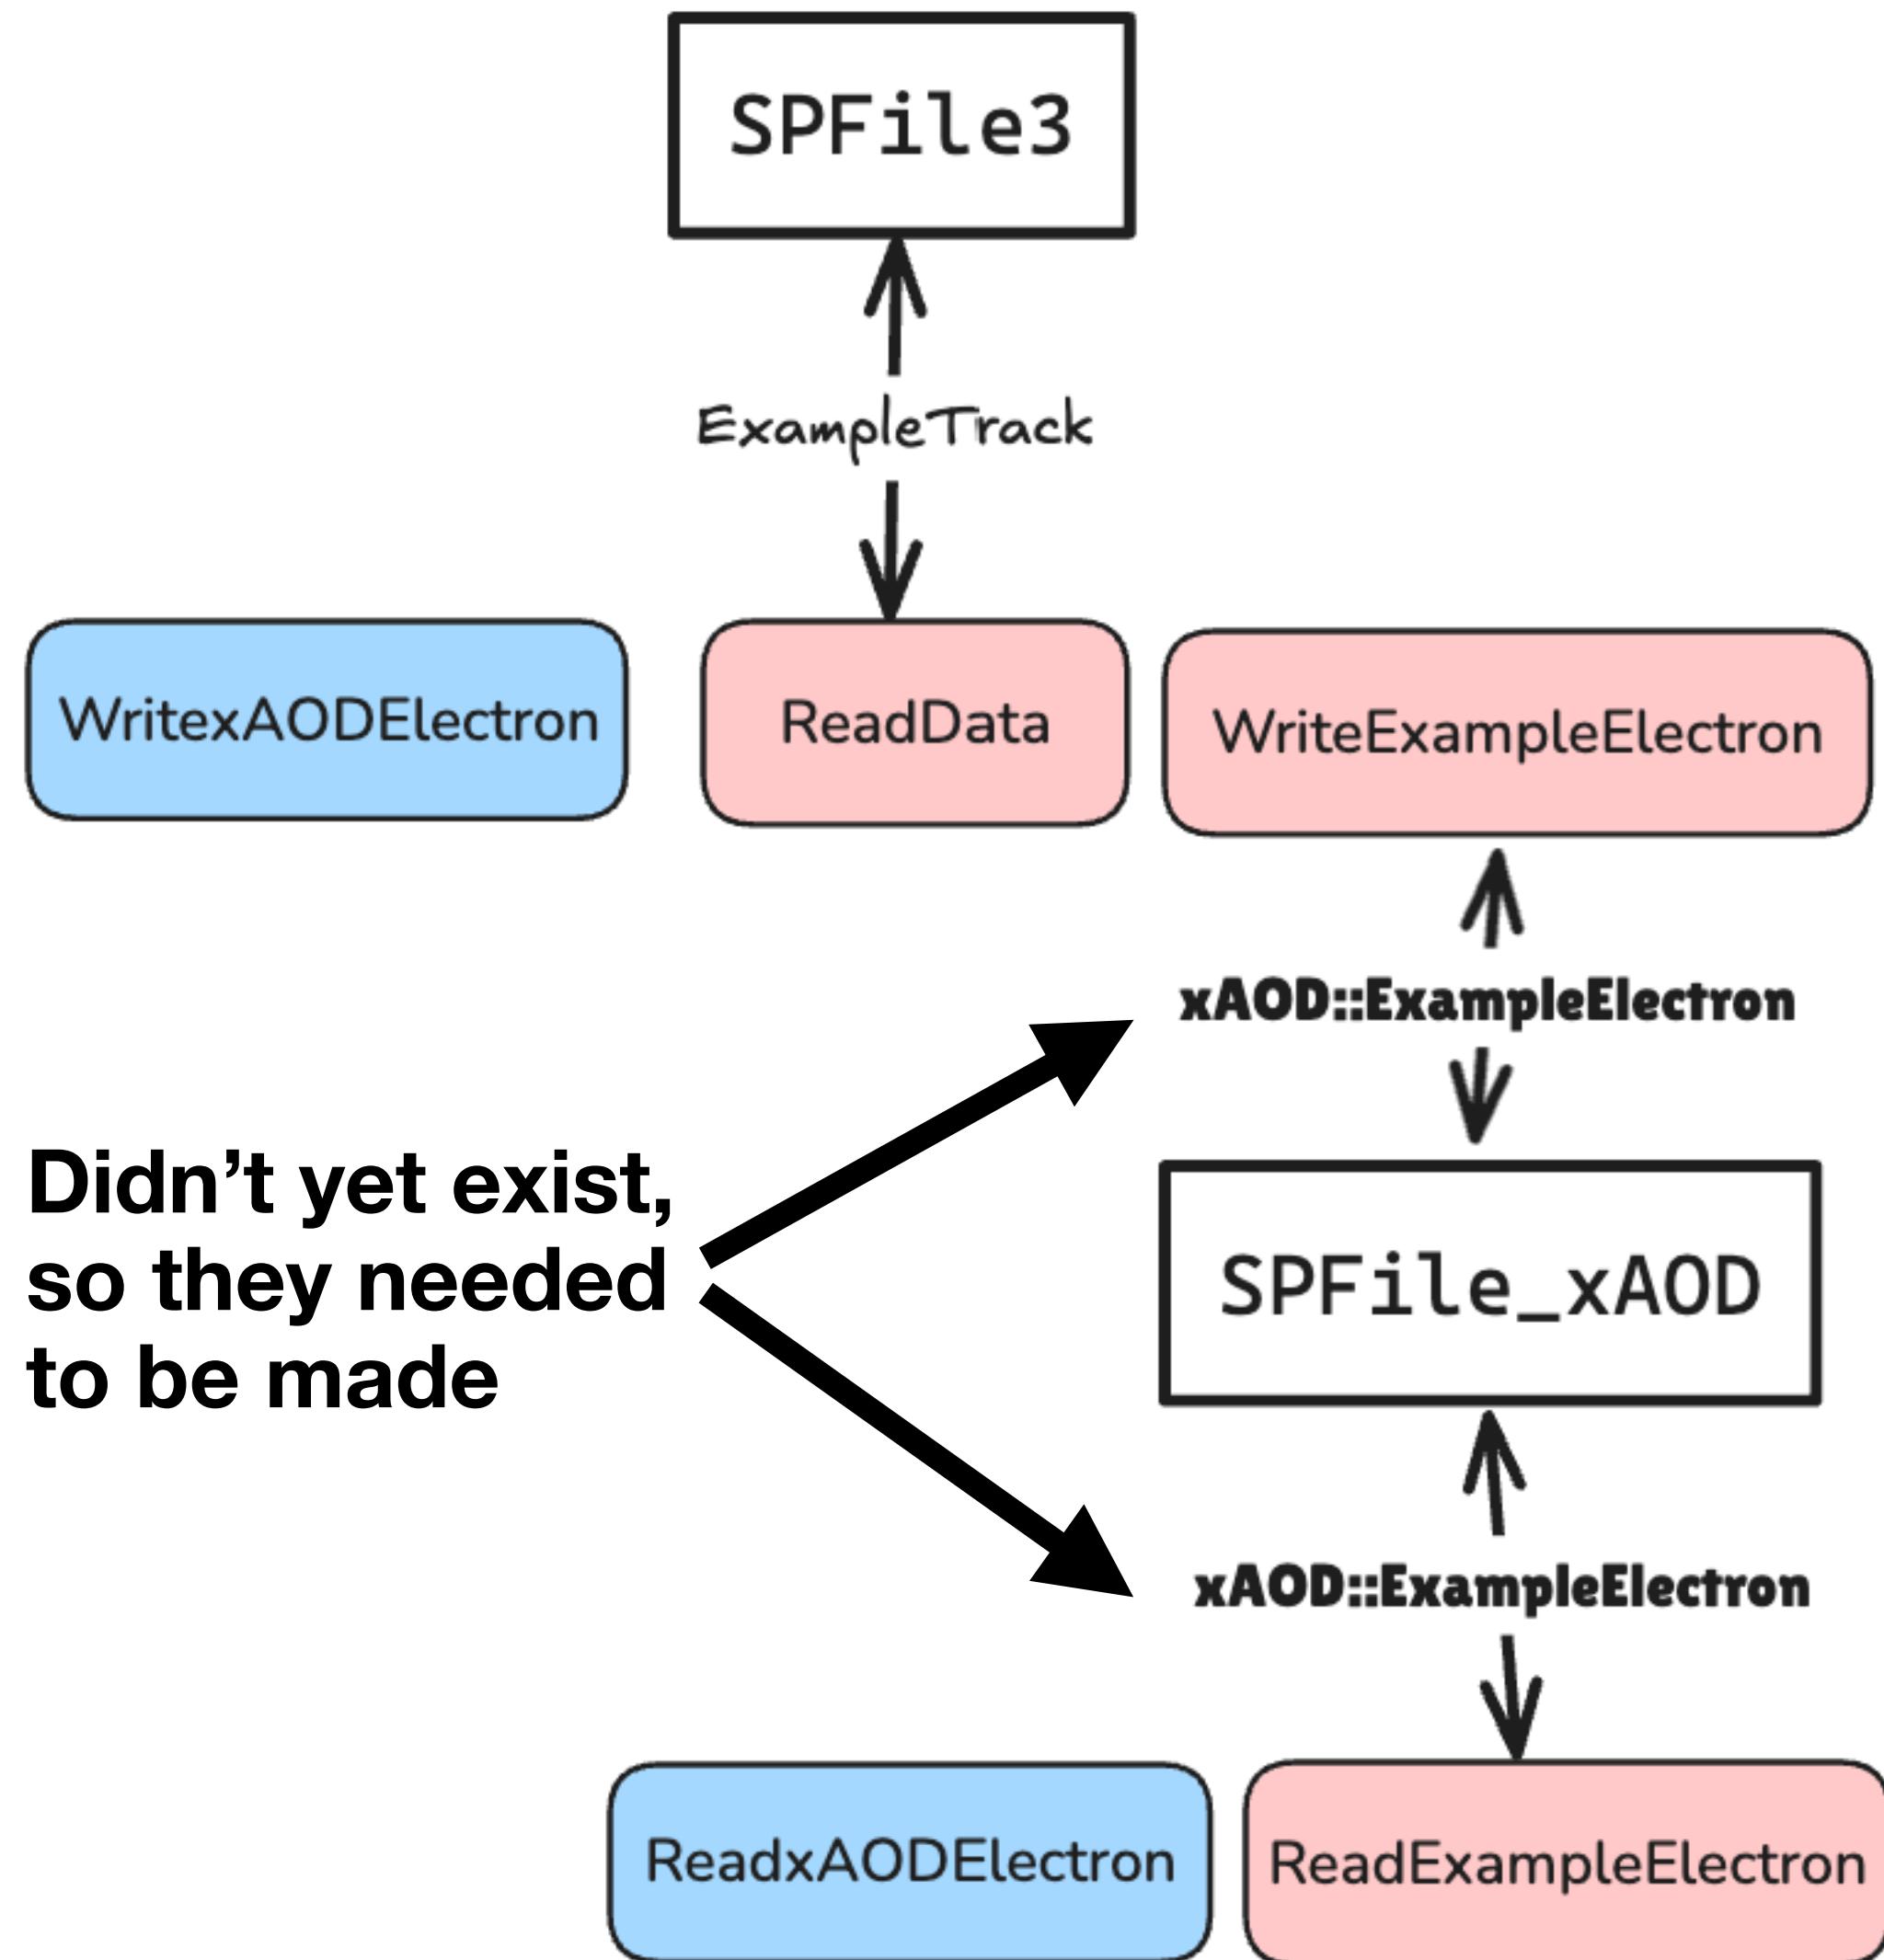
AthenaPoolExample

- There are a number of unit tests in the AthenaPoolExample package that test core I/O functionality, but we want to include one that tests the xAOD EDM specifically.

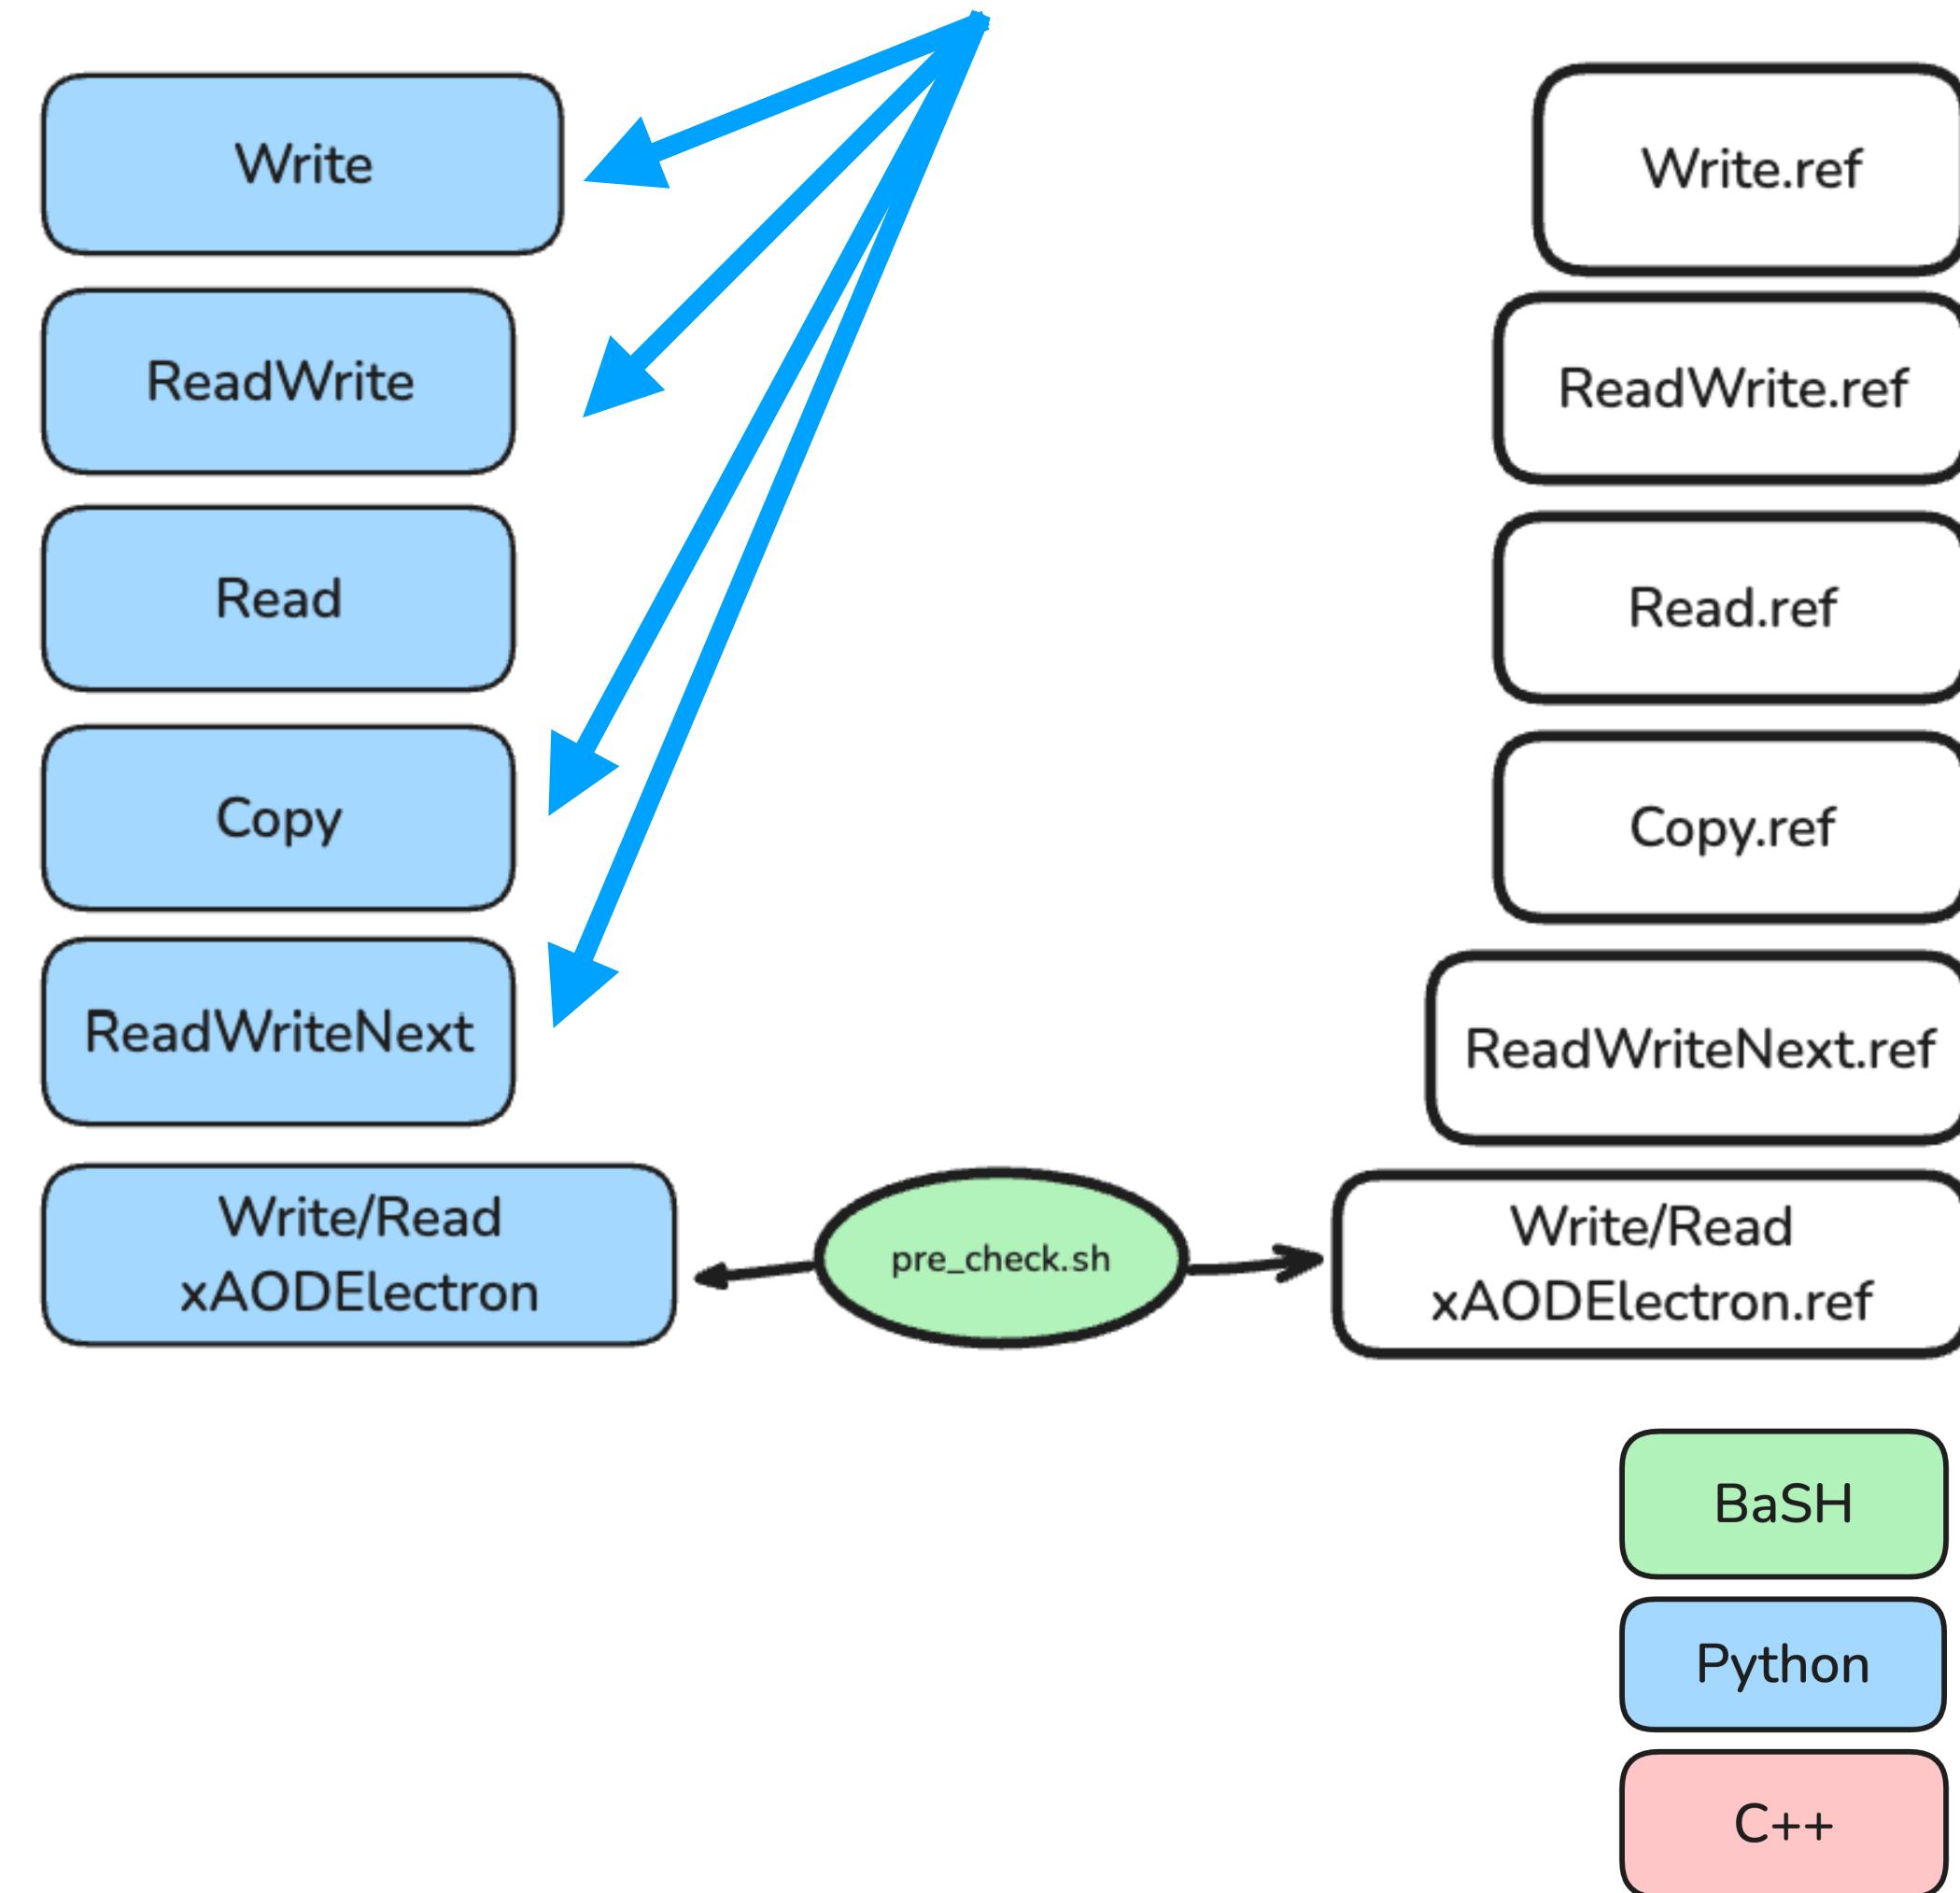
AthenaPoolExample Unit Tests



Write/ReadxAODElectron



All these SimplePool(SP) Files are created by the preceding unit tests.



xAOD::ExampleElectron Architecture

xAOD::ExampleElectrons are bespoke xAOD objects that try to achieve simplicity

Each of these have unique identifiers that link their 'xAOD::ExampleElectron____' to the dictionary that's accessible by ROOT

They also identify versioning by typed which version of the class to use. I.e.:

```
namespace xAOD {  
  
    typedef ExampleElectron_v1 ExampleElectron;  
  
} // namespace xAOD
```

[\[LXR link\]](#)

[\[Athena Gitlab link\]](#)

AthenaPoolExampleData
/src/

ExampleElectronContainer_v1.cxx
ExampleElectron_v1.cxx
/dict/

ContainerProxies.cxx
/AthenaPoolExampleData/
selection.xml
AthenaPoolExampleDataDict.h
ExampleElectron.h
ExampleElectronContainer.h
ExampleElectronAuxContainer.h
/versions/

ExampleElectron_v1.h
ExampleElectronContainer_v1.h
ExampleElectronAuxContainer_v1.h

xAOD::ExampleElectron Architecture

These could be enveloped in the non-versioned headers above, but this setup works and helps with versioning control a bit.

These versioned header files declare all necessary object types for the xAOD EDM. I.e.

Electron: getter/setters for pt and charge

ElectronContainer:

```
namespace xAOD {  
    typedef DataVector<ExampleElectron_v1> ExampleElectronContainer_v1;  
}
```

ElectronAuxContainer: Declaring AUX_VARS

```
private:  
    AUXVAR_DECL(double, pt);  
    AUXVAR_DECL(float, charge);
```

[\[LXR link\]](#)

[\[Athena Gitlab link\]](#)

AthenaPoolExampleData

/src/

ExampleElectronContainer_v1.cxx
ExampleElectron_v1.cxx
/dict/

ContainerProxies.cxx
/AthenaPoolExampleData/
selection.xml

AthenaPoolExampleDataDict.h
ExampleElectron.h

ExampleElectronContainer.h
ExampleElectronAuxContainer.h

/versions/

ExampleElectron_v1.h
ExampleElectronContainer_v1.h
ExampleElectronAuxContainer_v1.h

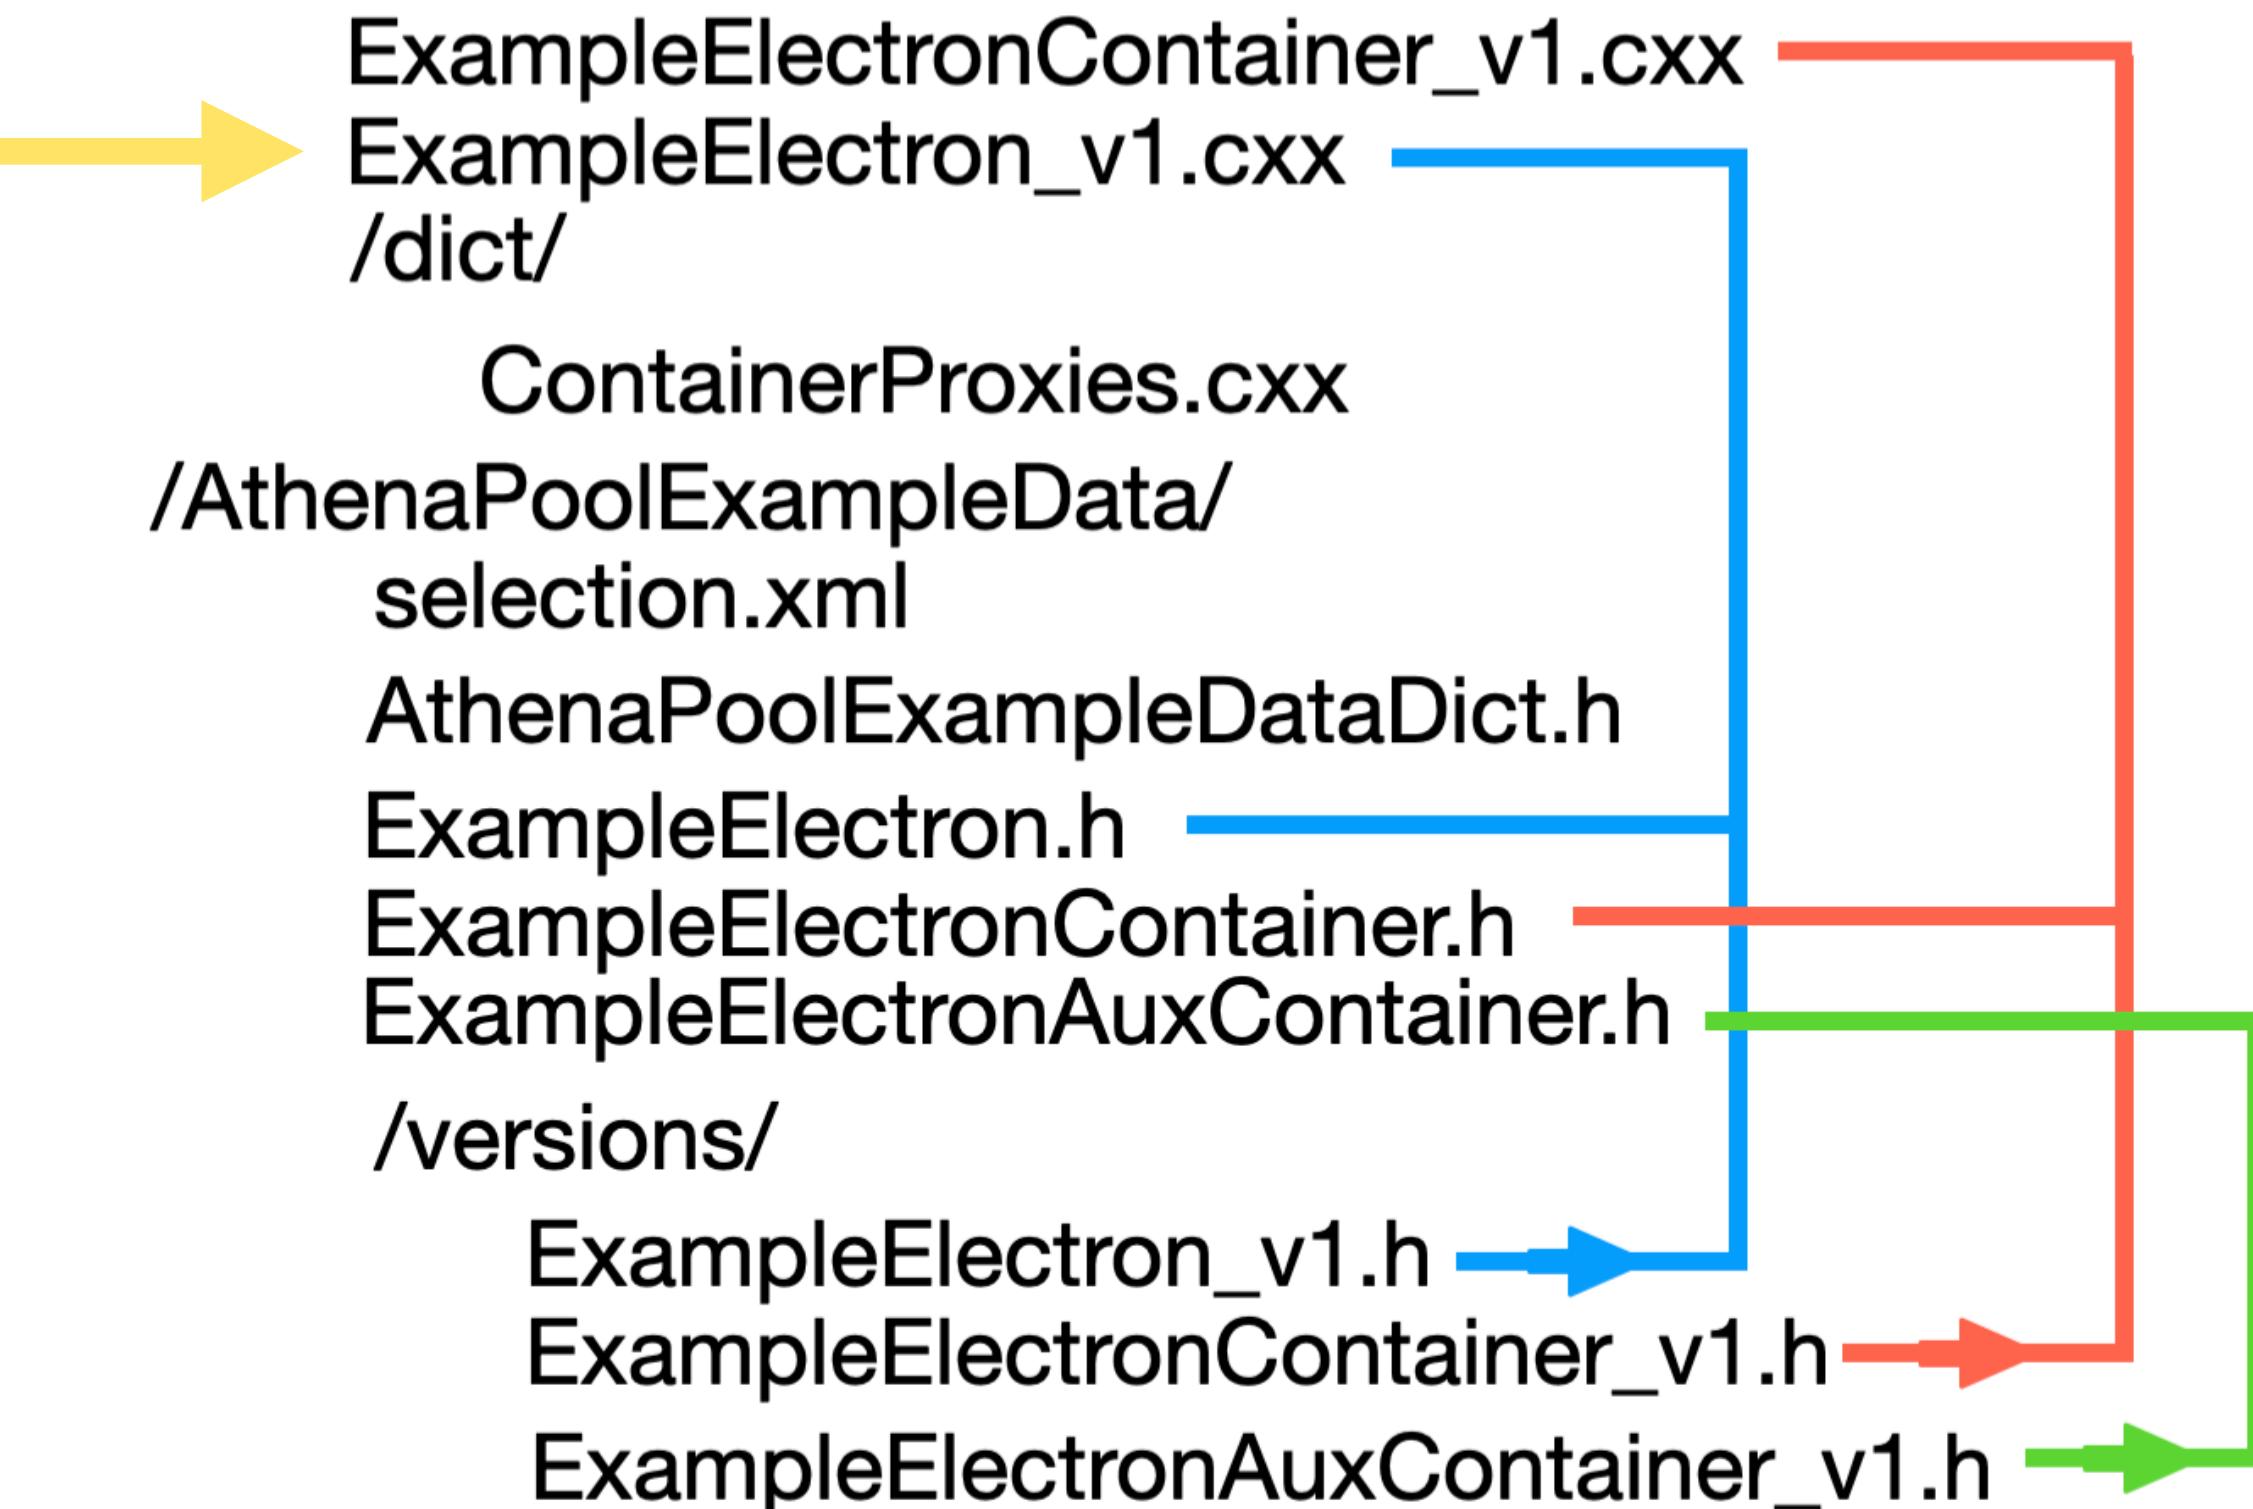
xAOD::ExampleElectron Architecture

You would probably need a **setter** and **getter** specific to the AuxStore.

```
#include "AthenaPoolExampleData/versions/ExampleElectron_v1.h"
#include "xAODCore/AuxStoreAccessorMacros.h"

namespace xAOD {
AUXSTORE_PRIMITIVE_SETTER_AND_GETTER(ExampleElectron_v1, double, pt, setPt)
AUXSTORE_PRIMITIVE_SETTER_AND_GETTER(ExampleElectron_v1, float, charge,
                                     setCharge)
} // namespace xAOD
```

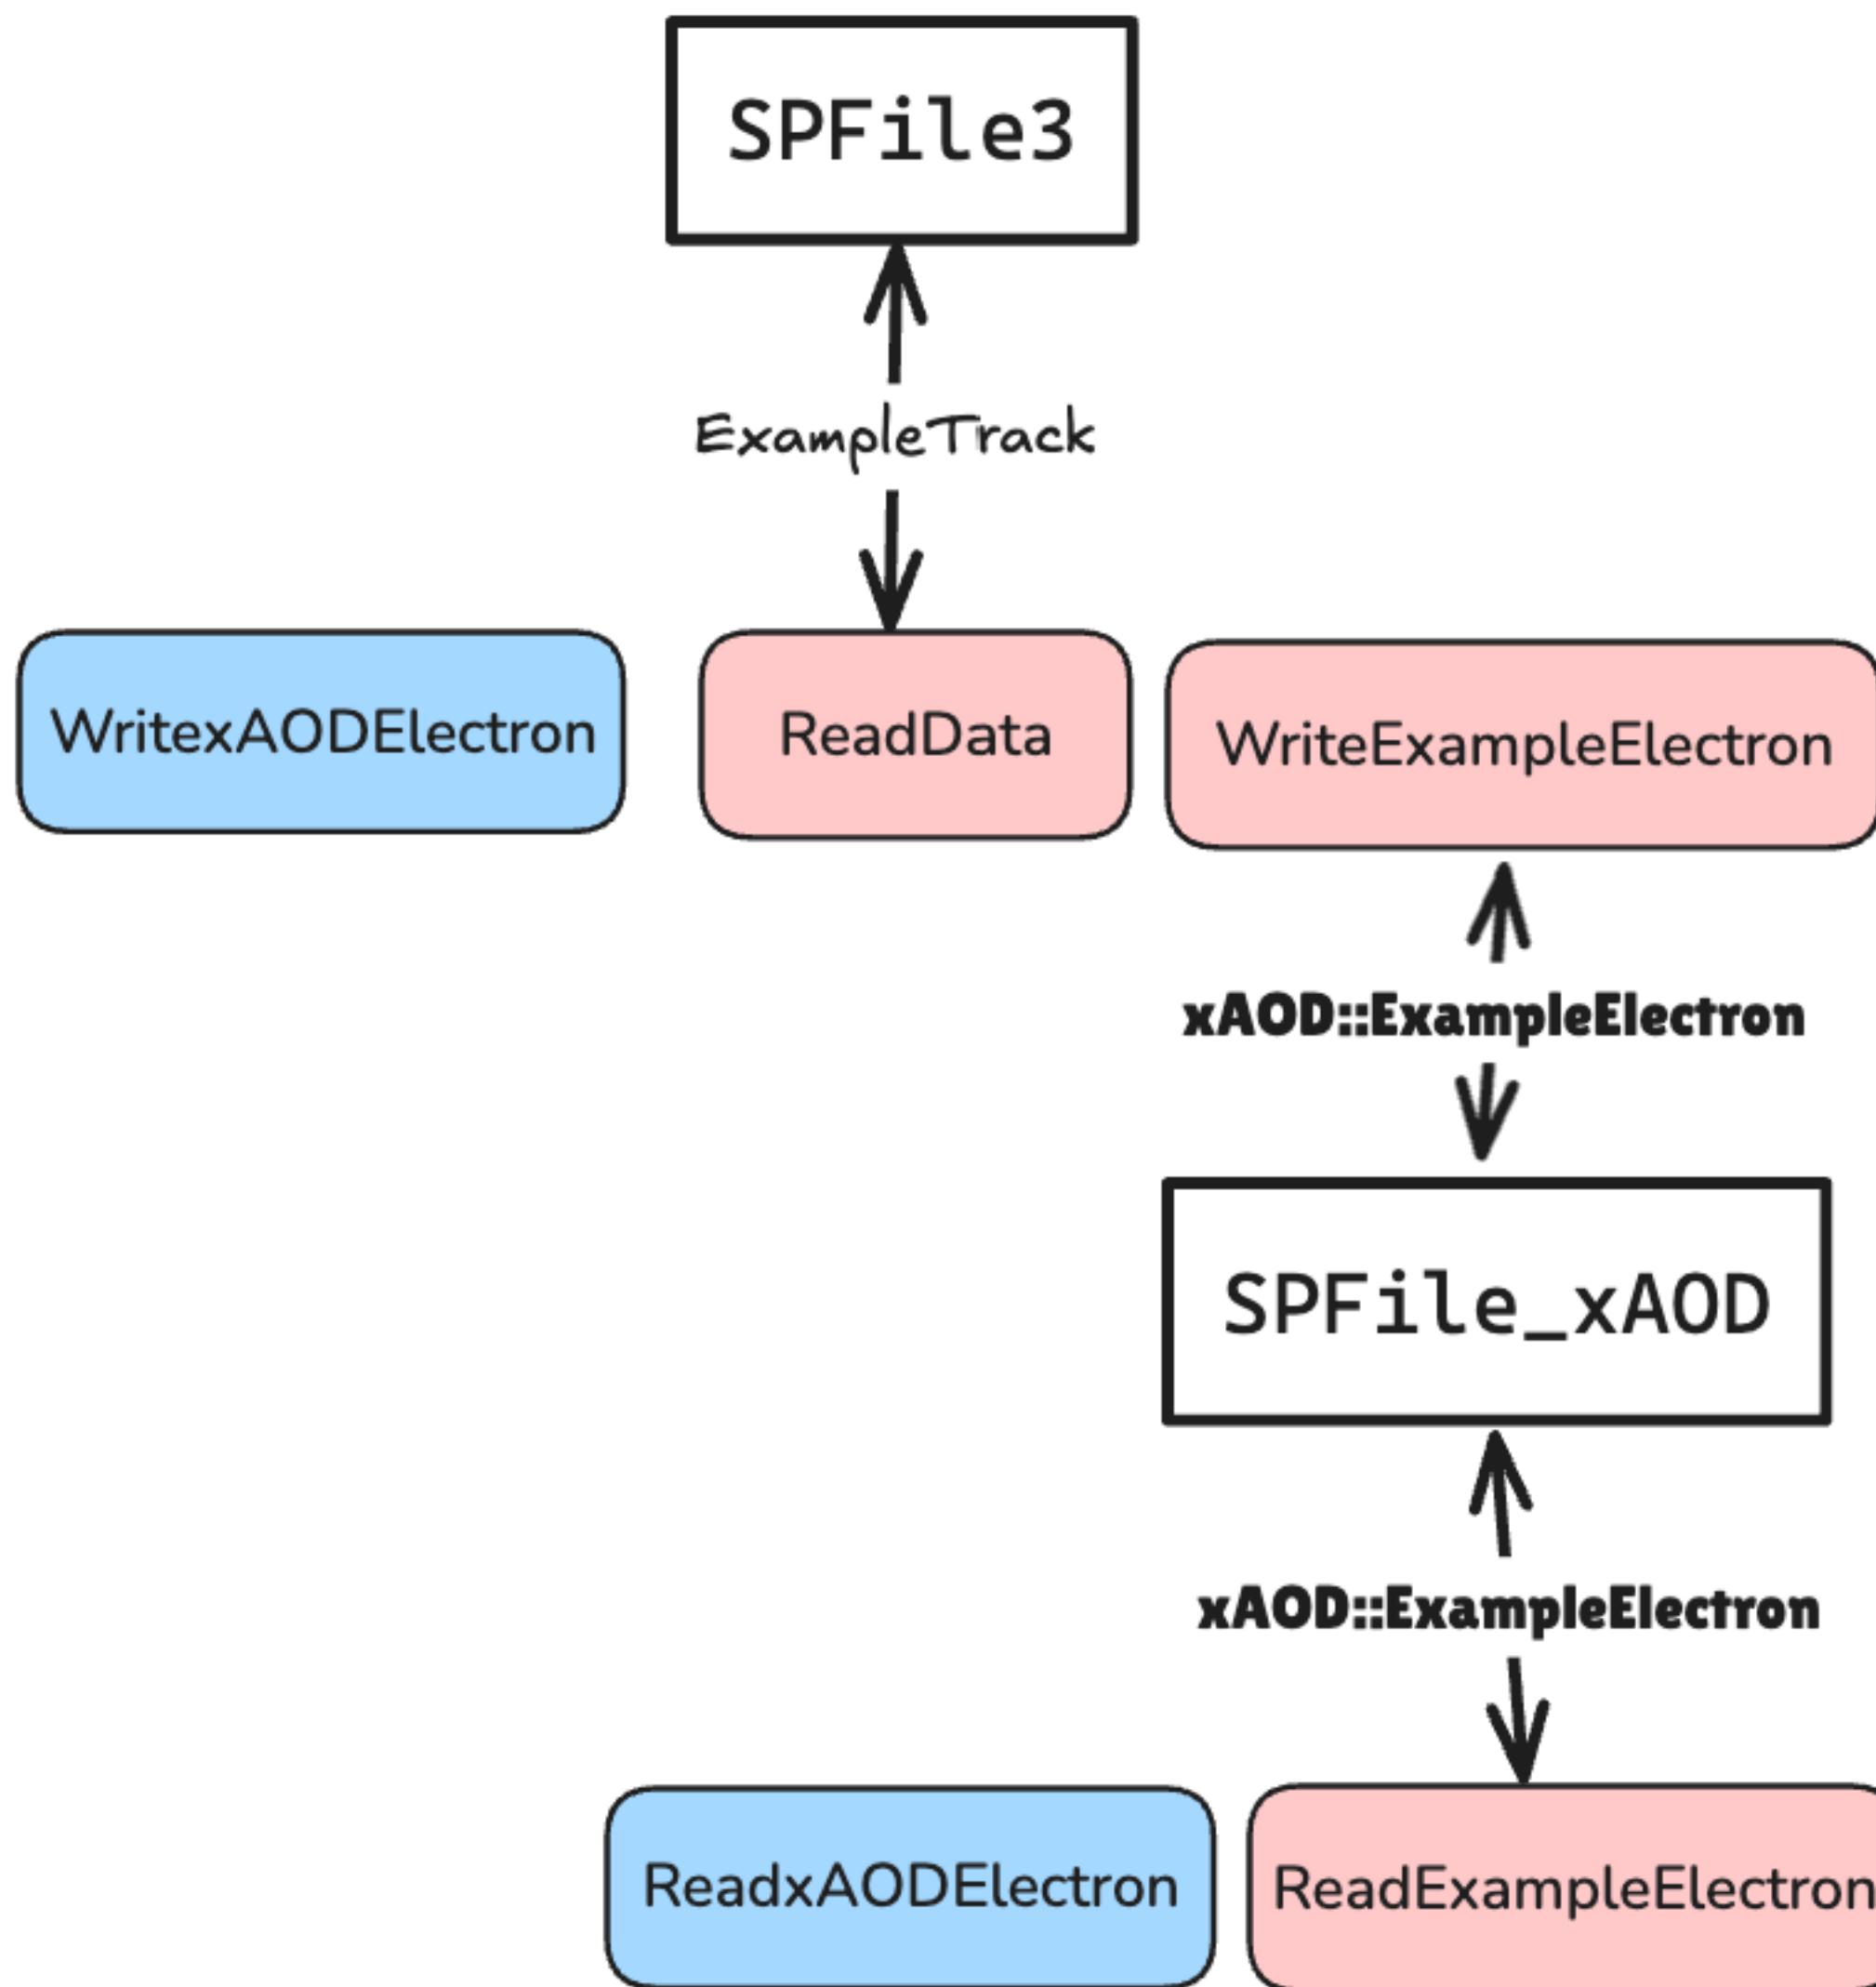
AthenaPoolExampleData
/src/



[\[LXR link\]](#)

[\[Athena Gitlab link\]](#)

Write/ReadxAODElectron



WriteExampleElectron

ExampleTracks

For every ExampleTrack,
Create an ExampleElectron object and
electron→setPt(track→getPT());
and save those electrons into a
container

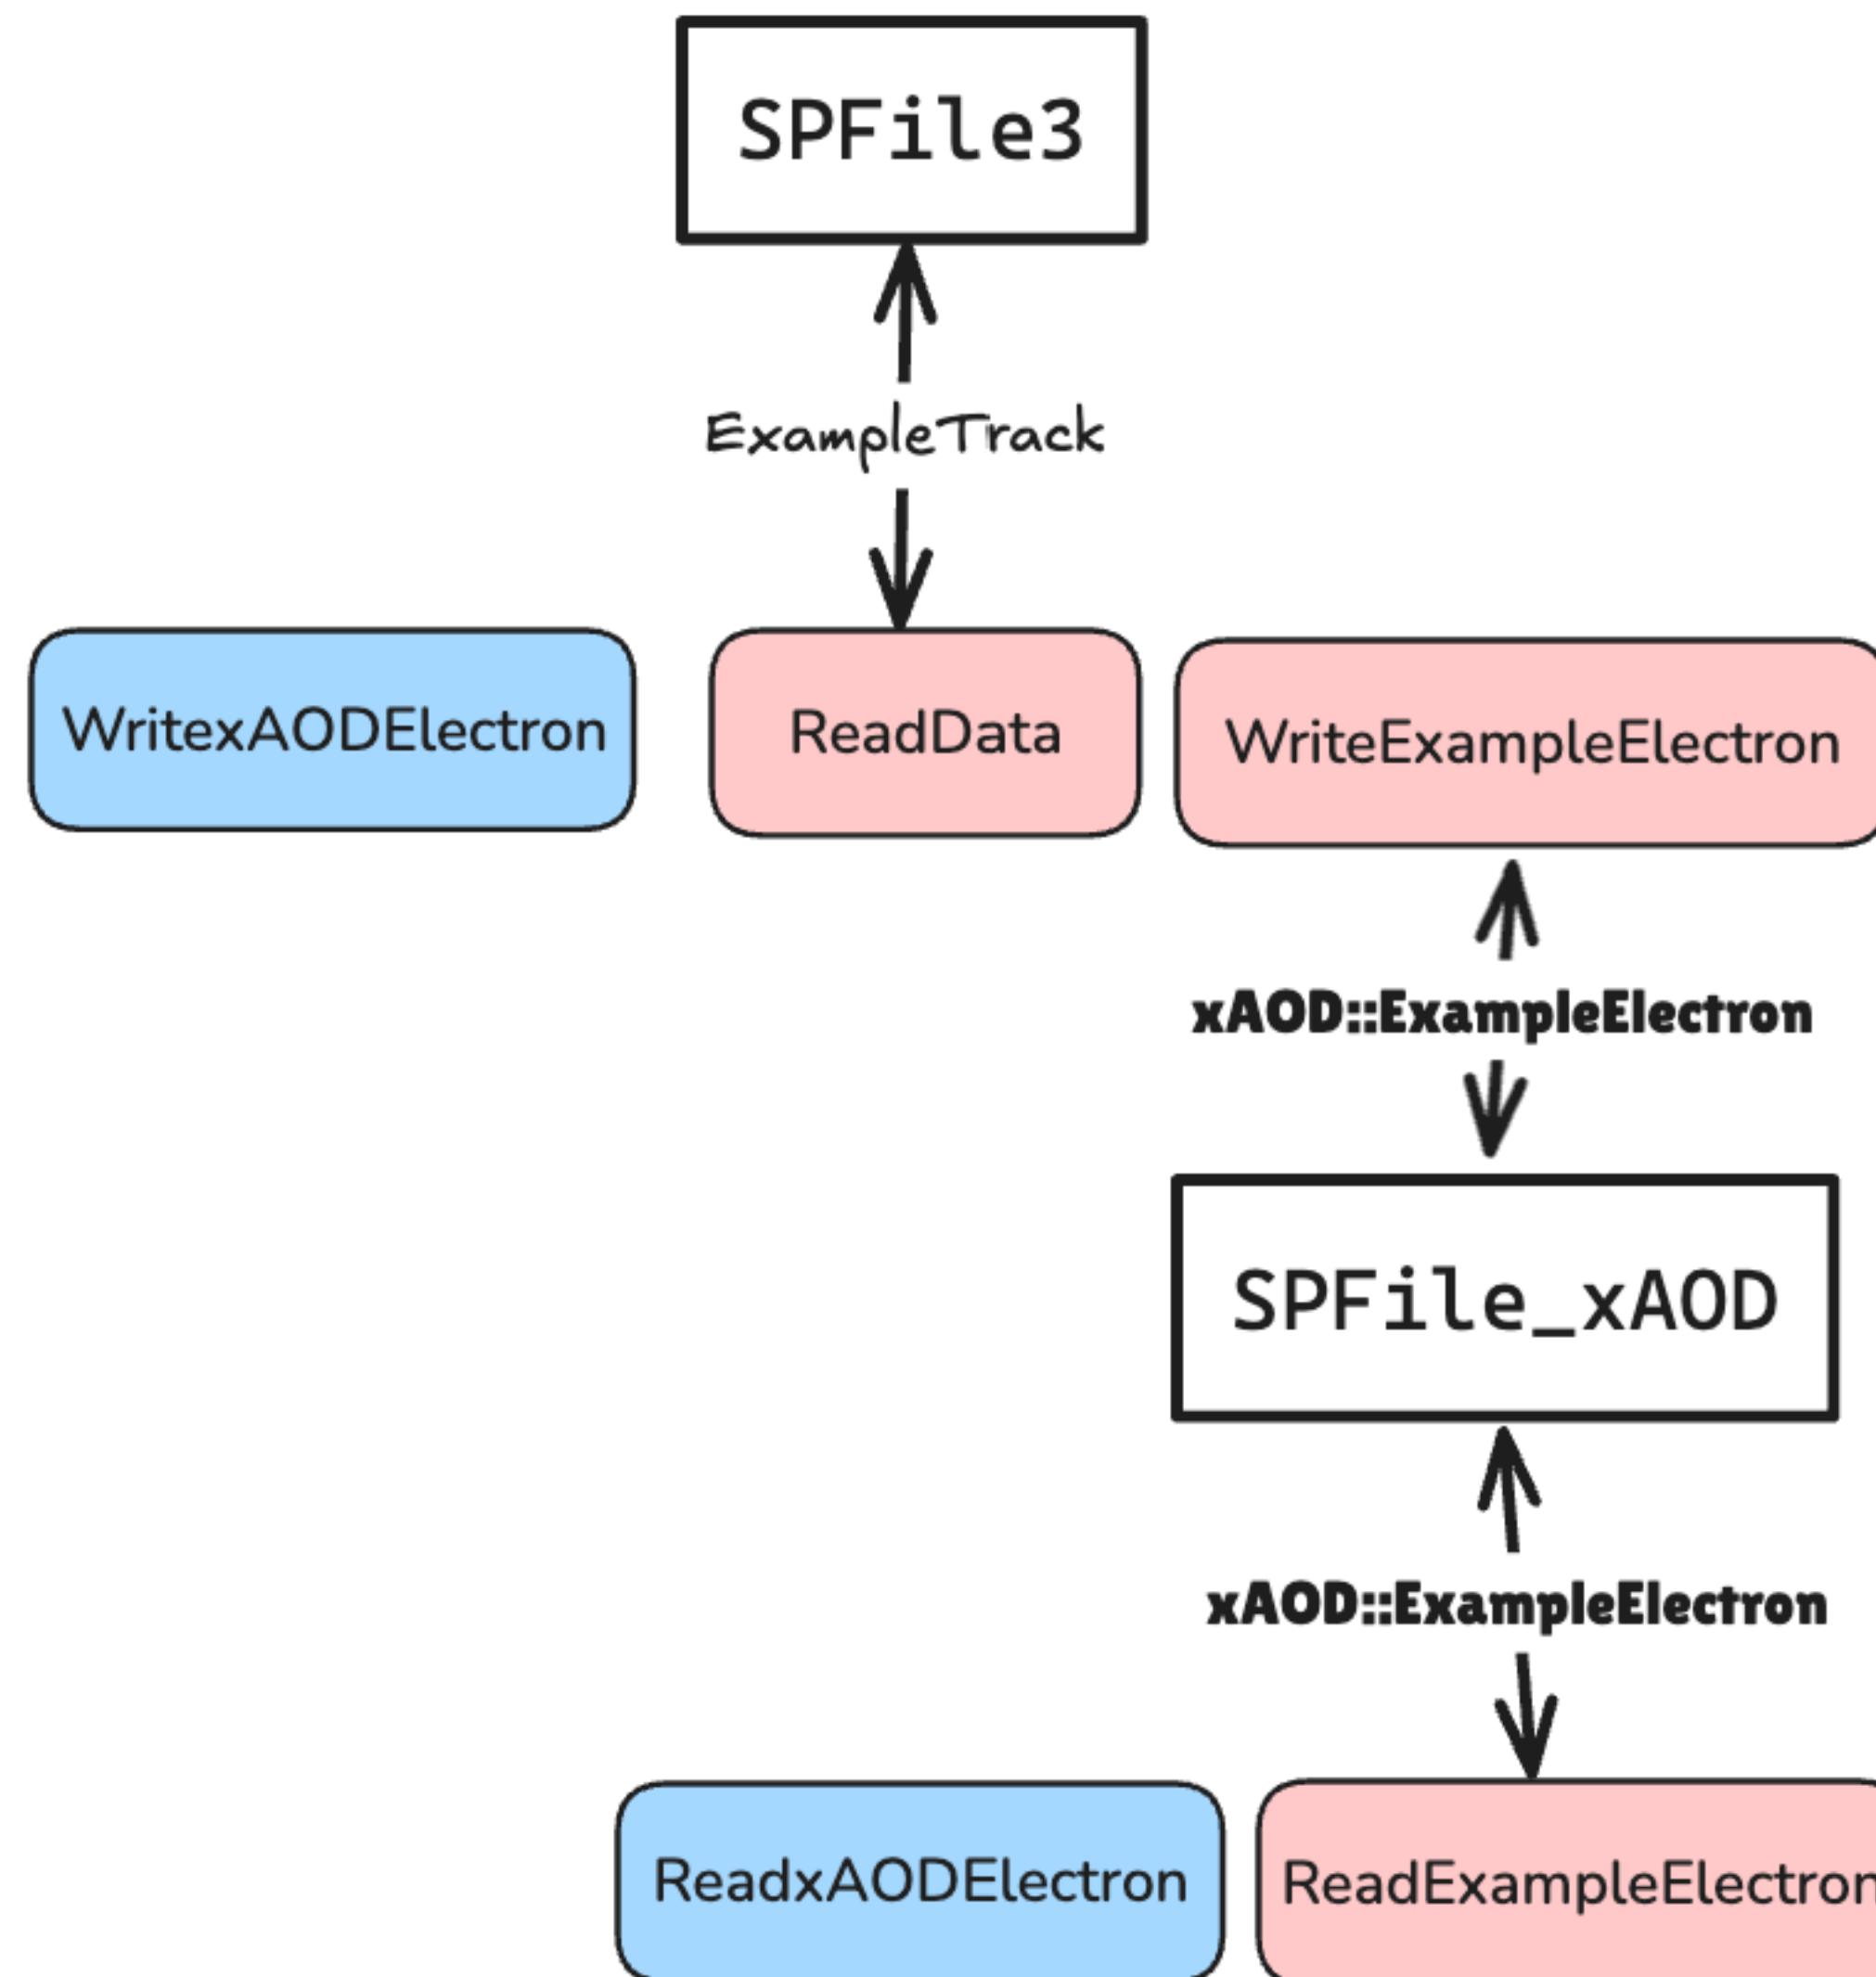
ExampleElectrons ←

For every ExampleElectron
access decor1 & decor2
set decor1 = 115.9 * index;
set decor2 = 114.9 + index;

ExampleElectrons.decor1
ExampleElectron.decor2 ←

Save this for later so we can test: access decor1 but not decor2.

Write/ReadxAODElectron



`ReadExampleElectron`

`ExampleElectrons`

For every `ExampleElectron`,
Print off every electron pt
& charge
Also see if `decor1` was written
print `decor1`

WritexAODElectron

```
# Create and attach the algorithms
acc.addEventAlgo( CompFactory.AthPoolEx.ReadData("ReadData", OutputLevel = DEBUG) )
acc.addEventAlgo( CompFactory.AthPoolEx.WriteExampleElectron("WriteExampleElectron", OutputLevel = DEBUG) )

# ----- Output Stream configuration
streamCA = OutputStreamCfg( flags, streamName, disableEventTag = True,
                            ItemList = [ "ExampleTrackContainer#MyTracks",
                                         "xAOD::ExampleElectronContainer#TestContainer",
                                         "xAOD::ExampleElectronAuxContainer#TestContainerAux.-decor2" ] )
```

```
# Create and attach the reading algorithm for the ExampleElectron and its decorations
acc.addEventAlgo( CompFactory.AthPoolEx.ReadExampleElectron("ReadExampleElectron",
                                                               OutputLevel = DEBUG) )
```

```
INFO WriteExampleElectron: track # 0 pT = 228.153.
INFO WriteExampleElectron: track # 0 is an electron with pT = 228.153;
INFO WriteExampleElectron: DecorationWriter: decor1 = 115.9, decor2 = 114.9.
```

⋮

```
INFO WriteExampleElectron: track # 0 pT = 2017.04.
INFO WriteExampleElectron: track # 0 is an electron with pT = 2017.04;
INFO WriteExampleElectron: DecorationWriter: decor1 = 115.9, decor2 = 114.9.
```

```
INFO ExampleElectron #0 {pt = 137.584, charge = 0}
INFO Decoration reader: decor1 = 115.9
```

```
INFO ReadExampleElectron: is executing ...
INFO ExampleElectron #0 {pt = 2017.04, charge = 0}
INFO Decoration reader: decor1 = 115.9
```

WritexAODElectron

```
# Create and attach the algorithms
acc.addEventAlgo( CompFactory.AthPoolEx.ReadData("ReadData", OutputLevel = DEBUG) )
acc.addEventAlgo( CompFactory.AthPoolEx.WriteExampleElectron("WriteExampleElectron", OutputLevel = DEBUG) )

# ----- Output Stream configuration
streamCA = OutputStreamCfg( flags, streamName, disableEventTag = True,
                            ItemList = [ "ExampleTrackContainer#MyTracks",
                                         "xAOD::ExampleElectronContainer#TestContainer",
                                         "xAOD::ExampleElectronAuxContainer#TestContainerAux.-decor2" ] )
```

ReadxAODElectron

```
# Create and attach the reading algorithm for the ExampleElectron and its decorations
acc.addEventAlgo( CompFactory.AthPoolEx.ReadExampleElectron("ReadExampleElectron",
                                                               OutputLevel = DEBUG) )

# Read
```

```
INFO WriteExampleElectron: track # 0 pT = 228.153.
INFO WriteExampleElectron: track # 0 is an electron with pT = 228.153;
INFO WriteExampleElectron: DecorationWriter: decor1 = 115.9, decor2 = 114.9.
```

⋮

```
INFO WriteExampleElectron: track # 0 pT = 2017.04.
INFO WriteExampleElectron: track # 0 is an electron with pT = 2017.04;
INFO WriteExampleElectron: DecorationWriter: decor1 = 115.9, decor2 = 114.9.
```

```
INFO ExampleElectron #0 {pt = 137.584, charge = 0}
INFO Decoration reader: decor1 = 115.9
```

⋮

```
INFO ReadExampleElectron: is executing ...
INFO ExampleElectron #0 {pt = 2017.04, charge = 0}
INFO Decoration reader: decor1 = 115.9
```

CI tests merged with Athena!

The screenshot shows a GitHub merge request page for a pull request titled "AthenaPoolExample xAOD EDM Read/Write CI build IO tests". The status bar at the top indicates the pull request is "Merged". The main content area contains a detailed description of the changes, including code snippets and JIRA ticket links. On the right side, there are sections for assignees, reviewers, labels, milestones, time tracking, and participants. A green circle highlights the "Merged" status in the top left.

AthenaPoolExample xAOD EDM Read/Write CI build IO tests

Merged Arthur Charles Kraus requested to merge [akraus/athena_old:jacks-](#) into [main](#) 5 months ago

Overview 81 Commits 37 Pipelines 15 Changes 30 All threads resolved! Add a to-do item

There are a number of CI build I/O tests in the Database/AthenaPool/AthenaPoolExample/AthenaPoolExampleAlgorithms that we would like to test more of the modern I/O functionality.

This MR adds in two tests within the package that read and write bespoke xAOD objects called ExampleElectrons that have a charge and pT. The test AthenaPoolExample_WritexAODElectrons.py takes existing pT values from ExampleTracks, which are generated from the AthenaPoolExample_ReadWrite.py test, and writes them to the xAOD::ExampleElectron object which is then put into TestContainer and saved to SimplePoolFile_xAOD.root. The AthenaPoolExample_ReadxAODElectrons.py test then takes that SimplePoolFile_xAOD.root and loops over all the ExampleElectrons and reads back the pT value.

These tests also write decorations to the corresponding ExampleElectrons, where, upon reading ignore the second decoration decor2 and only read the first decor1.

JIRA ticket: <https://its.cern.ch/jira/browse/ATEAM-959>

Edited 4 months ago by Arthur Charles Kraus

0 0

Pipeline #8420285 passed

Pipeline passed for 0863166f on [akraus:jacks-test-b...](#) 4 months ago

8 Approval is optional Assign reviewers

0 Assignees Edit None - assign yourself

0 Reviewers Edit Approval is optional Assign None - assign yourself

Labels Edit analysis-review-approved Database EDM Egamma main review-approved

Milestone Edit None

Time tracking No estimate or time spent

13 Participants + 5 more

Thanks!



Virginia Jamarillo, "Zero Point", 2023

Backup Slides:



Virginia Jamarillo, "Zero Point", 2023

AthenaPoolExample_WritexAODElectron.py

```
from AthenaConfiguration.AllConfigFlags import initConfigFlags
from AthenaConfiguration.ComponentFactory import CompFactory
from AthenaCommon.Constants import DEBUG
from OutputStreamAthenaPool.OutputStreamConfig import OutputStreamCfg, outputStreamName

streamName = "WritexAODElectrons"
outputFileName = "SimplePoolFile_xAOD.root"

# Setup flags
flags = initConfigFlags()
flags.Input.Files = ["SimplePoolFile3.root"]
flags.addFlag(f"Output.{streamName}FileName", outputFileName)
flags.Exec.MaxEvents = -1
flags.Common.MsgSuppression = False
flags.Exec.DebugMessageComponents = [outputStreamName(streamName),
                                         "PoolSvc", "AthenaPoolCnvSvc", "AthenaPoolAddressProviderSvc", "MetaDataSvc"]
flags.lock() # Setup flags

# Main services
from AthenaConfiguration.MainServicesConfig import MainServicesCfg
acc = MainServicesCfg( flags )

# Pool reading and writing
from AthenaPoolExampleAlgorithms.AthenaPoolExampleConfig import AthenaPoolExampleReadCfg, AthenaPoolExampleWriteCfg
acc.merge( AthenaPoolExampleReadCfg(flags, readCatalogs = ["file:Catalog1.xml"]) )
acc.merge( AthenaPoolExampleWriteCfg( flags, streamName,
                                         writeCatalog = "file:Catalog1.xml" ) )

# Main services
# Create and attach the algorithms
from AthenaConfiguration.MainServicesConfig import MainServicesCfg
acc.addEventAlgo( CompFactory.AthPoolEx.ReadData("ReadData", OutputLevel = DEBUG) )
acc.addEventAlgo( CompFactory.AthPoolEx.WriteExampleElectron("WriteExampleElectron", OutputLevel = DEBUG) )

# ----- Output Stream configuration
streamCA = OutputStreamCfg( flags, streamName, disableEventTag = True,
                            ItemList = [ "ExampleTrackContainer#MyTracks",
                                         "xAOD::ExampleElectronContainer#TestContainer",
                                         "xAOD::ExampleElectronAuxContainer#TestContainerAux.-decor2" ] )
stream = streamCA.getEventAlgo( outputStreamName( streamName ) )
acc.merge( streamCA )

#-----#
# Set output level threshold (2=DEBUG, 3=INFO, 4=WARNING, 5=ERROR, 6=FATAL)
#-----#
stream.WritingTool.OutputLevel = 3
stream.HelperTools[0].OutputLevel = 3

# Run
import sys
sc = acc.run(flags.Exec.MaxEvents)
sys.exit(sc.isFailure())
```

AthenaPoolExample_ReadxAODElectron.py

```
from AthenaConfiguration.AllConfigFlags import initConfigFlags
from AthenaConfiguration.ComponentFactory import CompFactory
from AthenaCommon.Constants import DEBUG

# Setup flags
flags = initConfigFlags()
flags.Input.Files = [ "SimplePoolFile_xAOD.root" ]
flags.Exec.MaxEvents = -1
flags.Common.MsgSuppression = False
flags.Exec.DebugMessageComponents = [ "ReadExampleElectron", "PoolSvc", "AthenaPoolCnvSvc",
                                         "AthenaPoolAddressProviderSvc", "MetaDataSvc", "EventSelector" ]
flags.lock()

# Main services
from AthenaConfiguration.MainServicesConfig import MainServicesCfg
acc = MainServicesCfg( flags )

# Configure AthenaPool reading
from AthenaPoolExampleAlgorithms.AthenaPoolExampleConfig import AthenaPoolExampleReadCfg
acc.merge( AthenaPoolExampleReadCfg(flags, readCatalogs = ["file:Catalog1.xml"]) )

# Create and attach the reading algorithm for the ExampleElectron and its decorations
acc.addEventAlgo( CompFactory.AthPoolEx.ReadExampleElectron("ReadExampleElectron",
                                                               OutputLevel = DEBUG) )

# Run
import sys
sc = acc.run(flags.Exec.MaxEvents)
sys.exit(sc.isFailure())
~
```

WriteExampleElectron.h

```
#ifndef ATHENAPOOLEXAMPLEALGORITHMS_WRITEEXAMPLEELECTRON_H
#define ATHENAPOOLEXAMPLEALGORITHMS_WRITEEXAMPLEELECTRON_H

/** @file AthenaPoolExampleAlgorithms/src/WriteExampleElectron.h
 *  @brief This file contains the class definition for the WriteExampleElectron
 *  class.
 */
#include "AthenaBaseComps/AthReentrantAlgorithm.h"
#include "AthenaPoolExampleData/ExampleElectronContainer.h"
#include "AthenaPoolExampleData/ExampleHitContainer.h"
#include "AthenaPoolExampleData/ExampleTrackContainer.h"
#include "StoreGate/WriteDecorHandleKey.h"
#include "StoreGate/WriteHandleKey.h"

namespace AthPoolEx {

/** @class AthPoolEx::WriteExampleElectron
 *  @brief This class provides an example for writing event data objects to
 *  *Pool.
 */
class WriteExampleElectron : public AthReentrantAlgorithm {
public: // Constructor and Destructor
    /// Standard Service Constructor
    WriteExampleElectron(const std::string& name, ISvcLocator* pSvcLocator);
    /// Destructor
    virtual ~WriteExampleElectron() = default;

public:
    /// Gaudi Service Interface method implementations:
    virtual StatusCode initialize() override final;
    virtual StatusCode execute(const EventContext& ctx) const override final;
    virtual StatusCode finalize() override final;

    // Read in ExampleTracks and ExampleHits
    SG::ReadHandleKey<ExampleTrackContainer> m_exampleTrackKey{this, "ExampleTrackKey", "MyTracks"};
    SG::ReadHandleKey<ExampleHitContainer> m_exampleHitKey{this, "ExampleHitKey", "MyHits"};

    // We will want to write some of those tracks as an ExampleElectron and place
    // them into a container
    SG::WriteHandleKey<xAOD::ExampleElectronContainer>
        m_exampleElectronContainerKey{this, "ExampleElectronContainerName",
                                      "TestContainer"};
};

#endif // ATHENAPOOLEXAMPLEALGORITHMS_WRITEEXAMPLEELECTRON_H
```

```
// Testing writing decorations, need keys
SG::WriteDecorHandleKey<xAOD::ExampleElectronContainer> m_decor1Key{
    this, "ExampleElectronContainerDecorKey1", "TestContainer.decor1",
    "decorator1 key"};
SG::WriteDecorHandleKey<xAOD::ExampleElectronContainer> m_decor2Key{
    this, "ExampleElectronContainerDecorKey2", "TestContainer.decor2",
    "decorator2 key"};
}

// Read in ExampleTracks and ExampleHits
#endif // ATHENAPOOLEXAMPLEALGORITHMS_WRITEEXAMPLEELECTRON_H
```

WriteExampleElectron.cxx

```
#include "WriteExampleElectron.h"

// the user data-class definitions
#include "AthenaPoolExampleData/ExampleElectron.h"
#include "AthenaPoolExampleData/ExampleElectronAuxContainer.h"
#include "AthenaPoolExampleData/ExampleElectronContainer.h"
#include "AthenaPoolExampleData/ExampleHitContainer.h"
#include "AthenaPoolExampleData/ExampleTrackContainer.h"
#include "GaudiKernel/EventContext.h"
#include "StoreGate/WriteDecorHandle.h" // for WriteExampleElectron.h
#include "StoreGate/WriteHandle.h" class definition for the WriteExampleElectron

using namespace AthPoolEx;
// _____
// WriteExampleElectron.h
// _____
// WriteExampleElectron(const std::string& name,
//                      ISvcLocator* pSvcLocator)
// : AthReentrantAlgorithm(name, pSvcLocator) {}

// _____
StatusCode WriteExampleElectron::initialize() {
    ATH_MSG_INFO(name() << ": in initialize()");
    // Check that all services needed for writing event data objects to
    ATH_CHECK(m_exampleElectronContainerKey.initialize());
    ATH_CHECK(m_exampleTrackKey.initialize());
    ATH_CHECK(m_exampleHitKey.initialize());
    ATH_CHECK(m_decor1Key.initialize());
    ATH_CHECK(m_decor2Key.initialize());
    ATH_CHECK(m_decor_floatKey.initialize());
    ATH_CHECK(m_decor_longdoubleKey.initialize());

    return StatusCode::SUCCESS;
} // void service interface method implementations:
    virtual StatusCode initialize() override final;
    virtual StatusCode execute(const EventContext& ctx) const override final;

// _____
StatusCode WriteExampleElectron::execute(const EventContext& ctx) const {
    ATH_MSG_DEBUG("WriteExampleElectron in execute()");
    std::size_t idx_trk = 0;
    SG::ReadHandle<ExampleTrackContainer> trackCont(m_exampleTrackKey);
    std::size_t MAX = 100;
    float idx_decor = 0.0f;
    SG::ReadHandle<ExampleHitContainer> hitCont(m_exampleHitKey);
    float trackPt = 0.0f;

    auto elecCont = std::make_unique<xAOD::ExampleElectronContainer>();
    auto elecStore = std::make_unique<xAOD::ExampleElectronAuxContainer>();
    elecCont->setStore(elecStore.get());
```

```
/*
 * Convert ExampleTrack to xAOD::ExampleElectron
 */
SG::ReadHandle<ExampleTrackContainer> trackCont(m_exampleTrackKey, ctx);
elecCont->push_back(std::make_unique<xAOD::ExampleElectron>());

for (const ExampleTrack* track : *trackCont) {
    trackPt = track->getPT();
    ATH_MSG_INFO(name() << ": track # " << idx_trk << " pT = " << trackPt
                 << ".");
}

// Push 10 electrons
for(size_t idx = 0; idx < MAX; ++idx) {
    // proceed to take this tracks pT
    elecCont->back()->setPt(trackPt);
    ATH_MSG_INFO(name() << ": idx: " << idx << " Pt= " << elecCont->back()->pt()
                 << ";");

}

// WriteExampleElectron::writeExampleElectron(const std::string& name,
//                                             ISvcLocator* pSvcLocator)
// : AthReentrantAlgorithm(name, pSvcLocator) {}

// Print out the pT that's being saved
ATH_MSG_INFO(name() << ": track # " << idx_trk
             << " is an electron with pT = " << elecCont->back()->pt()
             << ";");
idx_trk++;
}
ATH_CHECK(m_exampleElectronContainerKey.initialize());
ATH_CHECK(m_exampleTrackKey.initialize());
ATH_CHECK(m_exampleHitKey.initialize());
ATH_CHECK(m_decor1Key.initialize());
ATH_CHECK(m_decor2Key.initialize());
ATH_CHECK(m_decor_floatKey.initialize());
ATH_CHECK(m_decor_longdoubleKey.initialize());

SG::WriteHandle<xAOD::ExampleElectronContainer> objs(
    m_exampleElectronContainerKey, ctx);

ATH_CHECK(objs.record(std::move(elecCont), std::move(elecStore)));

// _____
// StatusCode WriteExampleElectron::execute(const EventContext& ctx) const {
/*
 * ATH_MSG_DEBUG("WriteExampleElectron in execute()");
 * Writing Decorations
 */
std::size_t MAX = 100;
float idx_decor = 0.0f;
// WriteDecorHandle for the decoration 'TestContainer.decor1'
SG::WriteDecorHandle<xAOD::ExampleElectronContainer, float> hdl1(m_decor1Key,
                     ctx);
// And for the second decoration 'TestContainer.decor2'
SG::WriteDecorHandle<xAOD::ExampleElectronContainer, float> hdl2(m_decor2Key,
                     ctx);
```